

Firewall = Which stops unauthorized access to the Network **Allow / Deny**
Load Balancer = Which distribute the traffic across multiple Servers

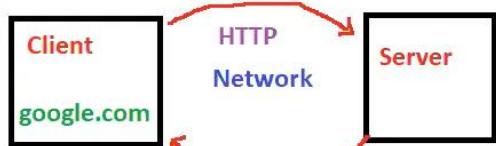
http = 80 / Not Secure connection to the Website
https = 443 / Secure connection to the website
SSH = 22 / Connect to Linux Machines
RDP = 3389 / Connect to Windows Machines

Browser ---> Local DNS ---> Root Name Server ---> Top Level Domain ----> Name Server ---> SOA
Browser ---> Firewall ---> Load Balancer ---> WebServer ---> App Server ---> DB Server

HTTP = HyperText Transfer Protocol

HTTP default Port Number is 80

For Customer it should be always 80 or 443



HTTP transfer the data to and fro from the browser to Server

HTTPS = 443

certificates

SSL / TLS / HTTPS

Data is Encrypted

HTTP is Not Secure

HTTPS is Secure

Status Codes

404 = Page Not Found

503 = Service Unavailable

500 = Internal Server Error

200 = Page Found / Success



HTTP

Plain Text / Clear Text

<http://192.168.10.20:80>

<http://192.168.10.20:8080>

<http://192.168.10.20>

Application

8080

Server

192.168.10.20

HTTP Over TCP/IP

Transmission Control Protocol

Ameerpet

Server 1

HOST

Miyapur

Server 2

HOST

Train
HTTP
TCP
(track)
IP

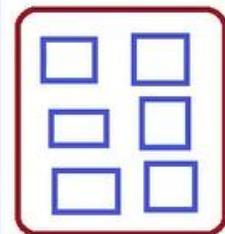
TCP establish the connection between 2 hosts

TCP is like a bridge / Messenger

TCP is reliable

UDP is not reliable

User Datagram Protocol



Store Room



Bharath (Operations Guys)

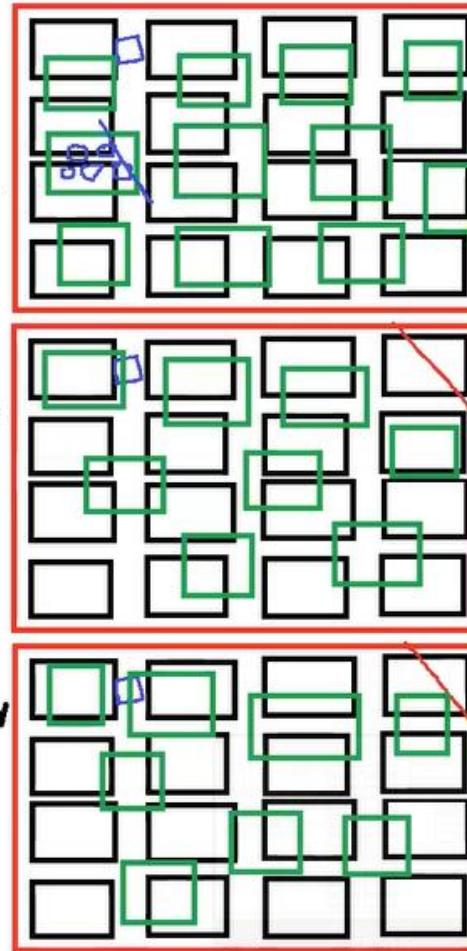
5
10
100
500
5000



Pravalika
(Business)

Data-Center
On-Premises
150 Racks
20 Racks

50



Traditional Architecture

Physical Machine

Virtualization

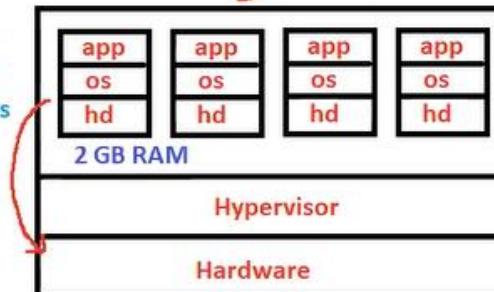
100 Machines

5 Big machines

95 are VM's

P2V Migration

V2C Migration



Host Machine

Bare Metal Virtualization

AWS



Host Based Virtualization

AWS has Global Infrastructure

AWS is providing Infrastructure as a Service

Cloud is present in the Remote Location

Remote Location = DataCenters

DataCenter = Infrastructure

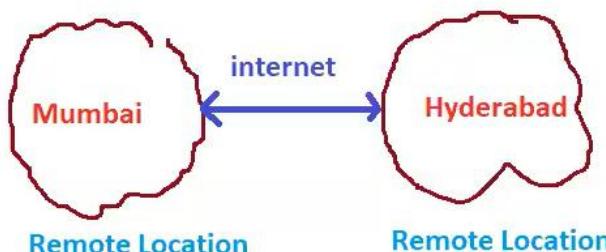
Infrastructure = Servers, DB's, Network etc

We need internet to connect to the Cloud

AWS is a Cloud Provider, who provides Infrastructure as a Service

**Amazon Web Services
(AWS Management Console)**

AWS is a Group of Services



AWS is Public Cloud Provider, who provides infrastructure as a Service

Key-Words

Virtualization, Host Machine, DNS, VM's, Infrastructure, DataCenters, Protocols, Load Balancer, Firewall, Hypervisor, Cloud, Remote Location.

Physical DC ----> Virtualization ----> Cloud --> AWS (Remote Location(DataCenter))

Cloud Computing

Instead of doing computing on local machine / On-premises, you will be now doing computing in the remote location (**Cloud**) that is called **Cloud Computing**

Deployment Models (Types of Cloud)

Public Cloud : The Providers Services which are accessed by everyone like AWS, Azure, GCP

Private Cloud : The Providers services which are accessed within the organization like Oracle

Hybrid Cloud : The combination of Public and Private Cloud

Service Models

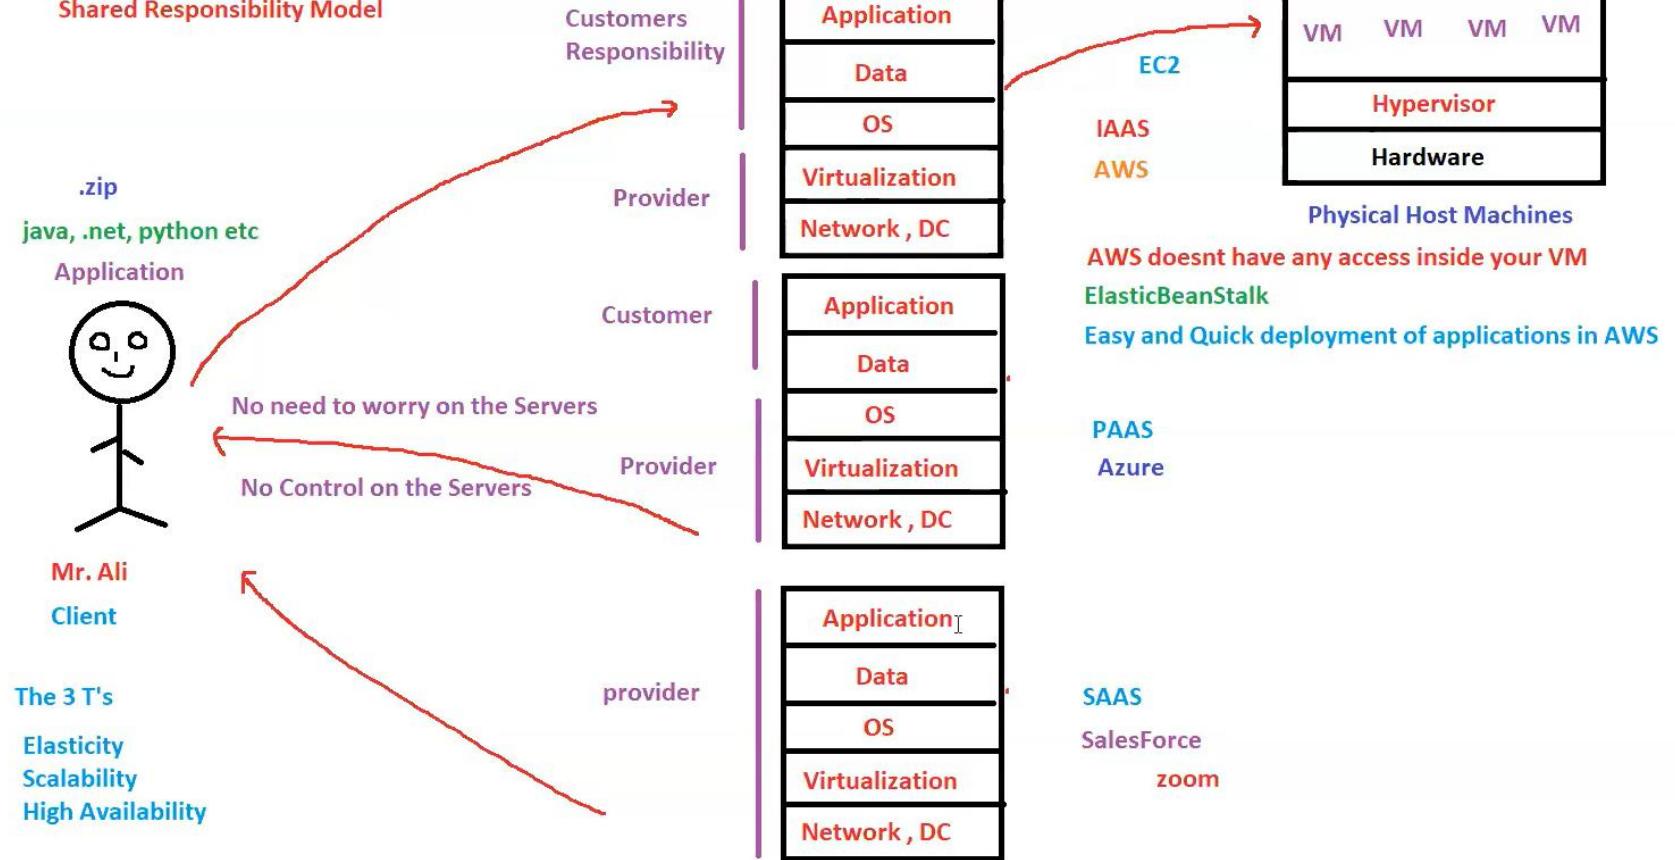
Infrastructure as a Service (IAAS)

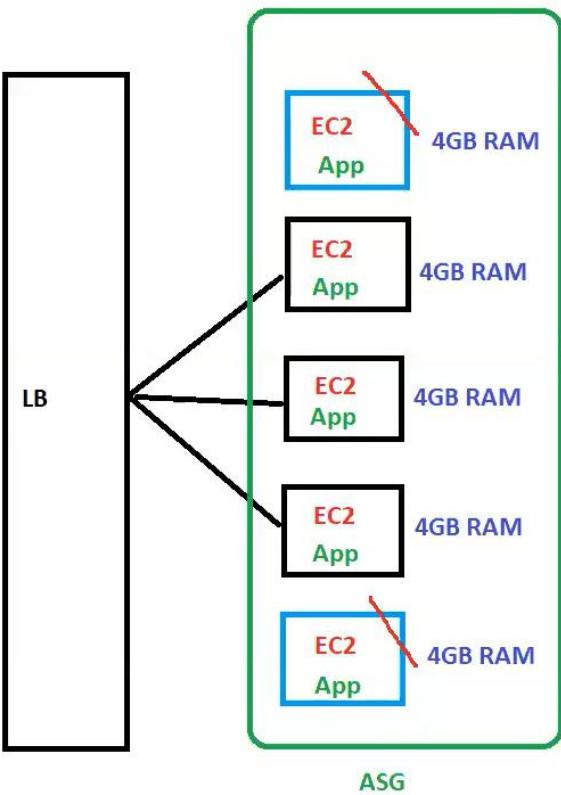
Platform as a Service (PAAS)

Software as a Service (SAAS)

AWS Management Console
AWS is a Group of Services
Elastic Compute Cloud
VM's = Instances
EC2 Instances
EC2 is a AWS Service where you can create VM's

Shared Responsibility Model





Elasticity

Increasing or decreasing the number of servers based on the demand

Elasticity is Short Term

Elasticity can be achieved in AWS using Auto-Scaling

Auto-Scaling = Scale Out and Scale In

increasing
adding decreasing
removing

Elasticity is also called as Horizontal Scaling

Use the same capacity of the Servers in AutoScaling Group

Scalability

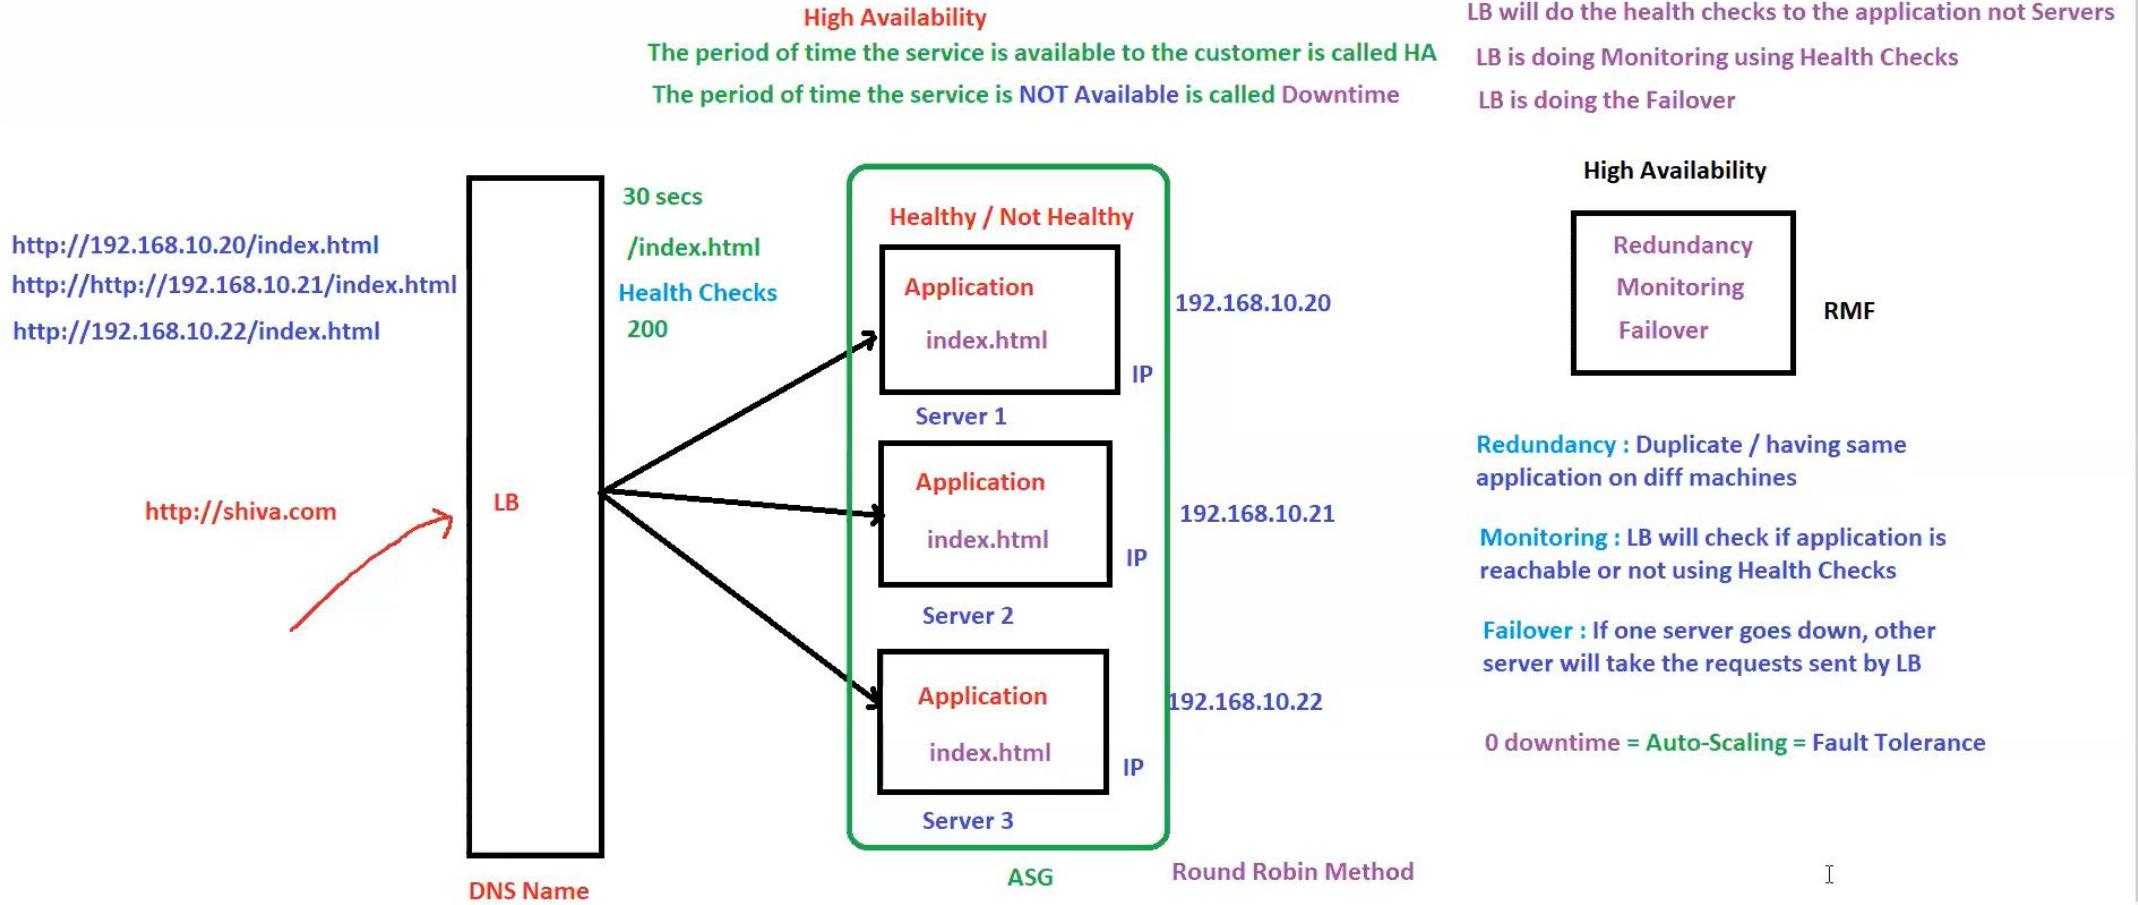


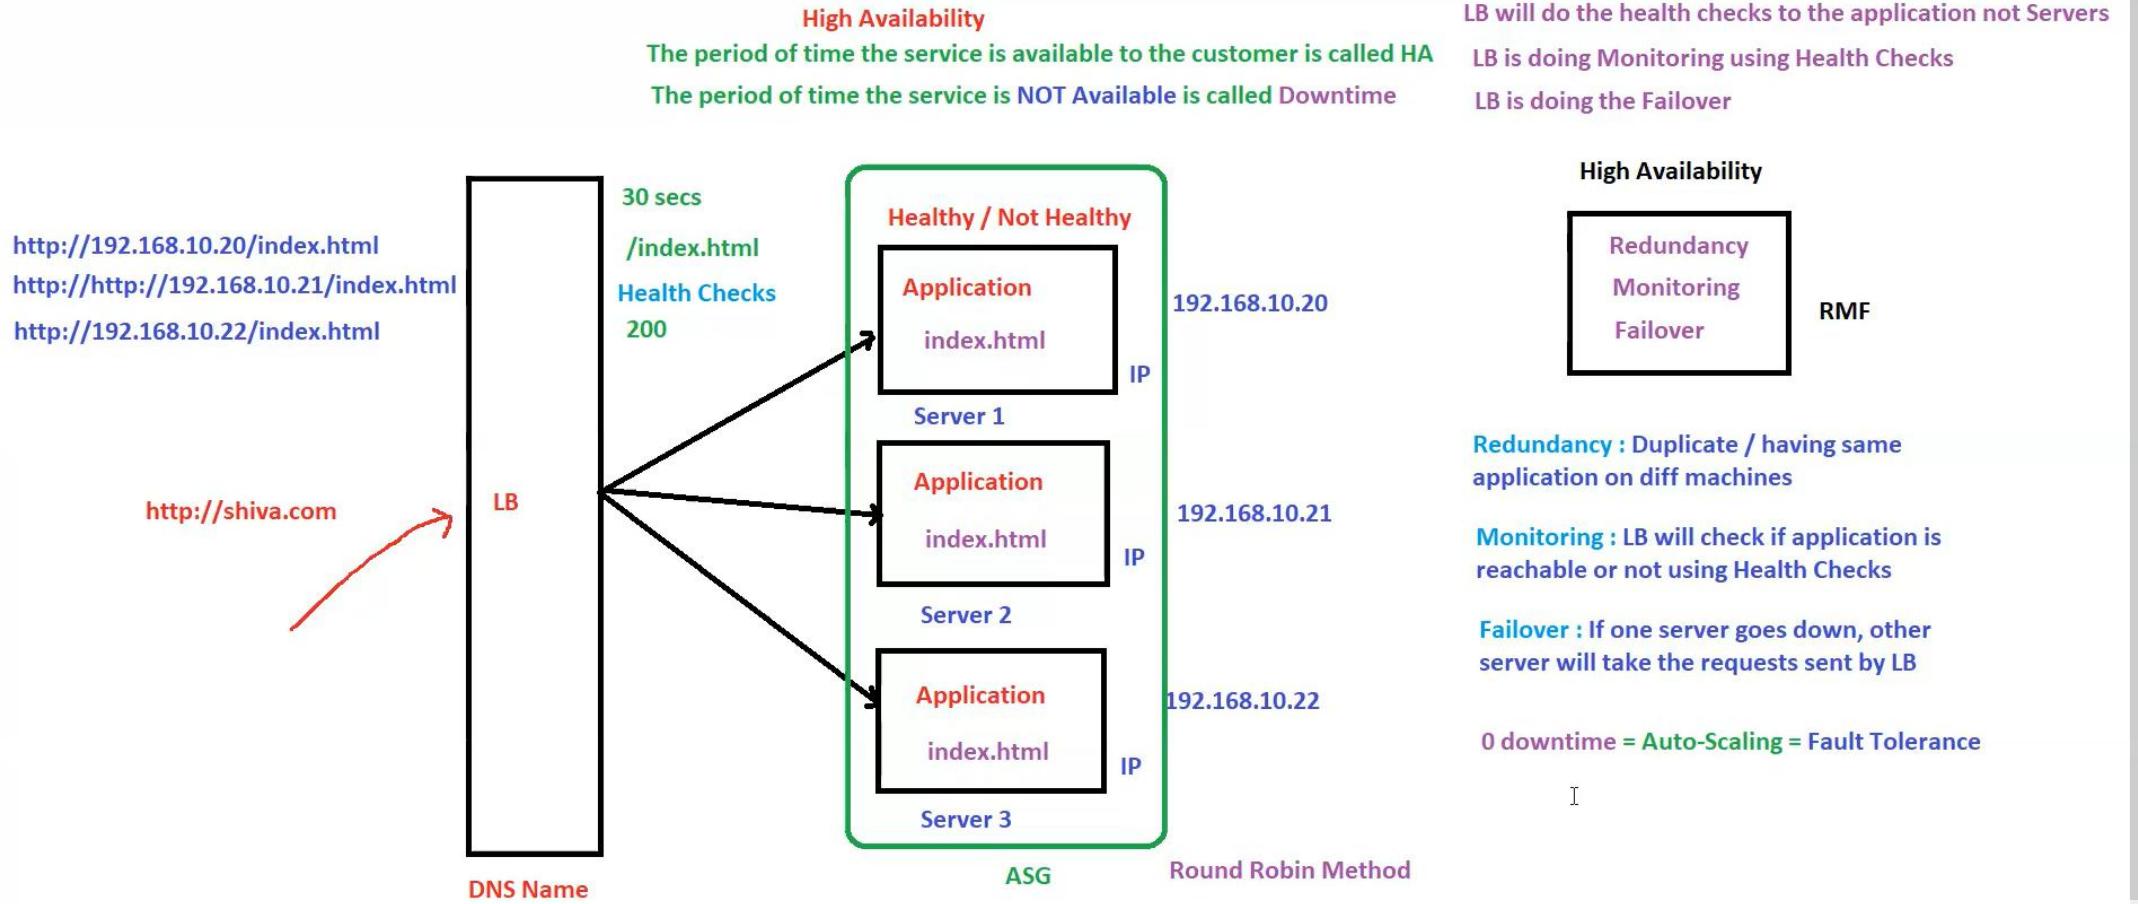
Increasing the Capacity of the Servers is called Scalability

Scalability = Scale Up and Scale Down

Scalability is Long Term

Scalability is also called as Vertical Scaling





AWS has Global Infrastructure

A Region is a place where AWS has its Infrastructure

A Region Contains multiple DataCenters

A Region has multiple AZ's

Servers = Instances

Servers / Instances are placed in AZ's

Best Practise is to distribute the instances across multiple AZ's

very very less chance that 1 AZ does down

1a or 1b or 1c = Group of DC's

1AZ is a group of DC's

Instances across AZ's can share the data if required as AZ's network are inter-connected

1a

1b

1c

dadar, andheri, aroli

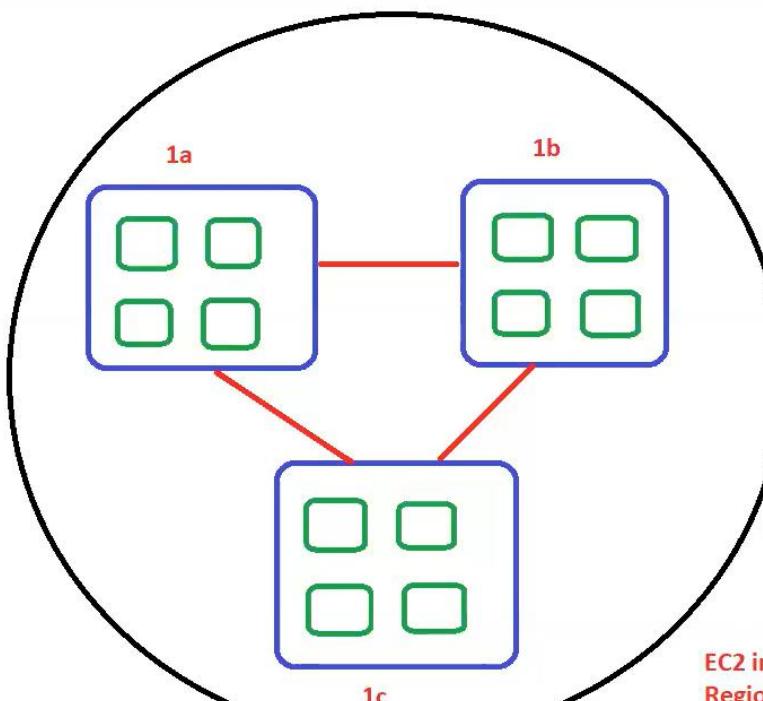
thane, bandra, kalyan

juhu, kurla, vashi

Regions and Availability Zones

Region = Its a Geo-Graphical area, Ex AWS-Region = Mumbai

Availability Zone = Simply a DataCenter (AZ)



AWS - Region = Mumbai

Mumbai = ap-south-1

AZ's = ap-south-1a
ap-south-1b
ap-south-1c

Regions and AZ's are managed by AWS

AZ's communicate with each other by default

AZ's network are inter-connected

There are 32 Regions in AWS at the moment

Low Latency = Good

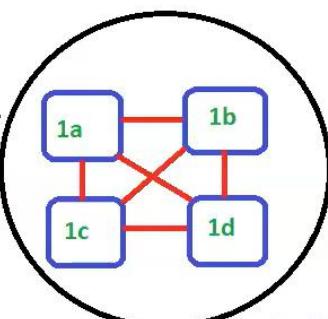
High Latency = Bad

Latency is depend upon the Response time

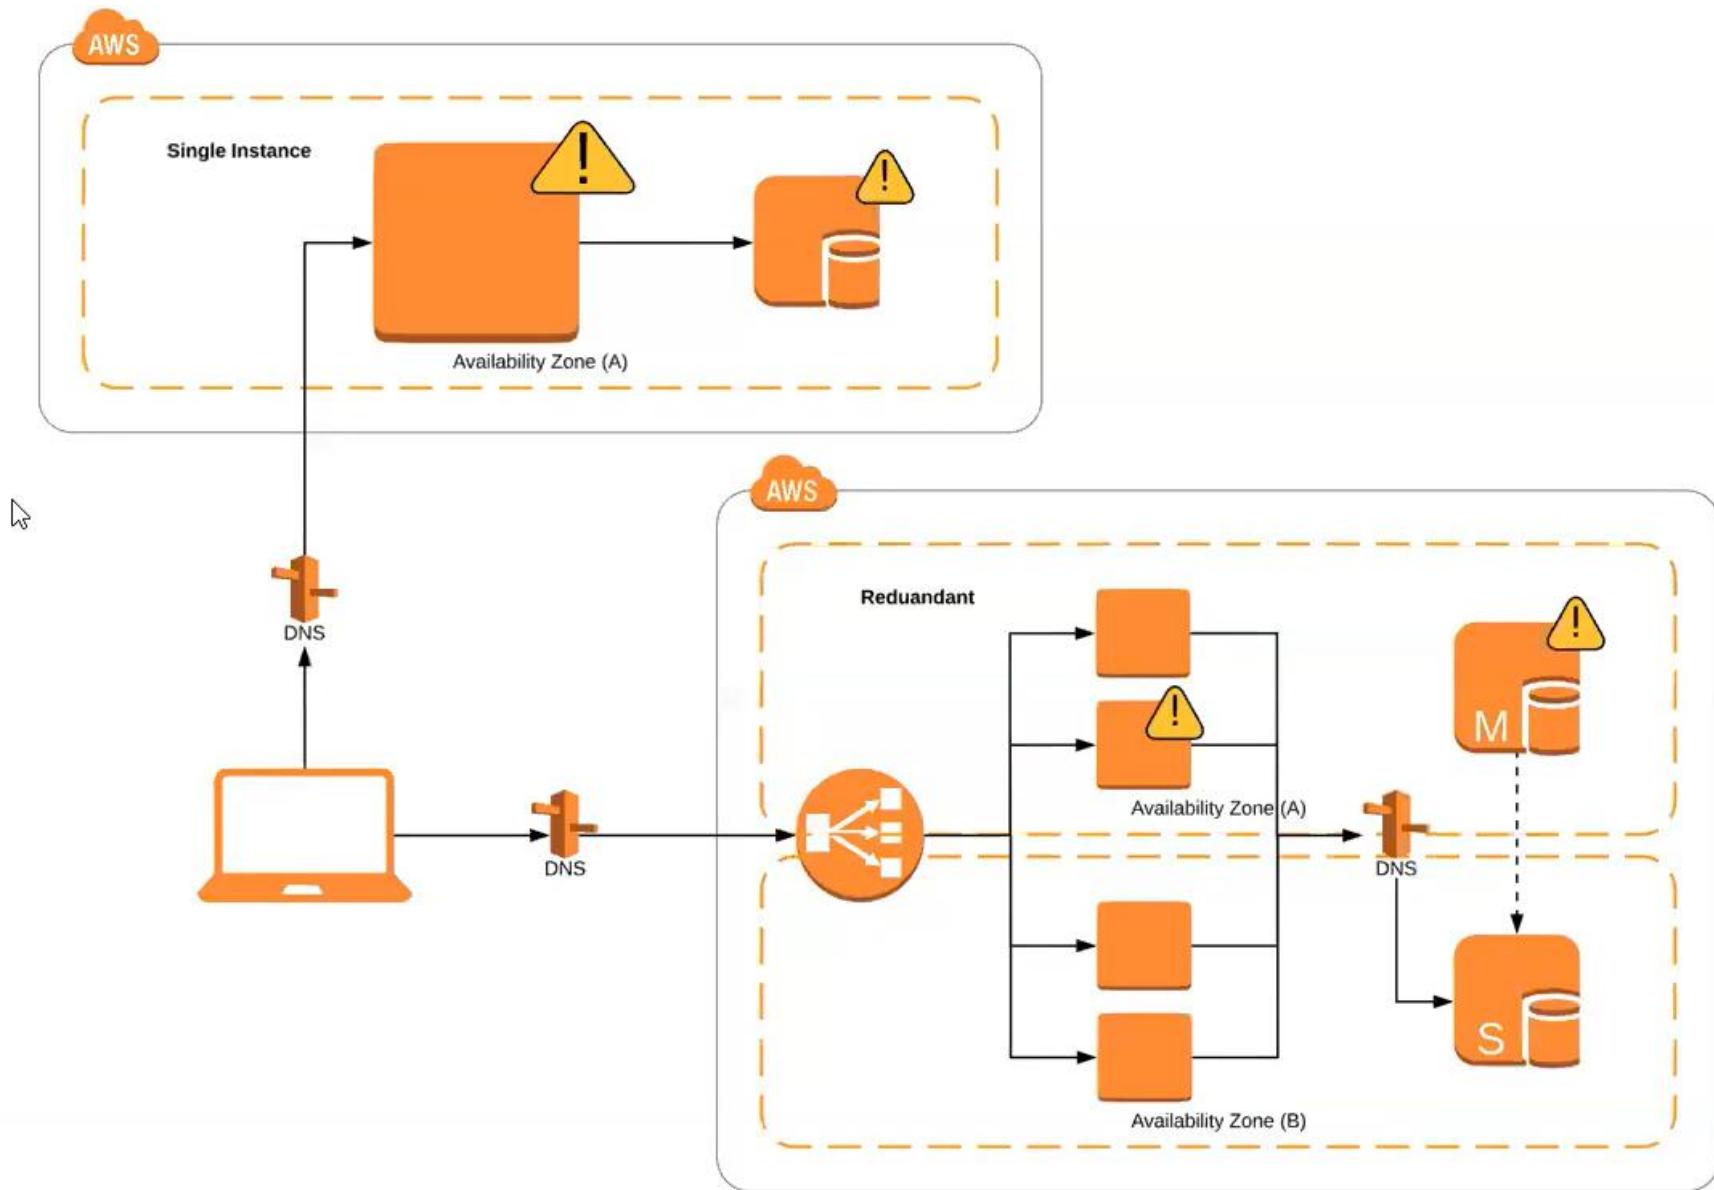
Regions dont communicate with each other by default, if required yes

Network
I NO
default

EC2 instance is specific to Region and AZ.



Ireland Region



AWS Services can be either **Regional or Global**

EC2 is Regional Service

EC2 = Elastic Compute Cloud

Servers = Instances / EC2 instances (VM's)

Load Balancer = Which distribute the traffic to the servers

Elastic Load Balancer (ELB) = Which distribute the traffic to multiple EC2 instances across AZ's

ELB is completely managed by AWS (HA, AS, Scalability, Performance etc)

ELB is a Service from AWS, not Server

You cannot login to ELB, you can access ELB with DNS Name

ELB doesn't have AZ's, it is created at Regional Level

ElasticBeanstalk = Easy and Quick deployment of applications in AWS

In General, PaaS ---> you don't have access and control on the Servers

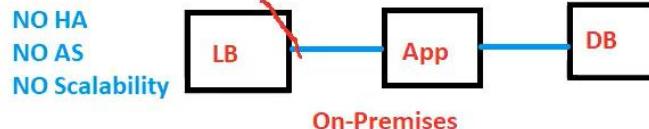
In AWS Beanstalk ---> You have full control on EC2 instances launched by Beanstalk

Beanstalk Handles EC2 instances (OS) behalf of us

The backbone of Beanstalk is EC2

LightSail = If you want to setup and create a virtual lightsail instance which already has everything installed (Wordpress, Gitlab, Nodejs, Joomla, Drupal, Redmine, nginx, Cpanel etc)

It doesn't support Auto-Scaling



EC2 - Launch EC2 instance, configure, deploy, maintain

Beanstalk

Give it some configurations

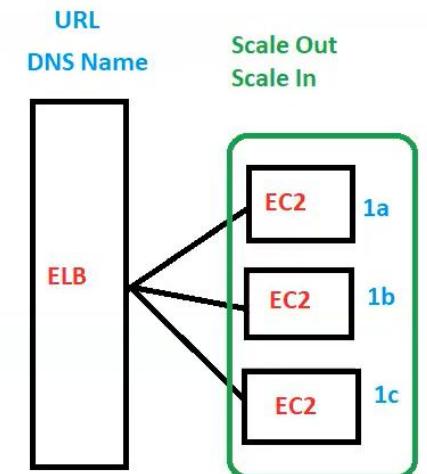
**Upload the application
select the Platform**

URL

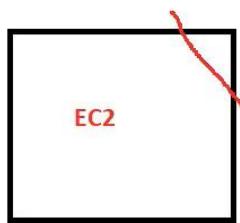
App

EC2

**Tomcat
.NET
Docker
Python
GO
etc**



Lambda is used for automation



Rule

Launch Event
Triggered

Event Bridge

Lambda

Lambda is Serverless

You can run the code without Server

Invoked

Lambda is invoked based on the trigger / Event

Ganesh

Create Functions

Lambda Functions Python, .net, java etc

TerminateFn

Python Code (To terminate the
EC2 instance)

Launch
Stop
Start
Reboot
Terminate ----> Destroy
Kill
Delete

Schedule



9PM -----> STOP

----> STOPFn (Stop all the EC2 instances)



9AM -----> START

----> STARTFn (Start all the EC2 instances)



Stop/Start EC2 instances

40 EC2's

Python Code

In AWS, all services will start with Simple and end with Service

SNS ---> Simple Notification Service

SES ---> Simple Email Service

S3 is Object Based Storage

Laptop

Windows

Folder

Files

jawan.mp4

S3

Bucket

Objects

KEY

S3 = Simple Storage Service

S3 is unlimited Storage

S3 is used to just store the files

S3 can store all FLAT files

With S3, we can upload, download and access files

You cannot execute any files in S3

Is it possible to install OS in S3 ? NO

Is it possible to install DB in S3 ? NO

Is it possible to install, run, execute any files in S3 ? NO

S3 is Serverless

AWS handles HA, Performance, Scalability etc for S3

Bucket is a container of Objects

Object is File

Name of the file / Object is a KEY

S3 is Global

Buckets are Regional

Floppy ---> 2MB

CD's ---> 700 MB

DVD's ---> 4.7 GB

Pen Drives ---> 128 GB

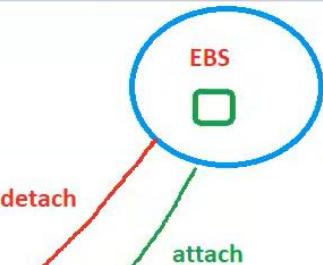
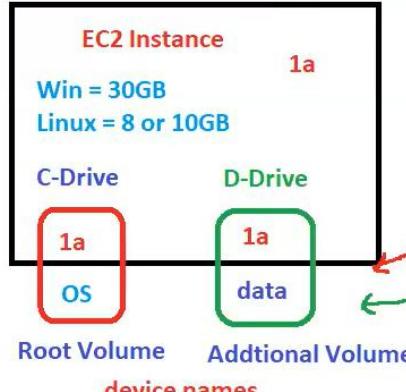
Hard Disks ---> 2 TB

S3 Support Static Website Hosting

Create a bucket, put all your files inside the bucket and enable static website hosting

No need to worry about HA, Performance, Scalability etc because S3 handles it

EBS is Block Based Storage



EBS - Elastic Block Storage

Hard Disk = Volume = EBS Volume

Volumes can be attached and detached

You can attach multiple Volumes to the EC2 instance

EC2 instance has default volume, that volume is called ROOT Volume

The ROOT volume always has OS (Windows, Linux)

EC2 supports only Server Side OS not Client Side OS

If you have OS on the Volume, the Volume is called ROOT Volume

EC2 instance can have only 1 Root Volume

EC2 instance can have multiple additional volumes

Max Size of the EBS Volume is 16TB

Volumes should be pre-provisioned like 50GB, 100GB Max 16TB

You cannot attach a volume to multiple EC2 instances at the same time *

Volume Size can be increased on FLY (NO need to stop the EC2 instance, no downtime)

Volume size cannot be decreased (delete the vol and re-create it based on your size requirement)

Is it possible to detach the ROOT volume while EC2 is running ? NO : Stop the EC2 instance first and then detach the ROOT Volume

Is it possible to detach the Additional volume while EC2 is running ? YES : It is not recommended to detach while running, Stop first

You cannot delete the volume while it is attached, detach first and then delete

EC2 instance has AZ, Volume also have AZ

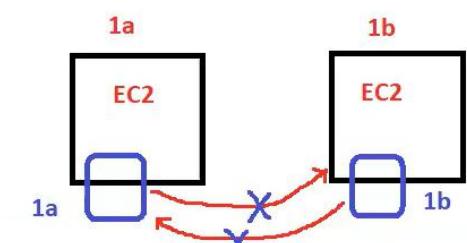
EC2 instance and Volumes should be in the same AZ

We cannot attach 1a volume to 1b EC2 instance (diff AZ)

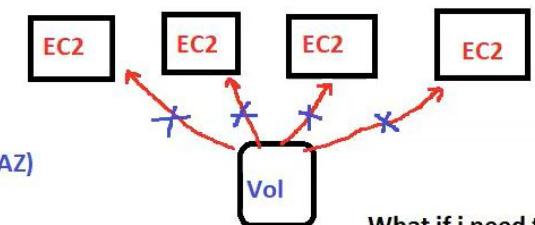
We can attach 1a volume to 1a another EC2 instance (Same AZ)

You cannot share the volumes among EC2 instances

OS
Client Side
win 10, 11
X
Server Side OS
win server 2020,
2021, 2022, 2021
RedHat, CentOS,
SUSE etc

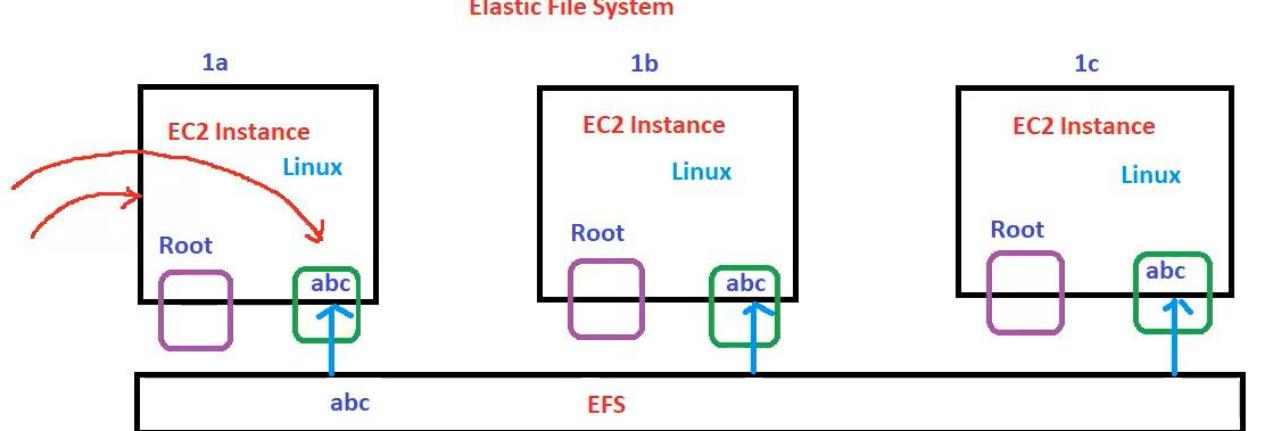
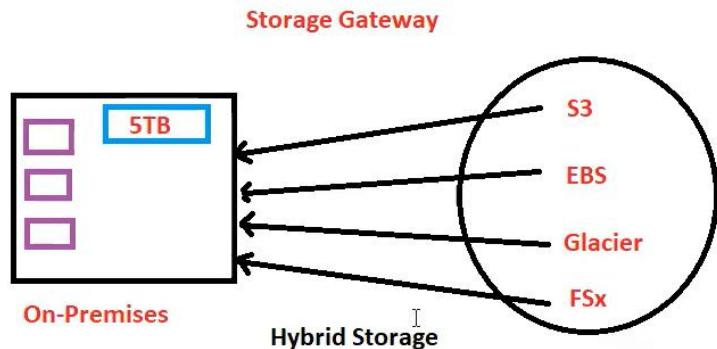


1a vol can be attached to 1b
EC2 instance ? NO (diff AZ)



What if i need to
have shared
storage ?

Is it possible to attach a single Volume to multiple EC2 instances at the same time ? NO



EFS is completely managed by AWS

EFS is only for Linux EC2 instances

FSx is for Windows EC2 instances

EFS works with NFSv4 Protocol

EFS is file based Storage

EFS is unlimited Storage

EFS doesn't require any pre-provisioning (it will automatically increase and decrease based on the data you put in EFS)

EFS can be mounted to multiple EC2 instances at the same time across AZ's

Glacier = Archive Purpose

Not Frequently used data

Cheaper than S3

SNOW Family

SnowCone ----> 8TB

SnowEdge ----> 100TB

SnowMobile ---> PB's

Import/Export Service
Snowball

-----> S3

SnowFamily is used to transfer huge data from on-prem to AWS and vice-versa

SnowFamily is a physical data transfer using devices

On-Premises



Physical Server

DBA

NO HA
NO Elasticity
Scalability issues
Performance issues
DB Crash
Backup issues
Restoration issues
Storage issues
Upgrade issues
Migration issues
replication issues
mirroring issues
loogin issues
security issues
failover issues
data issues netowrking
issues
maitnenace issues
cost issues
patching issues
Headache issues

Database Services

RDS = Relational Database Services

RDS is a Service where you can setup, configure, maintain , secure the RDBMS Databases

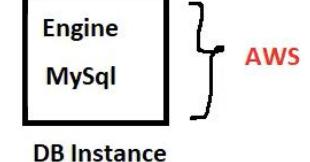
RDS is a Database Service, not database

RDS DB isntance, RDS Support only RDBMS Databases

RDS Supports 6 Engines

MySQL ---> Opensource
Oracle ---> Oracle
MSSQL ---> Microsoft
PostGreSql ---> Opensource
MariaDB ---> Community
Aurora ---> AWS

MOMPMA



Can we do RDP, SSH to
RDS DB Instance ? NO

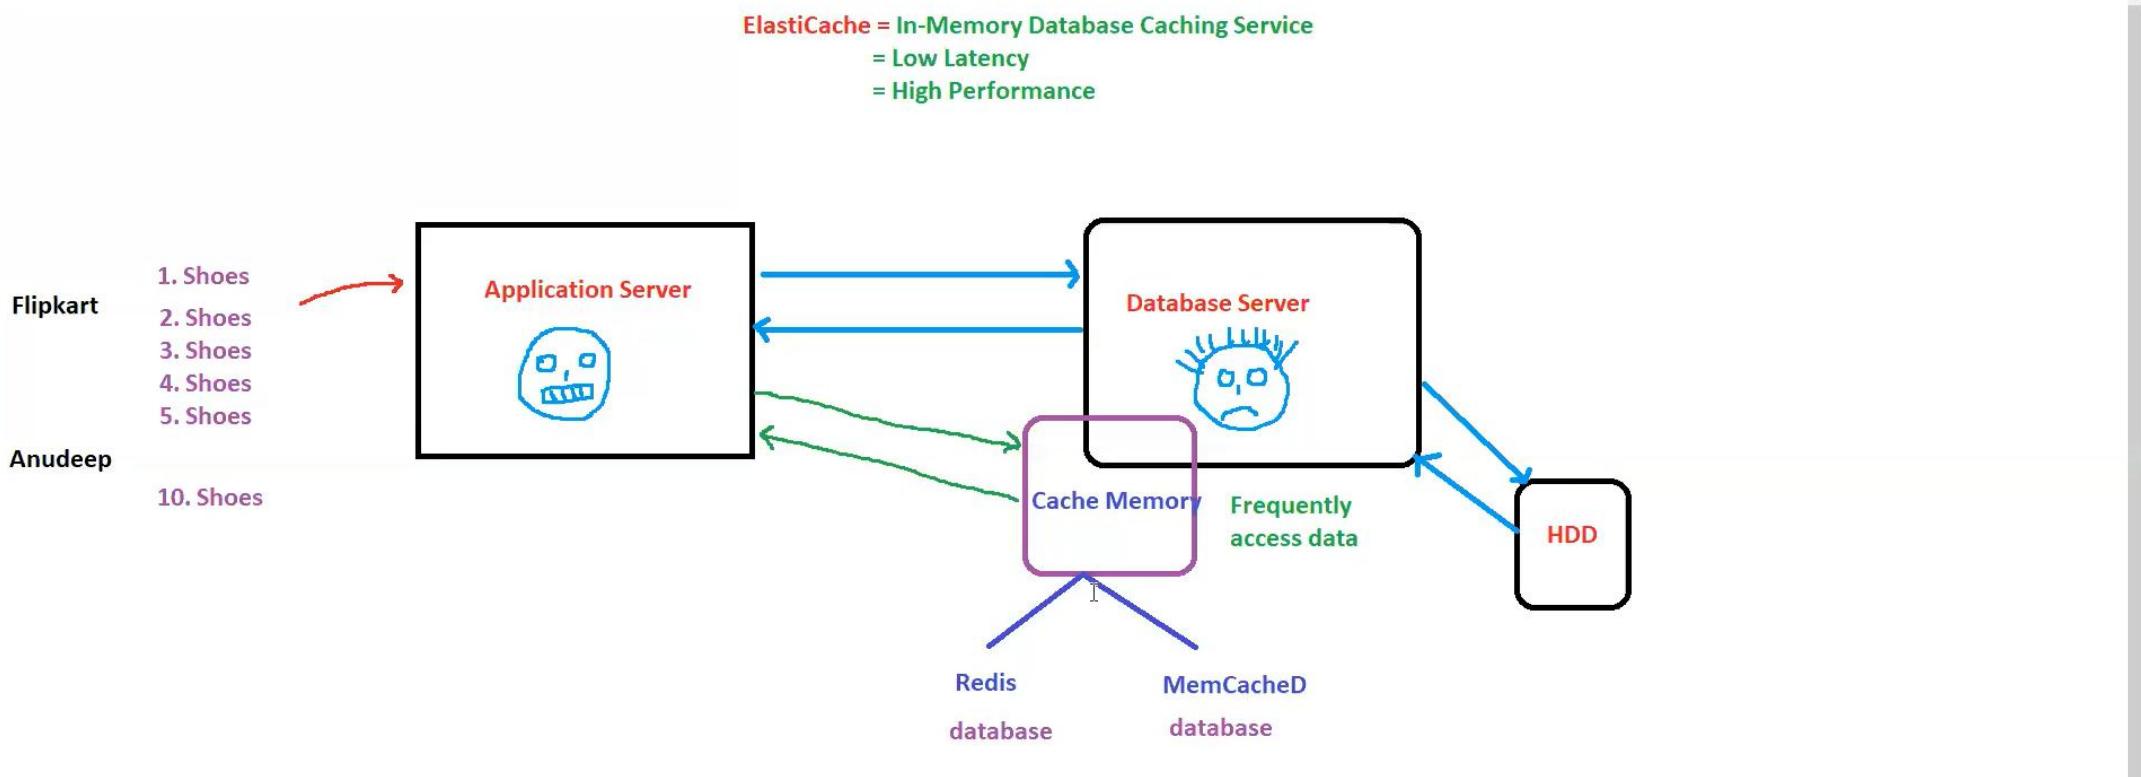
Platform is managed by AWS

DynamoDB ---> NoSql Database Service (Developers) [Non-Relational]

Database = It is used to store the data

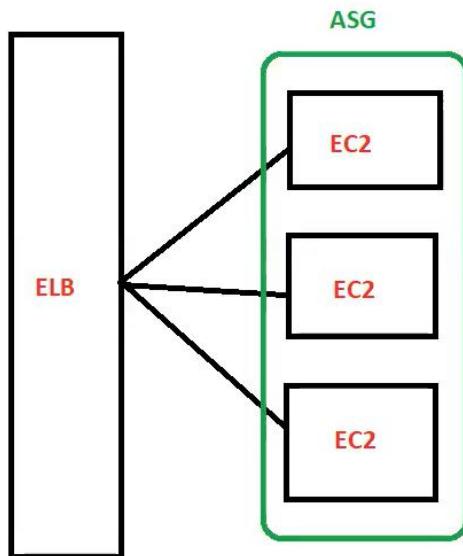
Datawarehouse = It is used to store the HUGE Data

RedShift = Datawarehouse in AWS



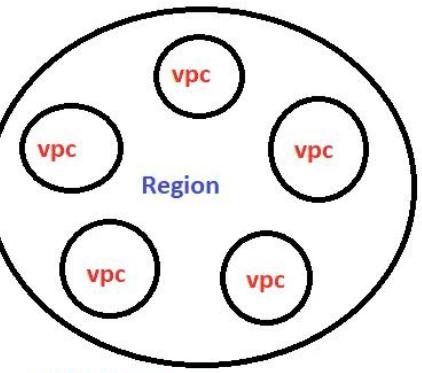
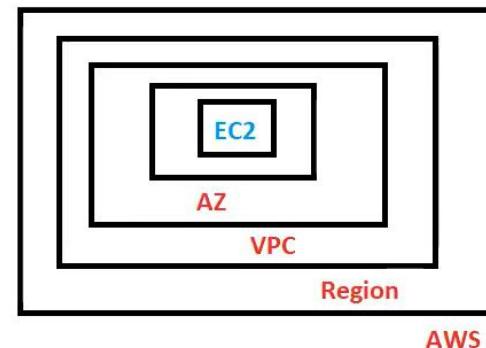
Route53
R53 is a DNS Service from AWS, DNS Port Number is 53

R53 contains Records

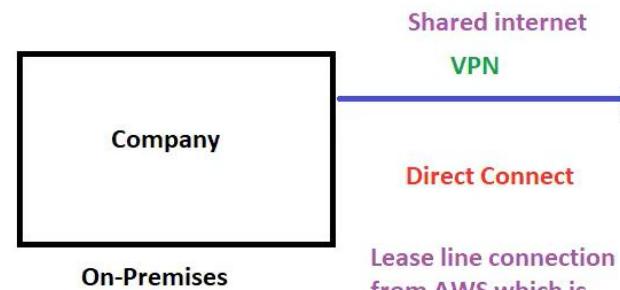


<http://myelb.876786876.elb.ap-south-1.amazonaws.com>
boom.com ----> Nasty URL mapping

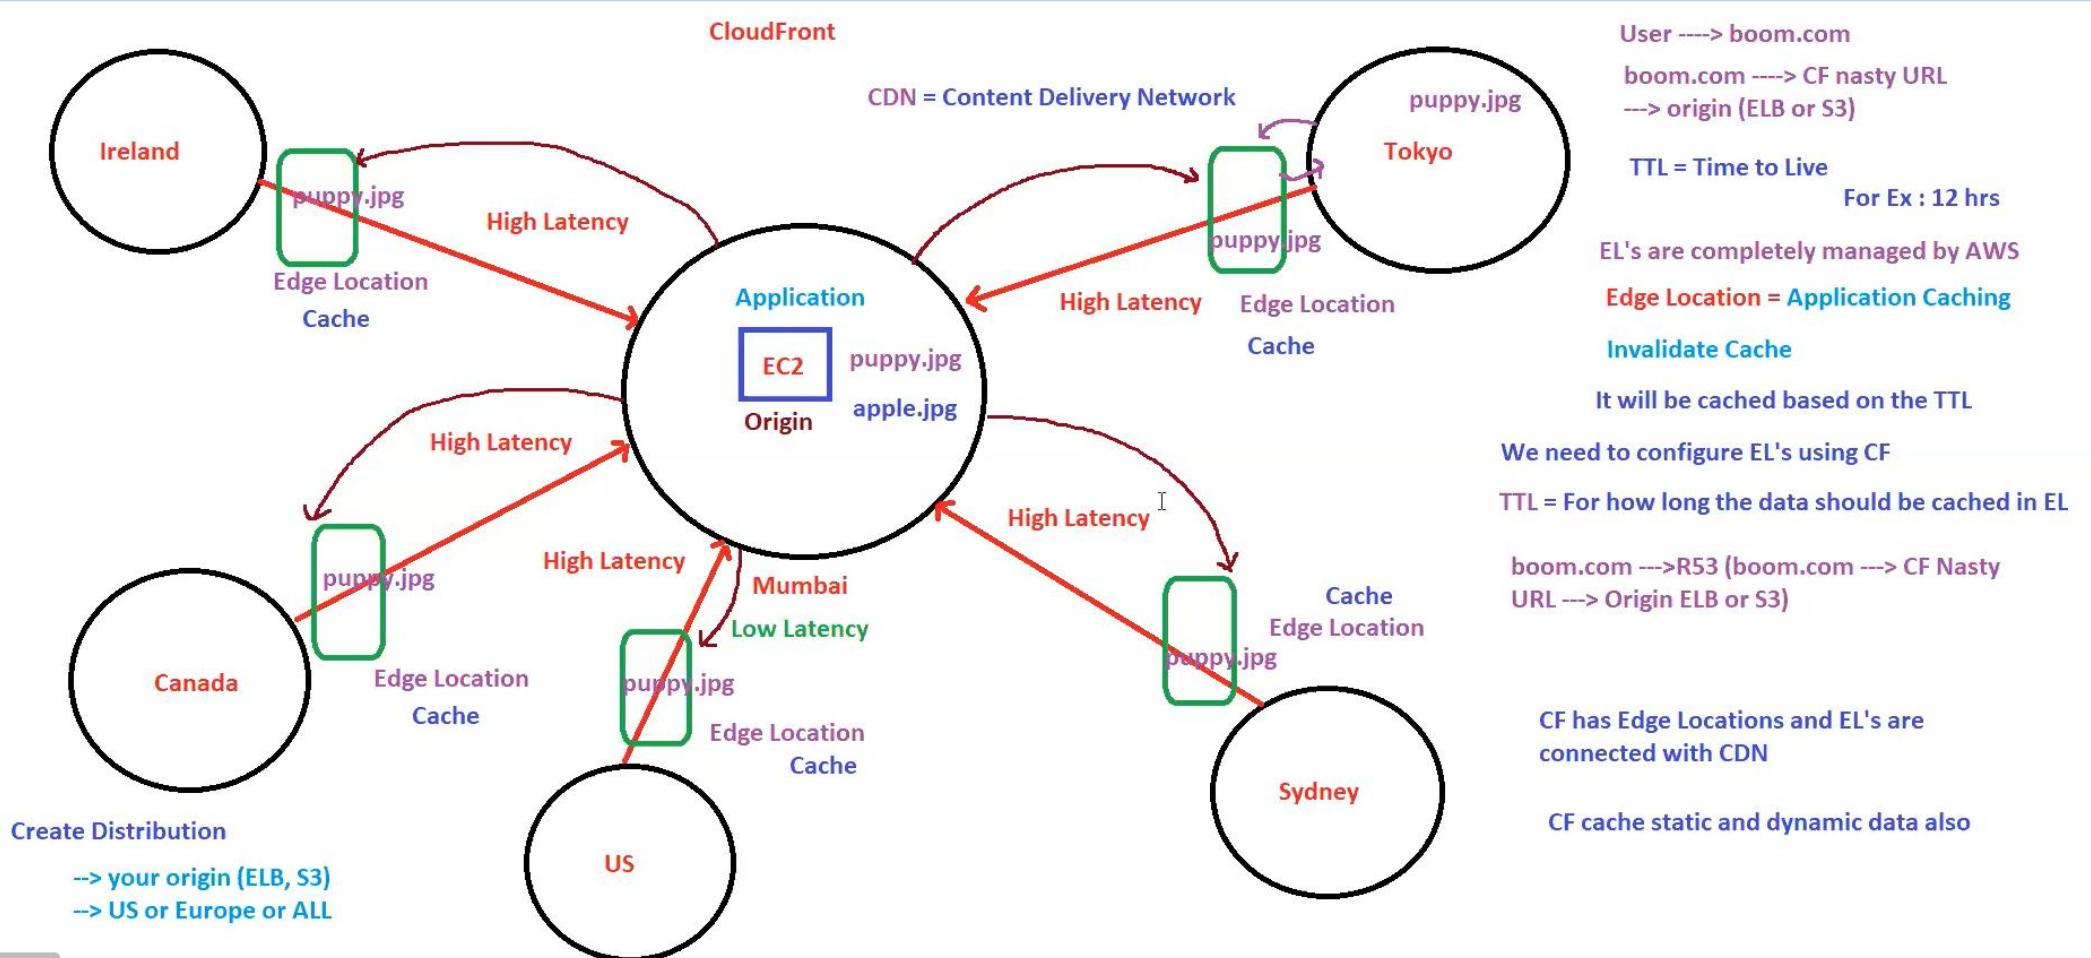
VPC = Virtual Private Cloud
It is like a virtual datacenter on Cloud

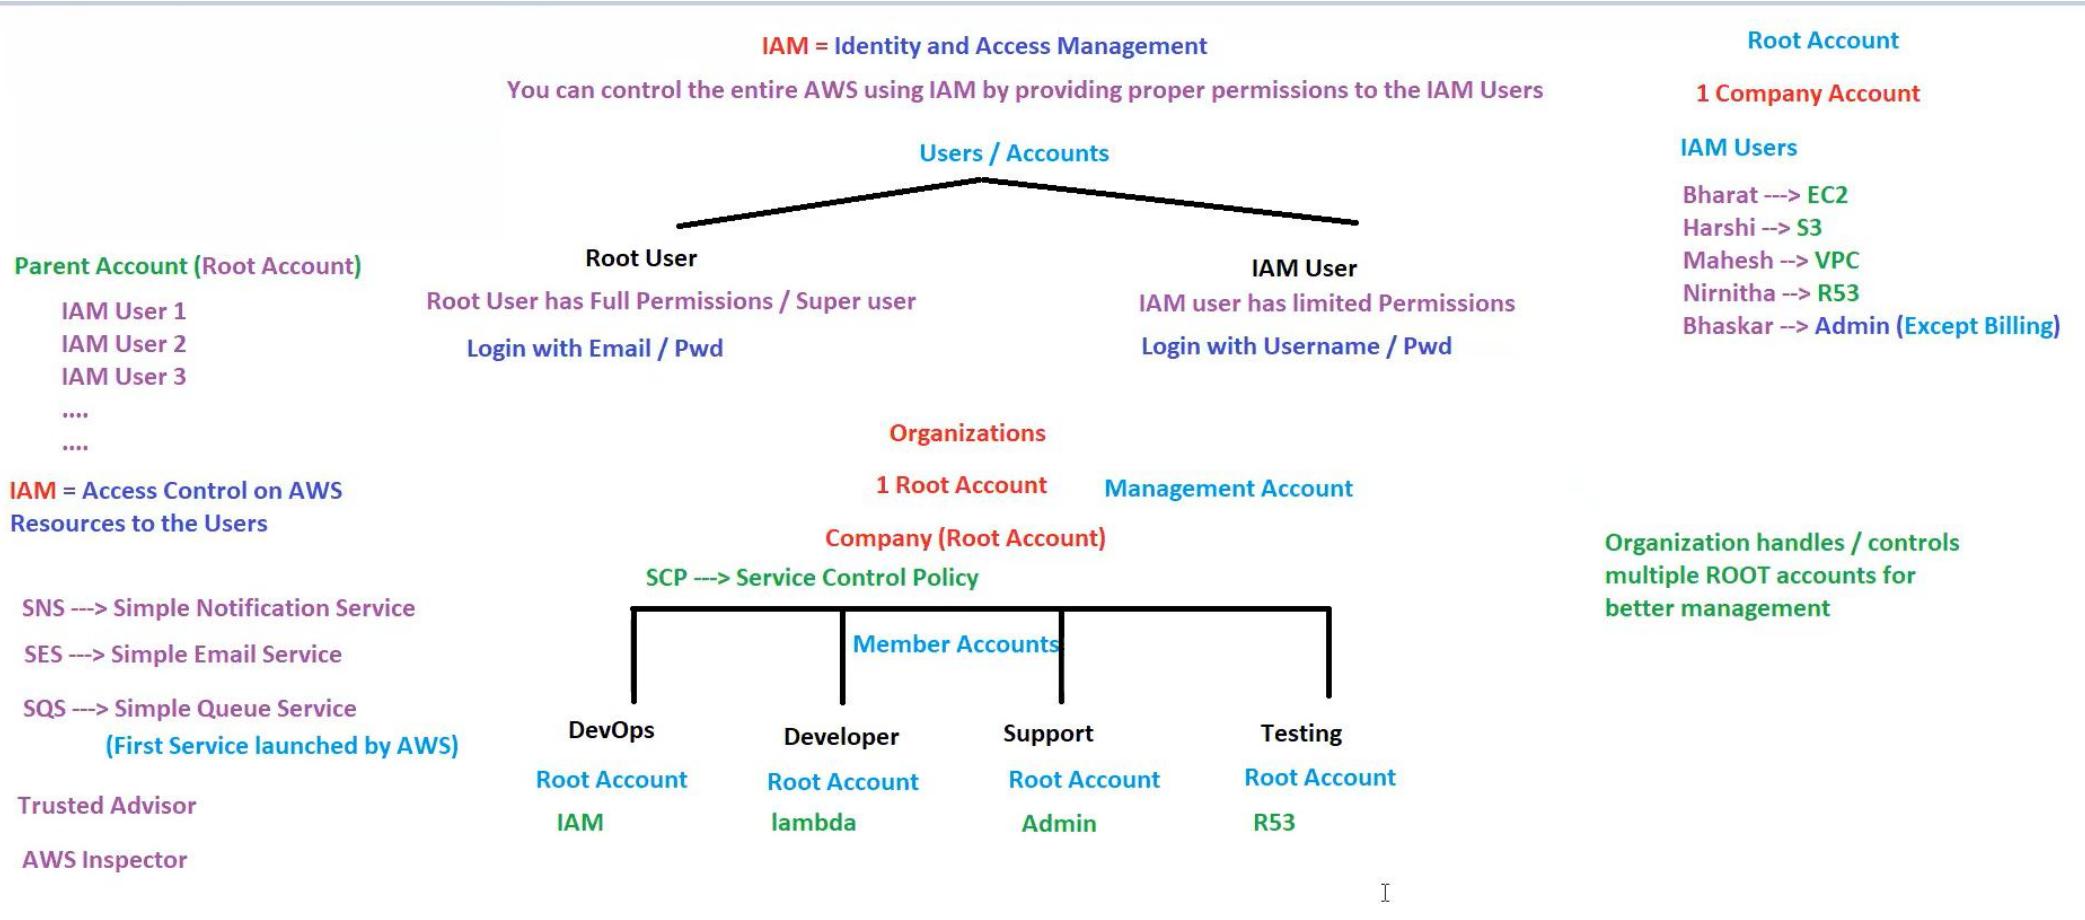


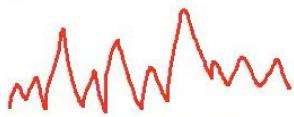
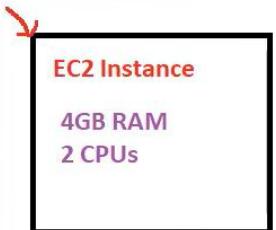
VPC is Regional
AWS provides a default VPC



\$16,000







CloudWatch
CPU > 90%
SNS
---> Email
---> SMS

CloudWatch

CloudWatch is used to monitor all AWS Resources (EC2, ELB , S3, RDS etc)

In CloudWatch we create Alarms to monitor AWS resources

CloudWatch monitors the performance

Basic Monitoring : You will get the data points every 5 mins, FREE [default]

Detailed Monitoring : You will get the data points every 1 min, Billable

CloudTrail : Records all activities happening in AWS account

Event Logger / Investigation Purpose

Config : Monitors the changes in AWS Resources

ALARMS

Metrics (CPU, Network, Disk, Requests etc)

AWS Support

Basic Support ---> FREE

Developer Support

Business Support

EnterPrise Support

\$15,000

IAM is Free

IAM is Global

Using IAM User, we can access the AWS Console

It is not recommended to use the ROOT account for daily activities or work, Instead use IAM User

MFA ---> Multi Factor Authentication (Google Authenticator)

MFA is highly recommended for ROOT and IAM users as well

We need to setup MFA for every individual IAM user

Open AWS Page ---> Login with Email/Pwd ---> MFA Code ---> Login to AWS Console

IAM User
Ganesh
EC2

MFA
(IAM User/Pwd)
Console
Programmatical
KEYS NO MFA

Once Keys are lost , it is lost , you cannot get the same keys back

But you can re-generate N Number of times

If you re-generate, you will get the new KEYS, you cannot get the old keys back

DONT Create KEYS and use the KEYS of ROOT Account

IAM - Deep Dive

IAM is used for Security Purpose

With IAM, you can control entire AWS resources centrally

Dont share your email/pwd to others

You can share the ROOT account by creating IAM user

Root User ---> Full Permissions

IAM User ---> Limited Permissions

Permissions ---> Policies

IAM Users

Sowmya ---> EC2

Venky ---> S3 , IAM

Surya ---> VPC

Rizwan ---> RDS

Mahesh ---> Admin (Except Billing)

Access Key : AKJDFHAKJDHFJKDHF

Secret Key : S^D&*SFDA%\$*D%FD

} 1 Set

2 Ways to Access AWS

Console Access
AWS Console Website (GUI)
(Email/Pwd , Username/Pwd)
Root IAM

Console access can also be Disabed and Enabled

Programmatical Access
(CLI, SDK's, Developer Tools)

Login with KEYS (Access Key and Secret Key)

KEYS are user specific, individual IAM User have their own KEYS

It is not recommended to share the KEYS to anyone

Create the KEYS based on the requirement, dont create it unnecessarily

KEYS also have same permissions like Console

Every IAM User can have max 2 Set of KEYS

We create the KEYS and AWS provides the KEYS

We can make the KEYS inactive based on the requirement

IAM Groups

IAM Group = Collection of IAM Users

Group under Group is not possible / Nested Groups are not possible

It is possible to attach multiple policies to the IAM user and IAM Groups also, Max 10

You can add/attach and remove/detach policies to the IAM users and Groups any time

If you attach any IAM user to the IAM Group, IAM user individual policies remains same, and the new policies will be inherited to the IAM User.

You cannot assign / Create KEYS to the Groups

KEYS are only for IAM users and not for IAM Groups

An IAM user can be attached to multiple IAM Groups at the same time

IAM groups are used to assign policies to the bunch of IAM user at the same time

A Brand new IAM user will not have any Permissions / Policies attached by default

User Based Permissions

EC2fullaccess

S3readaccess

Managed

Resource Level Permissions

Granular Level / Deeper Level

Customized Permissions

Inline

R53

IAM Group

DevOps

Admin

Mounish

Azar

IAM Users

Anudeep --> EC2 + R53

Sowmya --> S3 + R53

Krishna --> IAM + R53

Kiran --> RDS + R53

Siva --> Admin + R53

Policy ---> Policy contains Permissions

Permissions / Policies are written in JSON Format

Managed Policy : Created and managed by AWS (Predefined Policies)

EC2Fullaccess

EC2Readaccess

Inline Policy : Created and managed by Customer
(Customized Policies)

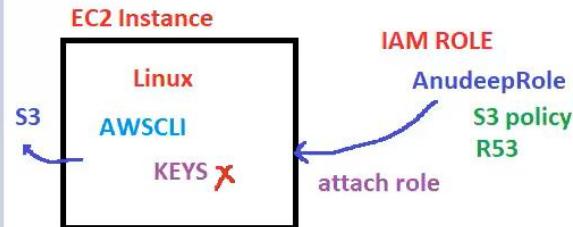
Customer Managed Policy

Visual Editor / Policy Generator

--> Json code generated automatically

ARN = Amazon Resource Name

To identify a AWS Resource to put in IAM Policy



Instance IP
Instance Username
Instance Password

IAM user is used to access AWS
Console Not Servers / EC2 instance

If you don't configure the KEYS on Linux EC2 instance , you cannot access AWS Services

If you configure the KEYS on EC2 instance,
KEYS are stored locally on the EC2 instance
which is not SAFE (Because might be hacked)

IAM ROLES

ROLE = Temporary access without Credentials

If you use the ROLES, we no need to configure the KEYS On the EC2 instance

Based on permissions, that you have attached to the ROLE, those permissions are available from the EC2 instance

1 EC2 instance can have only 1 ROLE attached at the same time

1 ROLE can be attached to multiple EC2 instances at the same time

Identity Provider / Federation / Identity Center

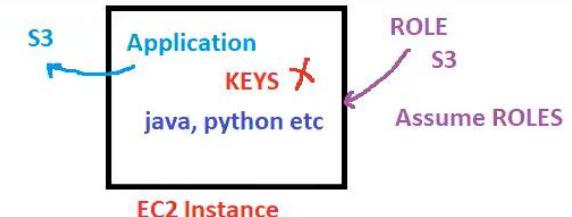
SSO - Single Sign On no need to create IAM user

Company Email / Pwd ---> AWS Credentials ---> Login to AWS

LDAP

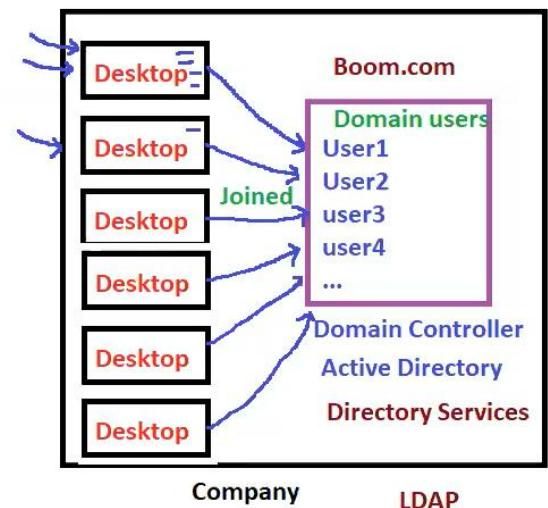
SSO

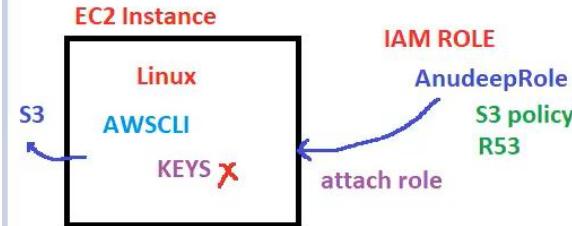
LDAP
OAuth
Okta



Some AWS Services will need to perform actions on your behalf

Boom





Instance IP
Instance Username
Instance Password

IAM user is used to access AWS
Console Not Servers / EC2 instance

If you don't configure the KEYS on Linux EC2 instance , you cannot access AWS Services

If you configure the KEYS on EC2 instance,
KEYS are stored locally on the EC2 instance
which is not SAFE (Because might be hacked)

IAM ROLES

ROLE = Temporary access without Credentials

If you use the ROLES, we no need to configure the KEYS On the EC2 instance

Based on permissions, that you have attached to the ROLE, those permissions are available from the EC2 instance

1 EC2 instance can have only 1 ROLE attached at the same time

1 ROLE can be attached to multiple EC2 instances at the same time

Identity Provider / Federation / Identity Center

SSO - Single Sign On no need to create IAM user

Company Email / Pwd ---> AWS Credentials ---> Login to AWS

LDAP

SSO

Federation to AWS (Identity Provider)
(Connection From LDAP to AWS)

OAuth

Okta

LDAP

Username/Pwd

IAM ROLES

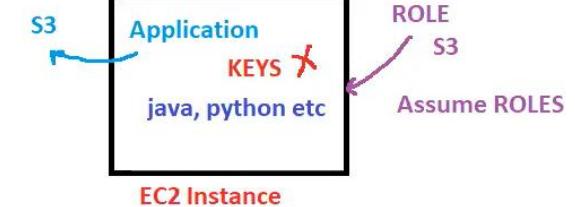
Developer (Lambda)
Admin
DevOps (EC2 etc)

Identity Center

Rahul

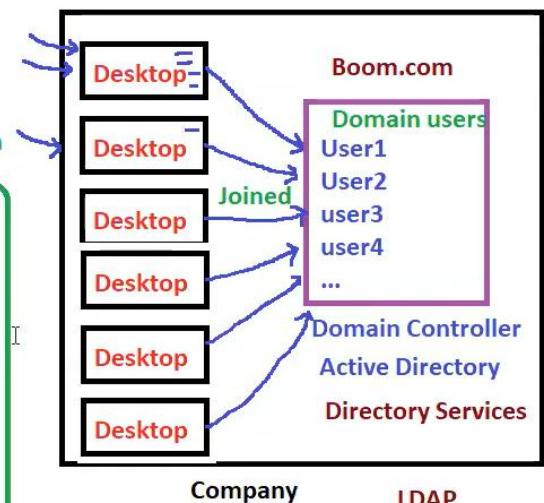
User

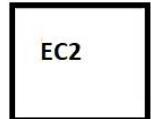
Has a URL



Some AWS Services will need to perform actions on your behalf

Boom





Name = WebServer

Key = value

500 IAM Users

Hardworkers

DevOps

Manually
500 users
delete one by one

Admins

Security

DBA X

100 EC2 instances

50 DevOps
Name = DevOps

TAGS are helpful for doing Cost Optimization also

Max Tags per resource is 50

IAM Credentials Report

A report that list all your accounts users and the status
of their various credentials

IAM - TAGS

TAGS are Key-Value Pair

TAGS are used for identification purpose

TAGS are not IAM Specific, It is through out AWS

TAGS are very important, but it is Optional

TAGS are helpful for doing automations in AWS

Jean

TAG

Price = 10k

Brand = Killer

Size =

Key = Value

It is a best practise to change/rotate the Passwords and
Keys for the IAM users for security purpose

Smart Workers

delete all IAM users where
Key = Name, Value = DBA

Python

Access Analyzer

Access analyzer analyze the access for all
IAM users and gives actions to act on it

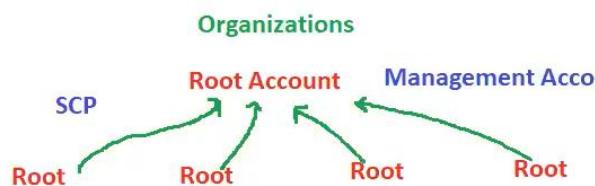
Access Advisor

- . It shows the service permissions granted to the users and when those were last used
- . It is used to revise your policies to the Users

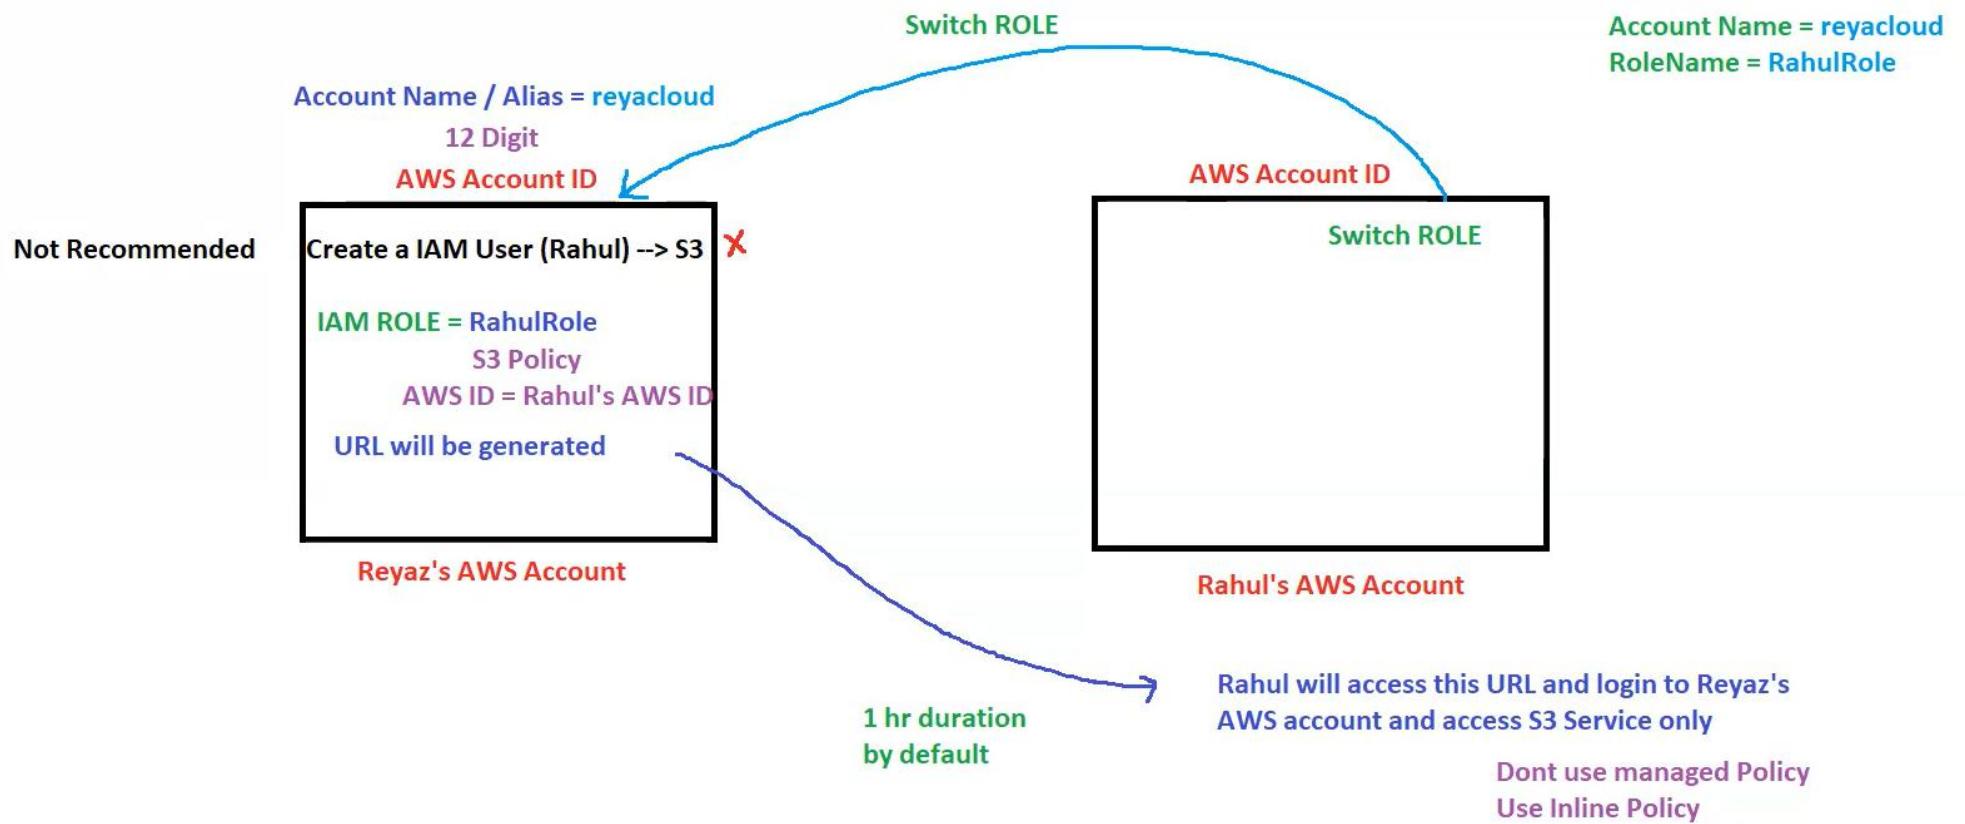
IAM Service

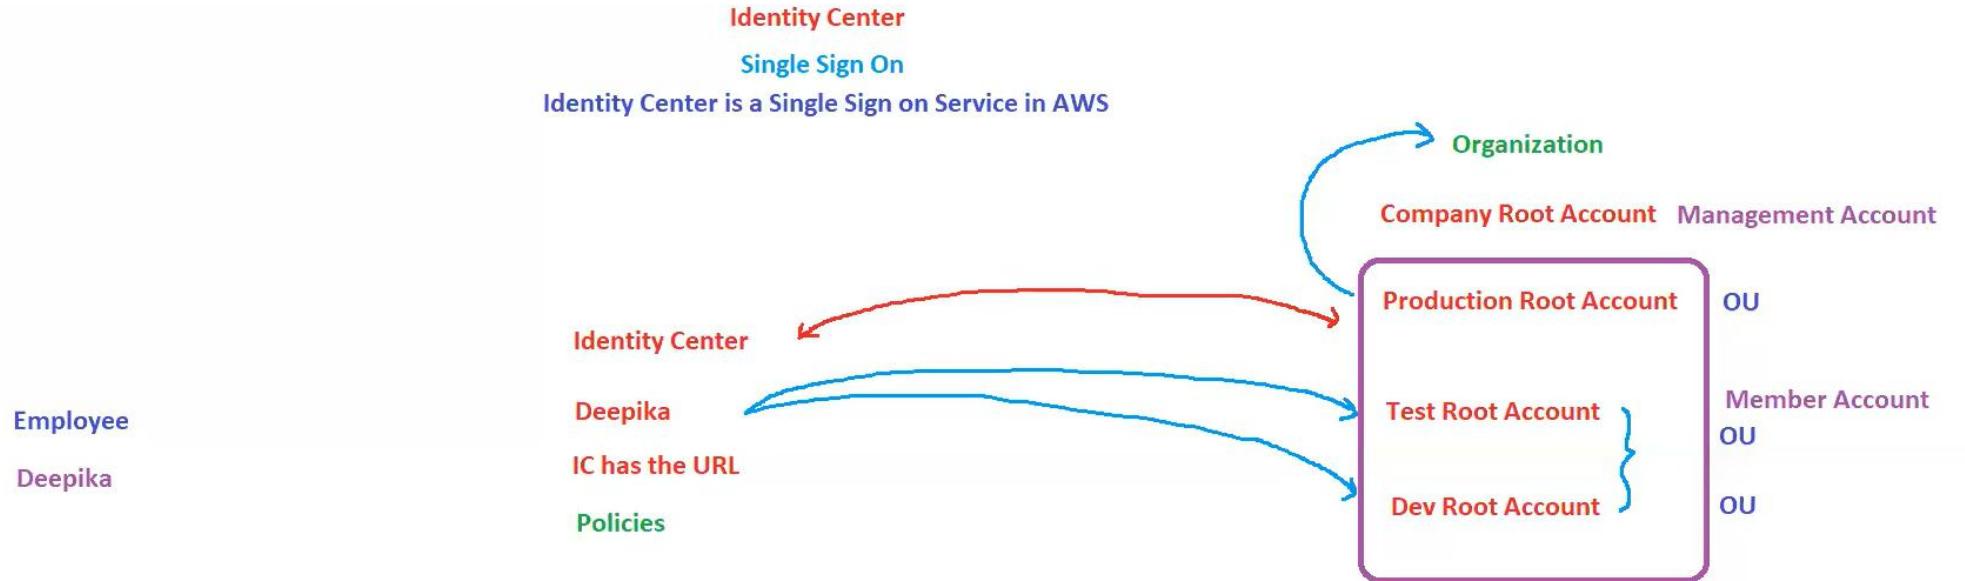
IAM Users
IAM Groups
IAM Policies
IAM ROLES
Federation/ IC

For managing multiple
Root accounts



I





By default Organizations and Identity Center are not enabled [→](#)

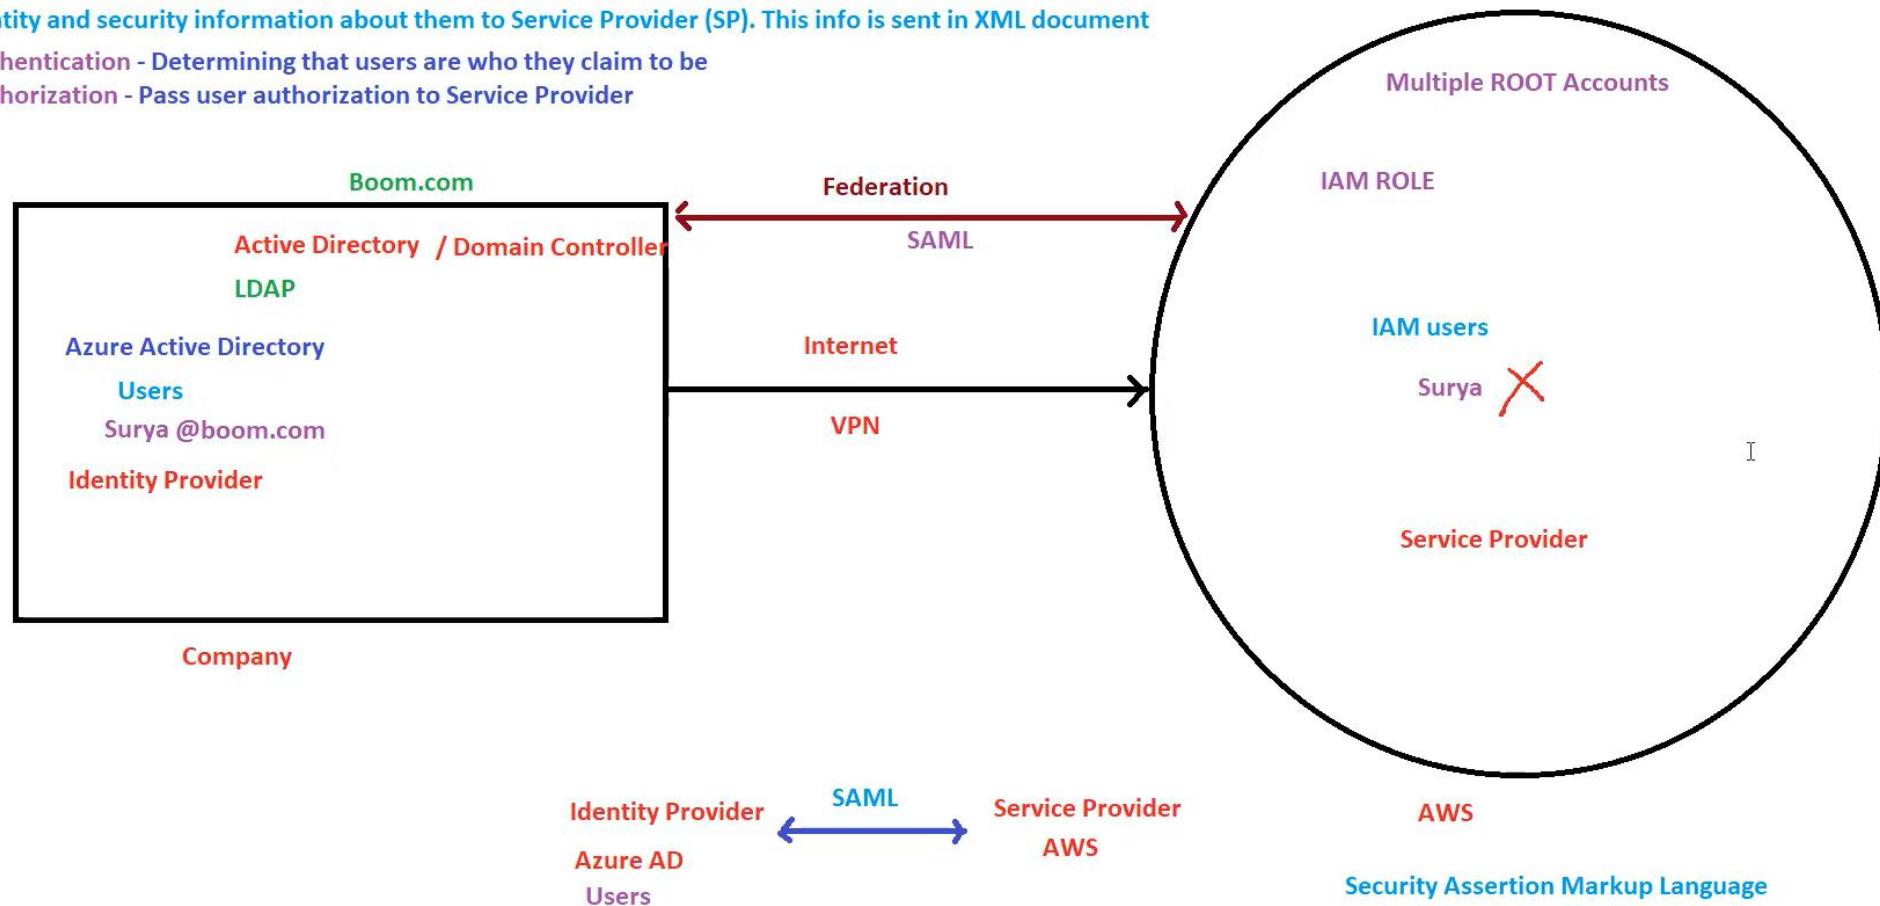
Organization Service is Global

1. You should have min 2 Root accounts
 1. management account 2. Member account
2. Enable Organizations and Identity Center
3. Go to Organization and Add AWS account (invite)
4. Go to IC, Create a Permission Set
5. Go to IC, Create User and assign Permission Set
6. Assign User to the Organization Units
7. Access AWS Access Portal URL and login with the user that you have created
Access at the AWS account

SAML is an Open Federation standard that allows an identity Provider (IdP) to authenticate users and pass identity and security information about them to Service Provider (SP). This info is sent in XML document

Authentication - Determining that users are who they claim to be

Authorization - Pass user authorization to Service Provider



EC2 is Regional

Servers = EC2 Instances

750 hours free per month

Price = Pay as you Go

Pay every hour

On-Demand Instances

Fixed Price (Hourly)

Pay for what you have used

Pay per Hour

No Commitment

No Upfront Payment

No Predictable Usage

Savings Plans

Same as RI but it has diff strategy

1 or 3 year , commitment to the amount of usage, Long Workloads

Get discount based on long term usage

Locked to specific instance family & Region

Flexible across instance Size, OS, Tenancy

EC2 - Elastic Compute Cloud
EC2 is a webservice from AWS that provides **resizable** compute services in the Cloud

Resizable = Scale Up / Scale Down, Scale Out / Scale In

Scalability **Elasticity**

Pricing Models

5 mins ---> Terminated
5 mins ---> Terminated ---> 2 hours

Reserved Instances

Long Term commitment

1 or 3 years

Upfront Payment (Full, Partial)

72% discount on hourly price

Predictable Usage

Standard RI : 72% discount, Longer Workloads

Convertible RI : Long Workloads with flexible instances , 66% discount

(Change instance type, OS, Tenancy etc)

SPOT Instances

Bidding

Auctioning

Huge Capacity for cheaper price

90% Discount

Can loose EC2 instances

Most cost effective instances in AWS

Dedicated HOSTS

Book entire Physical Machine

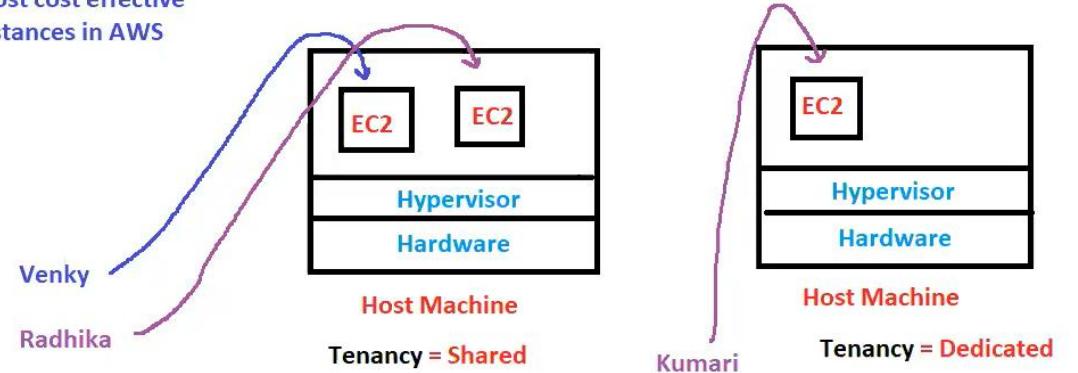
Control instance placements

Dedicated Instances

No other customer will share your hardware.
No Control on Instance Placements

Capacity Reservation

Reserve capacity in a specific AZ for any duration .
No Time Commitment



EC2 Instance Families

General Instances

Memory Instances

CPU Instances

Storage Instances

GPU Instances

Instance Types

Instance Type = CPU + Memory

t2.nano = 0.5GB RAM , 1 vCPU
t2.micro = 1GB RAM, 1 vCPU
t2.small = 2GB RAM , 2 vCPU
t2.medium = 4GB RAM
t2.large = 8GB RAM
t2.xlarge = 16GB RAM

Scalability = Scale Up and Scale Down

Scalability can be achieved by changing the instance type

Anytime Scale Up ---> No Data Lost

Anytime Scale Down ---> No Data Lost

You need to STOP the EC2 instance to change the instance type
(Downtime Required)

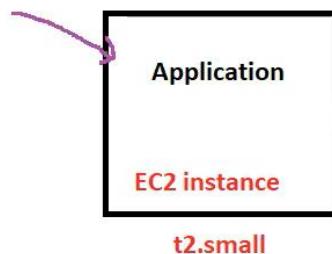
If you want to change the instance type, data in the instance will NOT BE LOST



Burstable Performance Instances

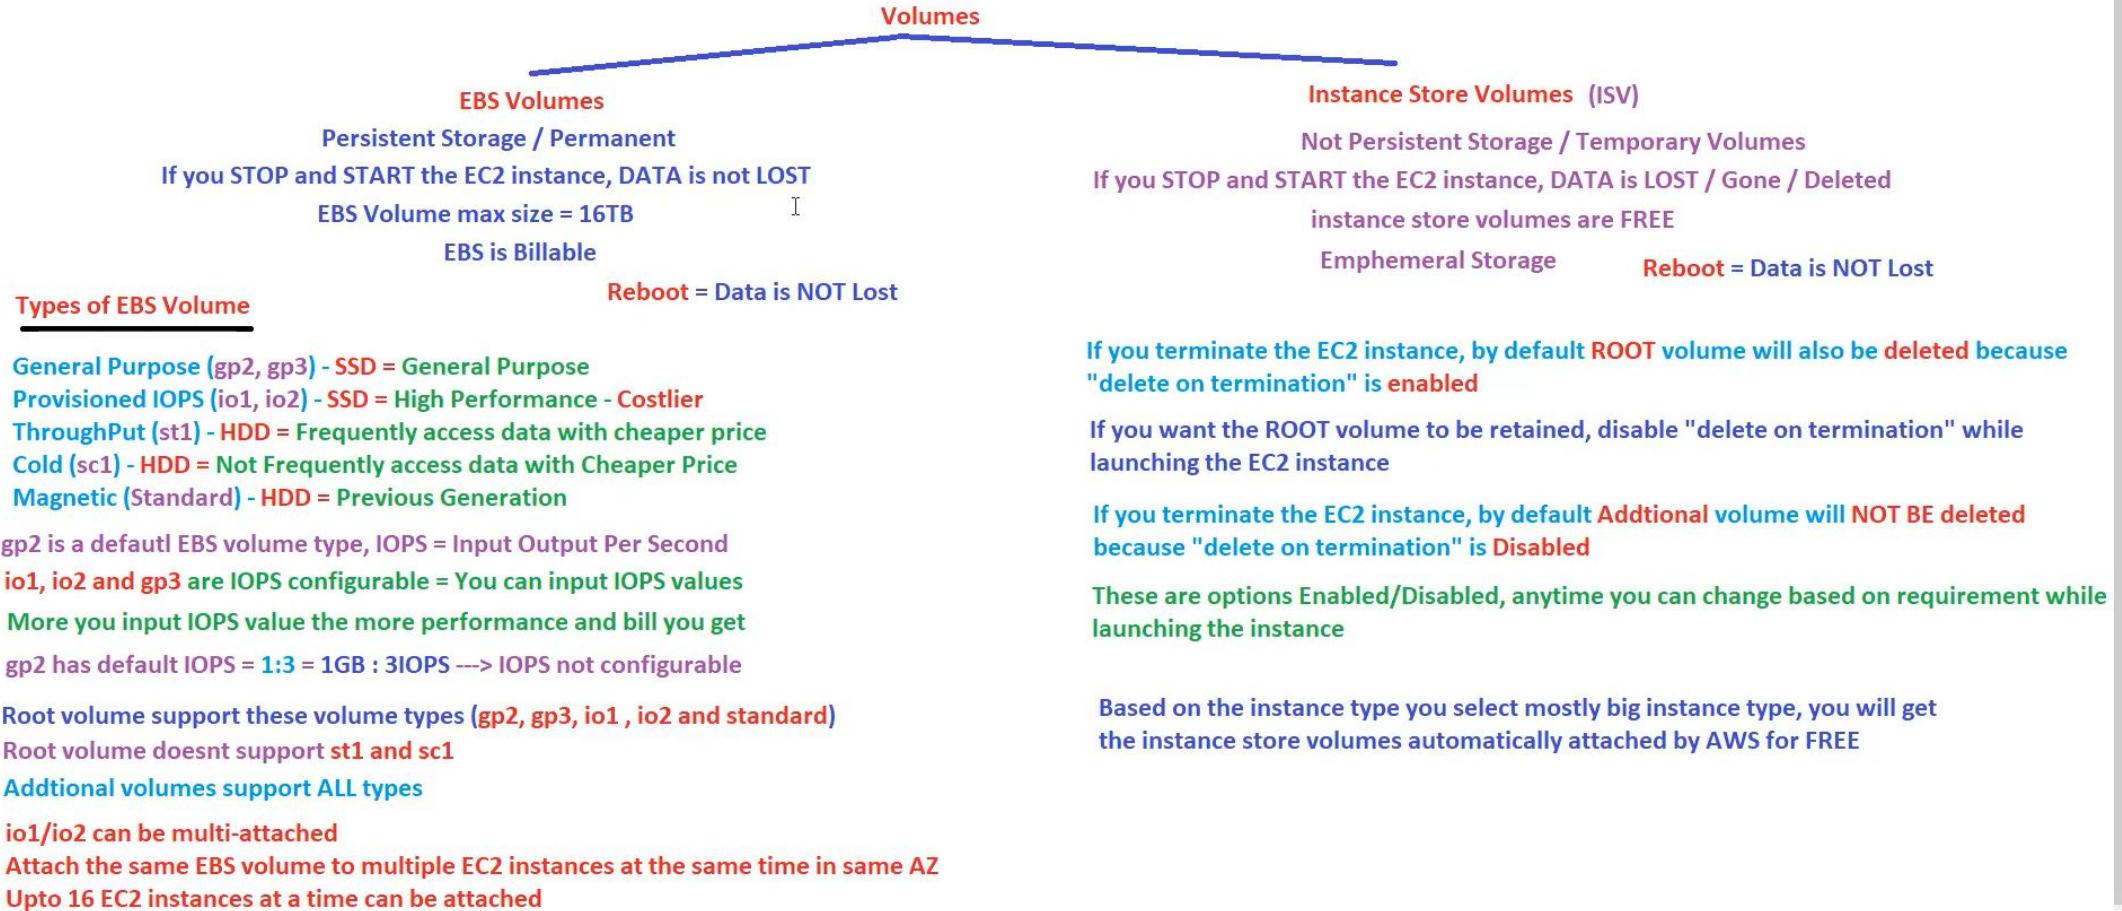
CPU Credits

EC2 instance will enter into the burstable mode and give high performance for limited period of time only



2 vCPU + 6 vCPU

CPU Credits depend upon the instance type

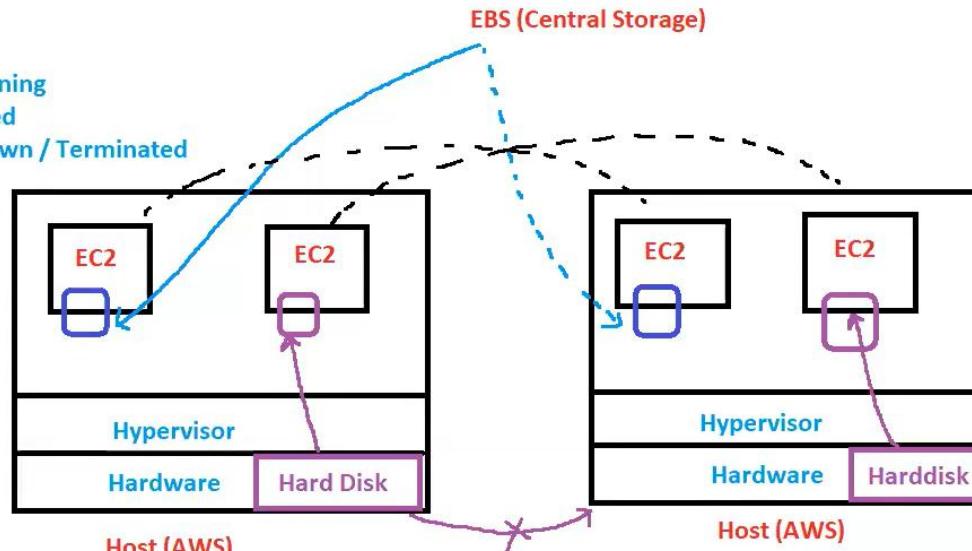


Instance 6 States

Launch --> Pending / Running

Stop --> Stopping / Stopped

Terminate --> Shuttingdown / Terminated



EBS

Stop and Start = JUMP = Data is not Lost

Reboot = No Jump = Data is NOT LOST

Terminate = Data is lost

Instance Store Volume

Stop and Start = Jump = Data is Lost

Reboot = No Jump = Data is NOT LOST

2 Types of Status Checks

Instance Status Check

System Status Check

2 Checks

1st hard check

2nd software check

2/2 checks are passed --> can login to EC2

1/2 checks are passed --> you cannot login to EC2

0/2 checks are passed --> you cannot login to EC2

AWS will do the Status Checks for Us

Solution : Stop and Start (jump) to new host machine and your EC2 will start working with 2/2 checks are passed

I

Solution 2 : Terminate the EC2 if you dont have any data, not recommended and launch new EC2 instance

- . EBS Volumes are network drives with good but limited performance
- . If you need a high performance hard disk, use instance store volumes
- . But Risk of data loss if hardware fails
- . Backups is our responsibility not AWS

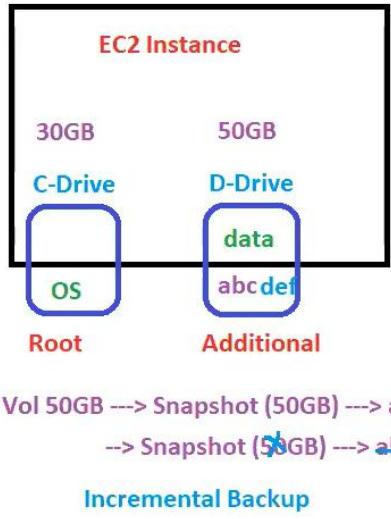
EC2 Hibernate

- . In Memory (RAM) is Preserved
- . The instance boot is much faster (the OS is not Stopped / Restarted)
- . The RAM state is written to a file in the ROOT EBS Volume
- . The Root EBS Volume must be Encrypted

Use Cases

- . For Long Running Process
- . Save RAM State
- . Services that take time to initialize

- . Instance RAM Size = Must be less than 150GB
- . Root volume must be EBS and Encrypted , NOT ISV
- . An instance cannot be hibernated for more than 60 days



To create a snapshot we no need to STOP the EC2 instance

Data Life Cycle Manager: It is used to take the snapshots automatically / Schedule volumes that should get snapshots will be identified by TAGS

Life Cycle Policy/Rule should be created

Retention Period = 7 days

Example

SNAPSHOTS

1. Snapshot is a point in time copy of the Volume
2. Backup of the Volume is called Snapshot
3. EBS Snapshots are Created from EBS Volumes
4. You can create Snapshots from volumes
5. EBS Volumes -----> EBS Snapshots -----> EBS Volumes
6. **You cannot attach a snapshot directly to the EC2 instance, you have to create a volume out of it first and then attach it to the EC2 instance**
7. It is not possible to login or use the snapshot directly
8. Snapshots are Stored in S3 (Provider's S3)
9. Snapshots are visible from the EC2 Console
10. Snapshots doesn't have any AZ's
11. Snapshots are Regional
12. By default Snapshots are Private, if required we can make it public
13. You can copy the snapshot from one region to another region in the same account
14. Snapshots can be shared from one AWS account to another AWS account (AWS ID, Private)
15. EBS Volumes cannot be moved directly to any other regions, accounts --> use Snapshots

EBS Volumes are created from EBS Snapshots

Instance Store Volumes are created from a template stored in S3

EBS Snapshot Tiers - 2 Types

Standard Tier, Archive Tier

Default is Standard Tier

- . Move the snapshots to an Archive Tier that is 75% Cheaper
- . Takes 24 to 72 hours for restoring the archive --> Billable

RecycleBin

- . Setup rules to retain deleted snapshots so you can recover them after accidental deletions (Retention period 1 day to 1 yr)
- . Selected Snapshots can also be sent to recyclebin using TAGS

Fast Snapshot Restore (FSR)

- . Forcefull initialization of snapshot to have no latency for first use

By default, Volumes are Not Encrypted

Encryption/Decryption is managed by AWS

Not Encrypted --> Not Encrypted

Encrypted --> Encrypted

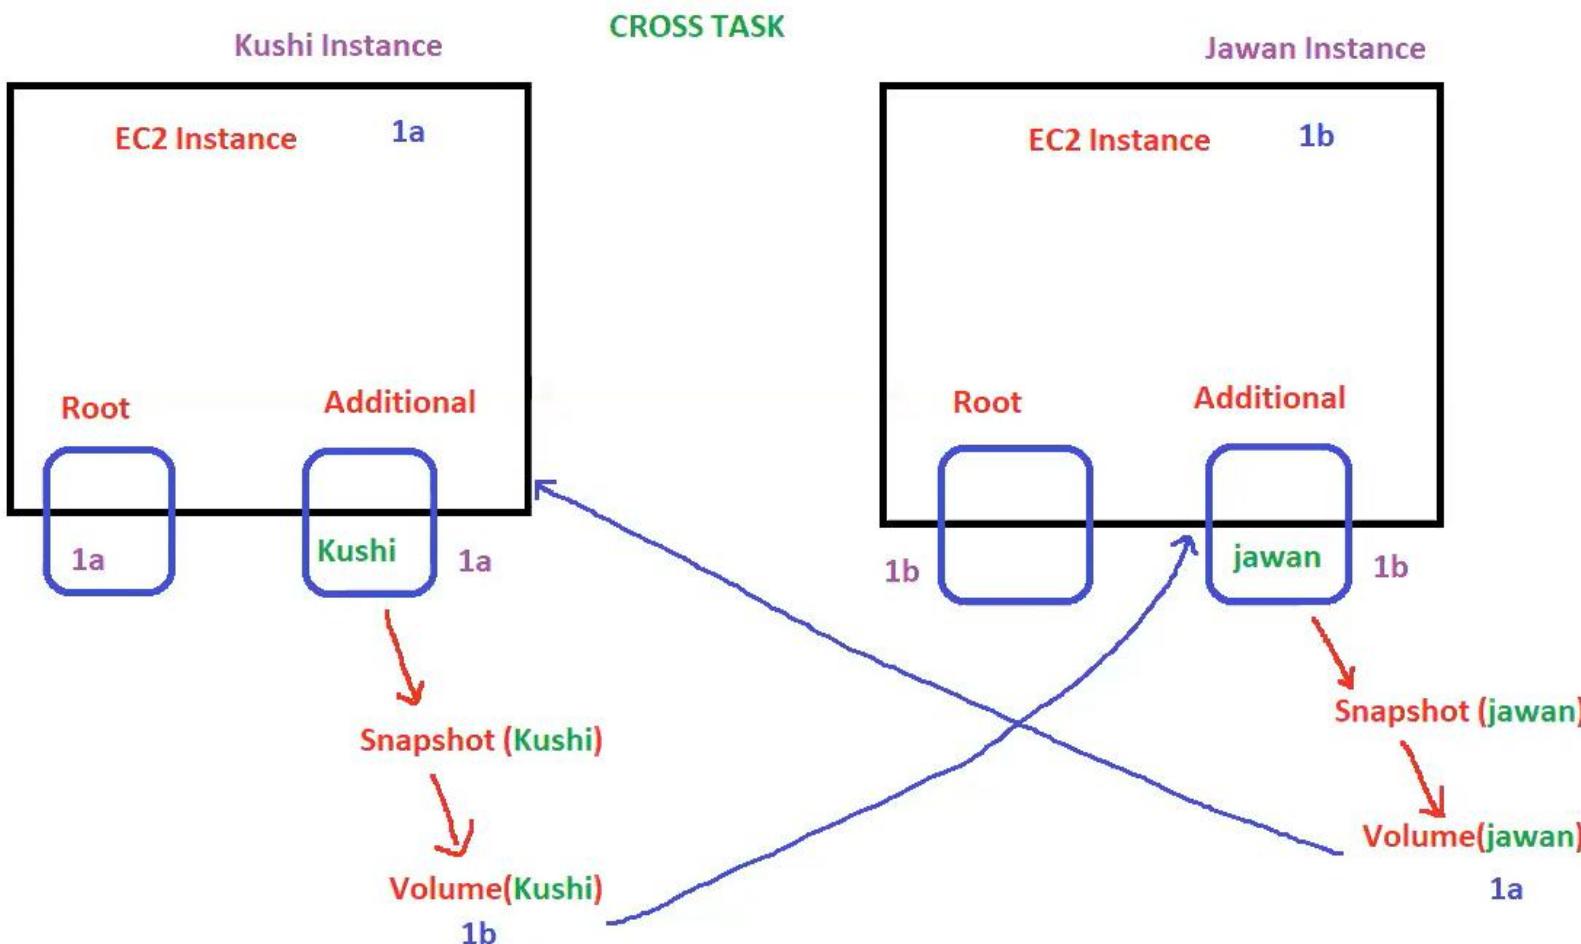
Not Encrypted --> Encrypted (Copy Option)

All Encryption KEYS are Stored in KMS (Key Management Service)

Access and Secret Keys are for access purpose / Encryption Keys are for Security Purpose

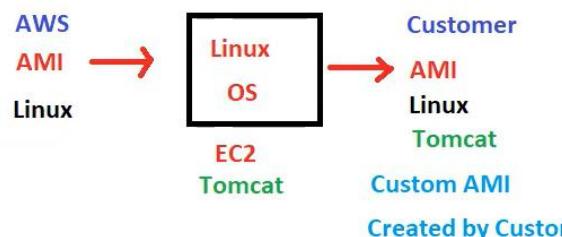
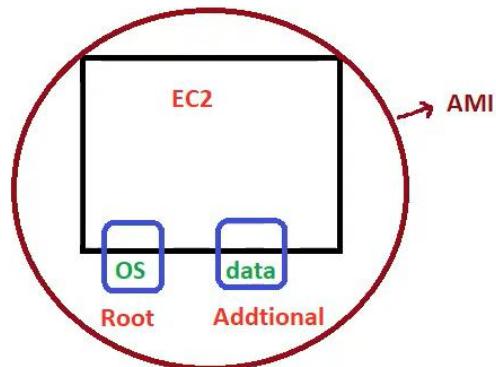
Encryption Snapshots cannot be Shared

How to move the volume from 1 AZ to another AZ



Snapshot ---> Copy of the Volume

AMI --> Copy of the entire EC2 instance (includes volumes also)



AMI's are stored in S3

IMAGES

Copy of the OS is called Image

Image = AMI = Amazon Machine Image

Template of the OS is called AMI

AMI contains OS or OS + Few Apps

Copy of the image includes all configurations that we did on original instance

EC2 instance ----> Image(AMI) ---> EC2 instance

1 AMI, can be used multiple times to launch multiple EC2 instances

AMI's are re-usable, AMI doesn't have any AZ's

You cannot directly use AMI to login, instead launch EC2 instance from the image and then login to the EC2 instance

By default AMI's are Private, If required make it public

AMI's are Regional

AMI's can be copied from one region to another region in the same account

AMI's can be shared from one account to another AWS account using AWS ID

All public images are located at AWS Market Place

Image Contains OS ---> Root Volume (EBS)

---> Root Volume (ISV)

Images are backed by either EBS volumes or ISV

NO need to stop the EC2 instance to create image (but recommended is to STOP)

OS

Windows

**Win Server
2021, 2022, 2023**

**Plain OS
OS + Few Apps**

Linux

**RedHat, CentOS
Ubuntu / SUSE**

If you customize the app on OS, and then take the image ---> Custom AMI
(Or)
Manually
Automatically

Creating images can be automated using EC2 image builder ---> Golden AMI's

EC2 instance

2 Volumes

**Root
Additional**

→ AMI(Image) (2 Volumes)

**↓
2 Snapshots**

When ever you are creating a image, Snapshots will also get created based on how many volumes you have in the EC2 instance

EC2

Windows / Linux

IP: will be provided by AWS

Username : Windows - Administrator

Password : You will get it through KEY-PAIR

IP: will be provided by AWS

Username : Linux - ec2-user

Password : You will get it through KEY-PAIR

KEY-PAIR

Key-Pair is used to retrieve the password of the EC2 instance

We dont have any key-pair by default, we have to create it

We have to create a key-Pair ---> Key-Pair is also called as **PEM file**

Key-Pair extention is .PEM

When ever we launch EC2 instance, the console will ask you to create and attach a key-pair to the EC2 instance

You can create multiple key-Pairs

1 Key-Pair can be attached to multiple EC2 instances at the same time

EC2 instance can have only 1 Key-Pair attached at any point

For Every EC2 instance the password is unique / Different

Once the PEM file is attached, you cannot change the PEM file to the EC2 instance

Keep the PEM file in safe and secure it

Every time you retrieve the password usign pem file, you will get the same password from the EC2 instance

Key-Pair = Combination of Public Key + Private Key

AWS has public key, Customer has private key(Pem file)

For Windows use Remote Desktop Protocol Client

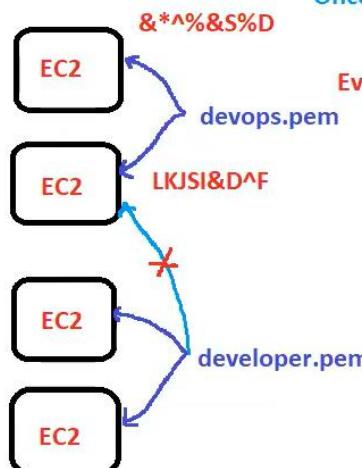
For Linux

Putty ---> doesnt support pem file

---> ppk file



DevOps



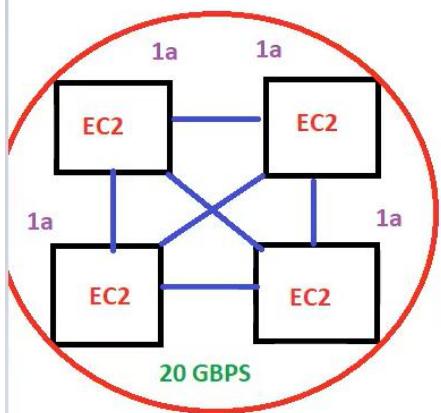
Developer

Windows EC2 instance ---> administrator / Pem file

Protocol ---> RDP / 3389

Linux EC2 instance ---> ec2-user / pem file

Protocol ---> SSH / 22



Cluster Networking Instances

Cluster = Group of Servers / EC2 instances ---> This group is called Placement Group

When you launch a new EC2 instance, the EC2 service will place the instance such a way that all your instances are spread out across dif hardwares

Cluster Placement Group : Grouping the instances in same / single AZ / Same Rack, High Performance, Low HA

Spread Placement Group : EC2 instances are spread across AZ's. High HA, Critical Applications

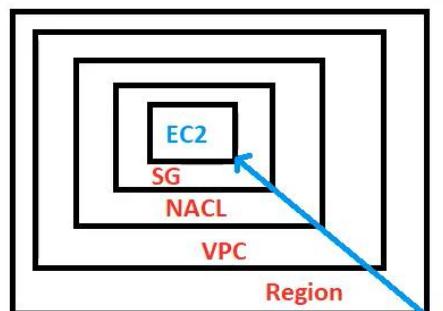
Per 1 AZ = 7 instances

Partition Placement Group : Across AZ's, Max Partitions = 7

1 Partition can contains 100's of EC2 instances

Placement Group recommended to have same /
homogenous instance type

When you are launching the EC2 instance, you can select which placement group you want to keep the instance

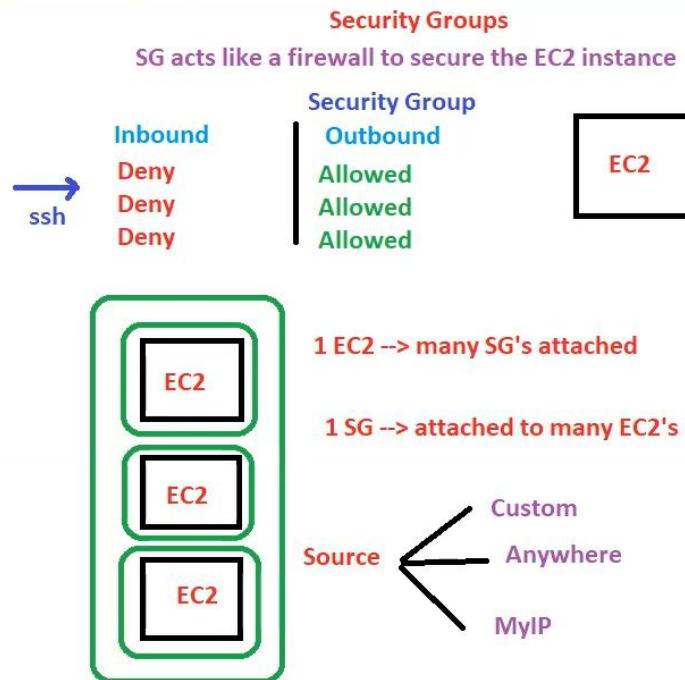


Allow http / 80
inbound rules
Allow SSH / 22
Rest are all DENY

RDP ---> is used to login / connect remotely to the Windows EC2 instance

SSH ---> is used to login / connect remotely to the Linux EC2 instance

http / https ---> is used to access the applications



If you allow any inbound rules, you MUST allow on outbound rule ---> STATELESS

I
NACL's are STATELESS

Firewall = Which stops unauthorized access to the network

Firewall = Security Group Allow / Deny

Security Group = SG stops unauthorized access to the EC2 instance

Security Group is used to secure the EC2 instance

Security Group has 2 Rules

Inbound Rules ---> Which allows the traffic towards EC2 instance

Outbound Rules ---> Which send the traffic outside EC2 instance

By default, Inbound rules are DENY / Outbound rules are ALLOWED

Is it possible to DENY protocol in SG ? NO (Because by default inbound rules are already DENY)

In SG, you can only ALLOW, you cannot DENY

Every EC2 instance must have atleast 1 SG attached

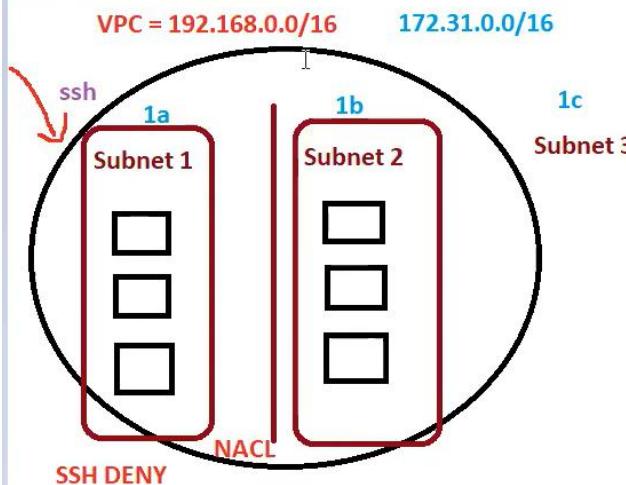
You can create multiple SG's and you can attach multiple SG's to a single EC2 instance

AWS EC2 has a default Security Group

A brand new SG , all inbound rules DENY and outbound rules are ALLOWED

If you allow any inbound rules, you no need to allow on outbound rules STATEFUL

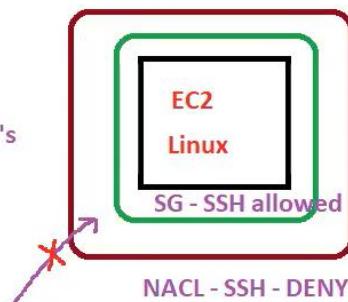
Security Groups are STATEFULL



1 Subnet = 1 NACL

1 Subnet should not be associated to multiple NACL's

1 NACL can have multiple Subnets



NACL = Network Access Control List
 Another layer of Security to the EC2 instance
 If you want to tight security then go with NACL
 Like SG, NACL also has inbound rule and Outbound rules
 NACL will hit first and then SG
 SG is EC2 instance level, NACL is Subnet Level
 NACL's are Optional

Security Group
 Inbound rules and Outbound rules

Default SG
 SG will hit after NACL
 By default , inbound rules are DENY
 You cannot DENY on SG
 If you create any new SG, inbound rules are DENY and Outbound rules are ALLOWED

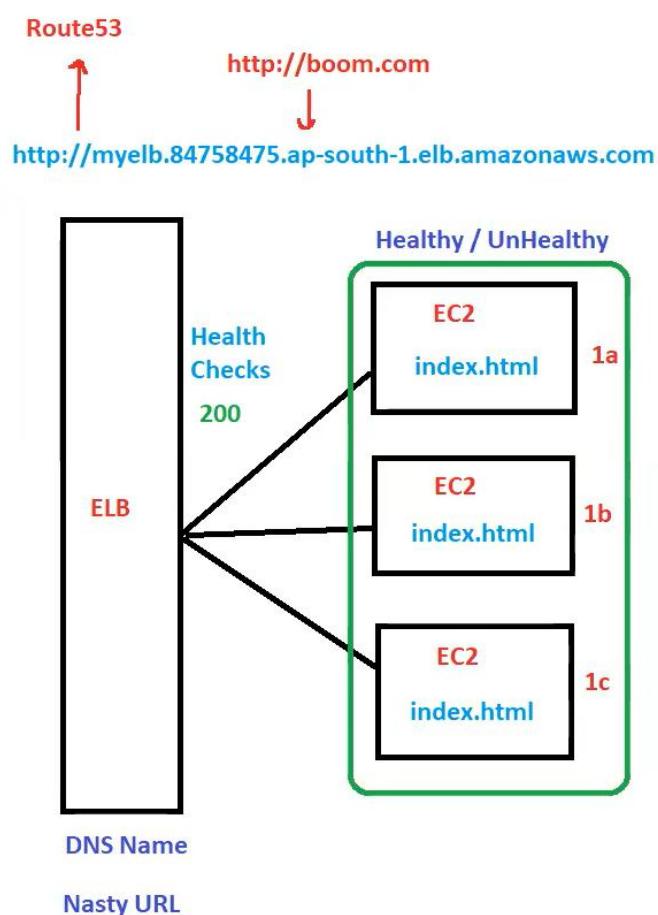
SG;s are Instance level
 SG's are STATEFULL

If you allow any inbound rules, you no need to allow on outbound rules

1 Subnet is associated to 1 AZ
 1 Subnet cannot be associated to multiple AZ's at the same time
 1 AZ can have multiple Subnets

NACL

Inbound rules and Outbound rules
 Default NACL
 NACL will hit first and then SG
 By default , inbound and outbound rules are ALLOWED
 On NACL, we can DENY and Allow Also
 If you create any new NACL, inbound rules are DENY and Outbound rules are DENY
 NACL's are Subnet Level
 NACL's are STATELESS
 If you allow any inbound rules, you MUST allow also on outbound rule



Elastic Load Balancer

- ELB which distribute the traffic to multiple EC2 instances across AZ's
- ELB is completely managed by AWS
- ELB is a Service by AWS, not Server for us
- ELB can be accessed by URL (DNS Name), you cannot login to ELB
- ELB has the IP address, but these IP's are Dynamic
- AWS always recommend to use the ELB DNS Name not IP address

Types of Elastic Load Balancers

- | | |
|---|---|
| Application Load Balancer
ALB
Latest Generation
default choose is ALB
HTTP and HTTPS
Best for micro-services

Routing Features
Host based routing
Path based routing
String parameter routing
etc

like?course=aws | Network Load Balancer
NLB
Latest Generation
TCP and UDP
Extreme High Performance

Network Level
1 Static IP per AZ |
|---|---|

Classic Load Balancer

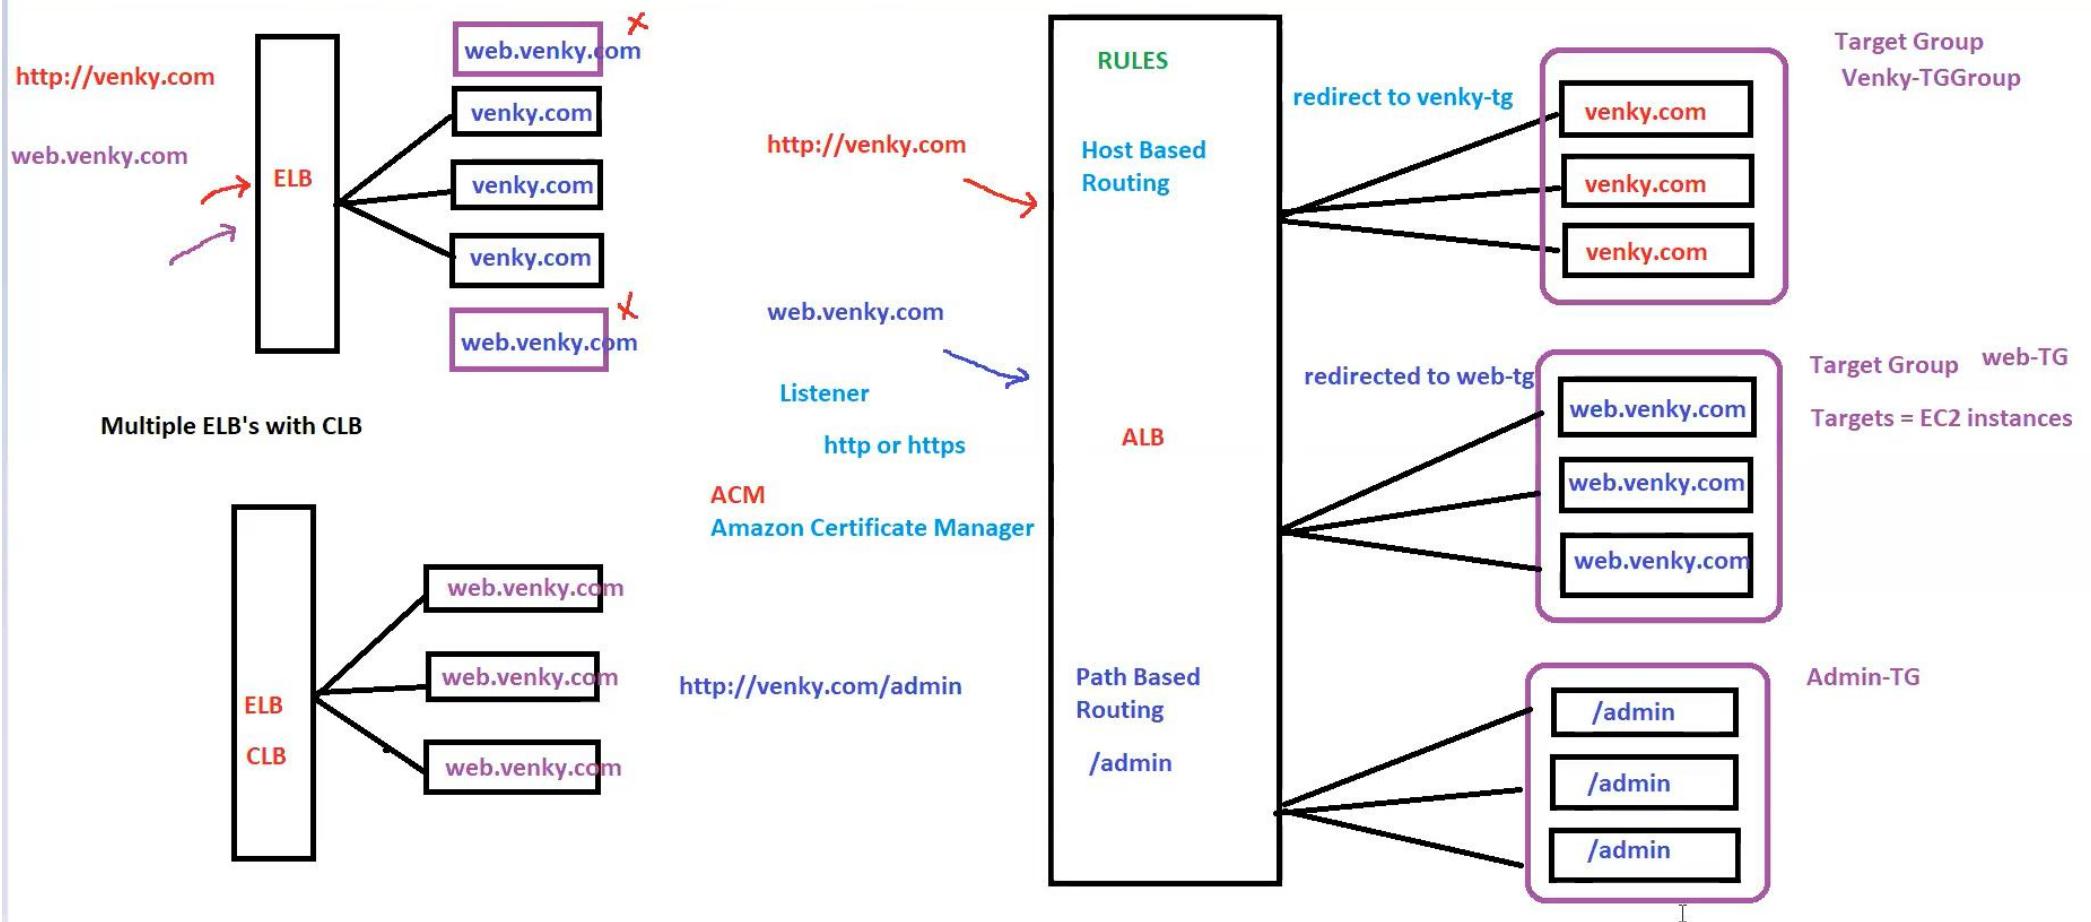
- CLB
- HTTP, HTTPS and TCP
- Previous Generation

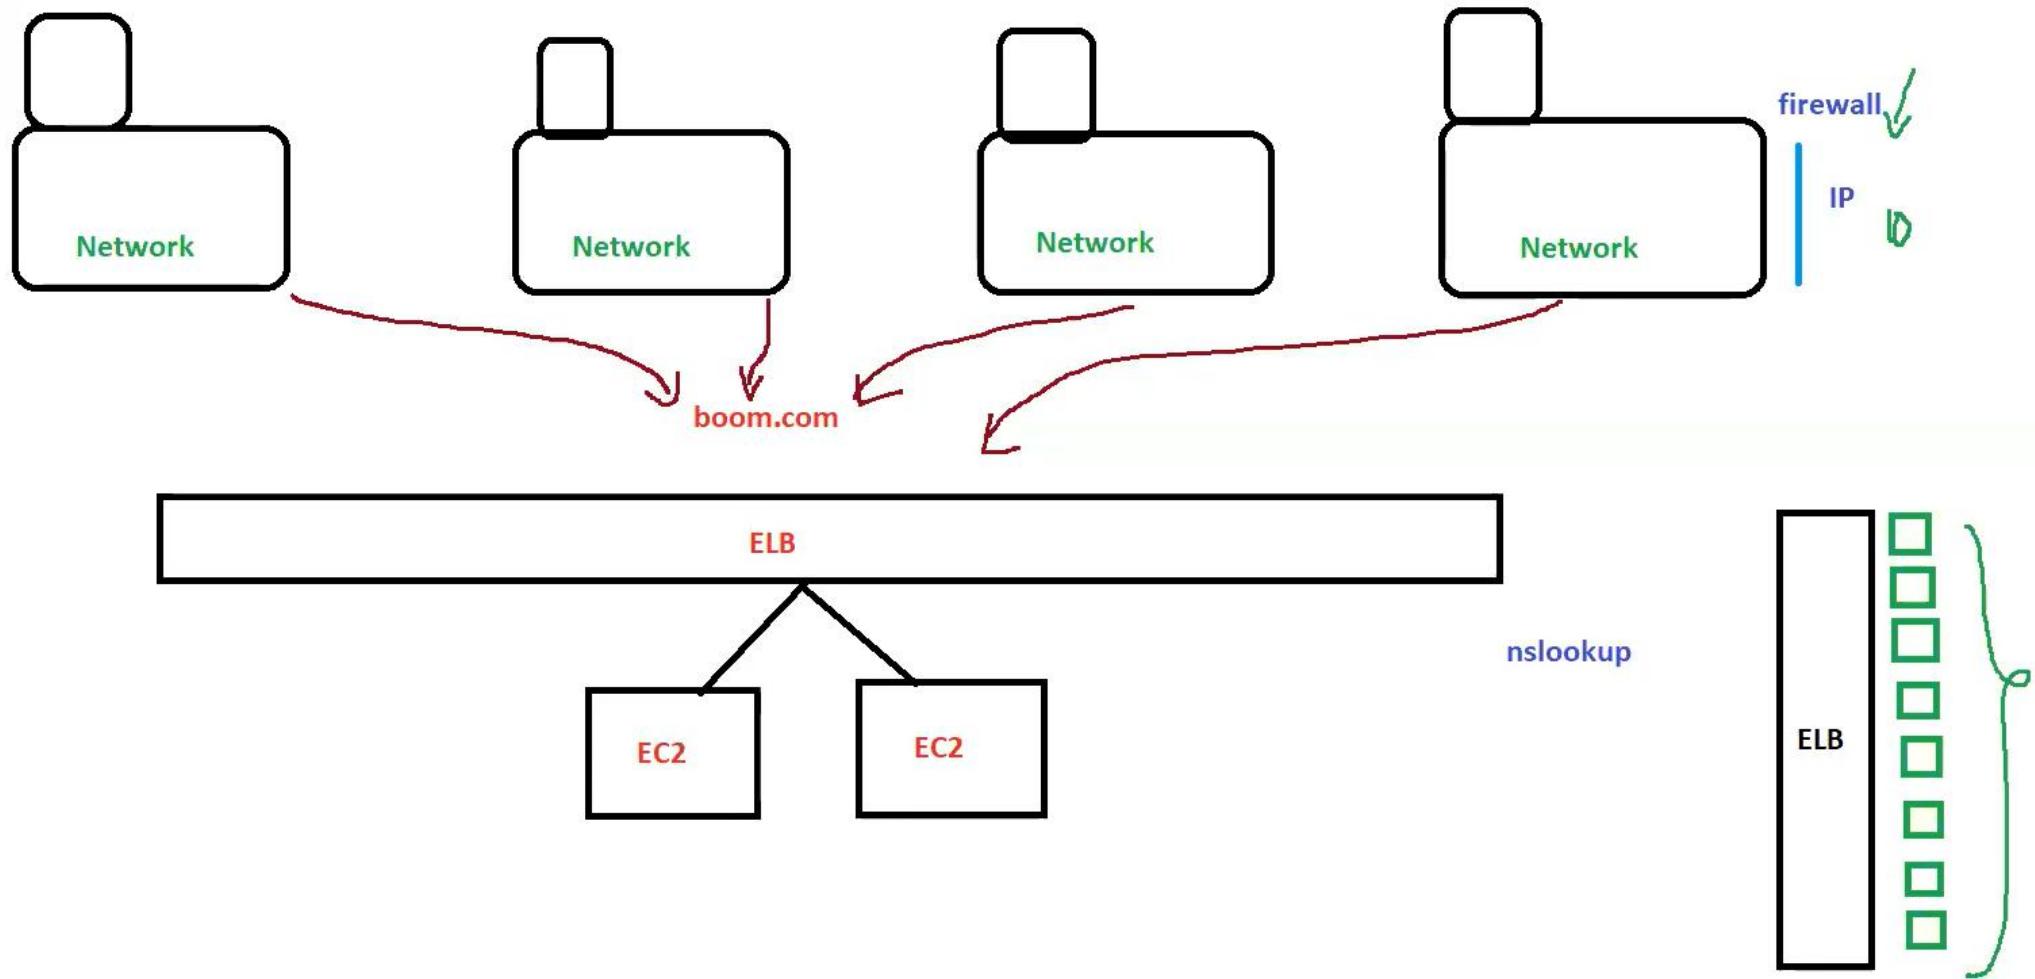
Gateway Load Balancer

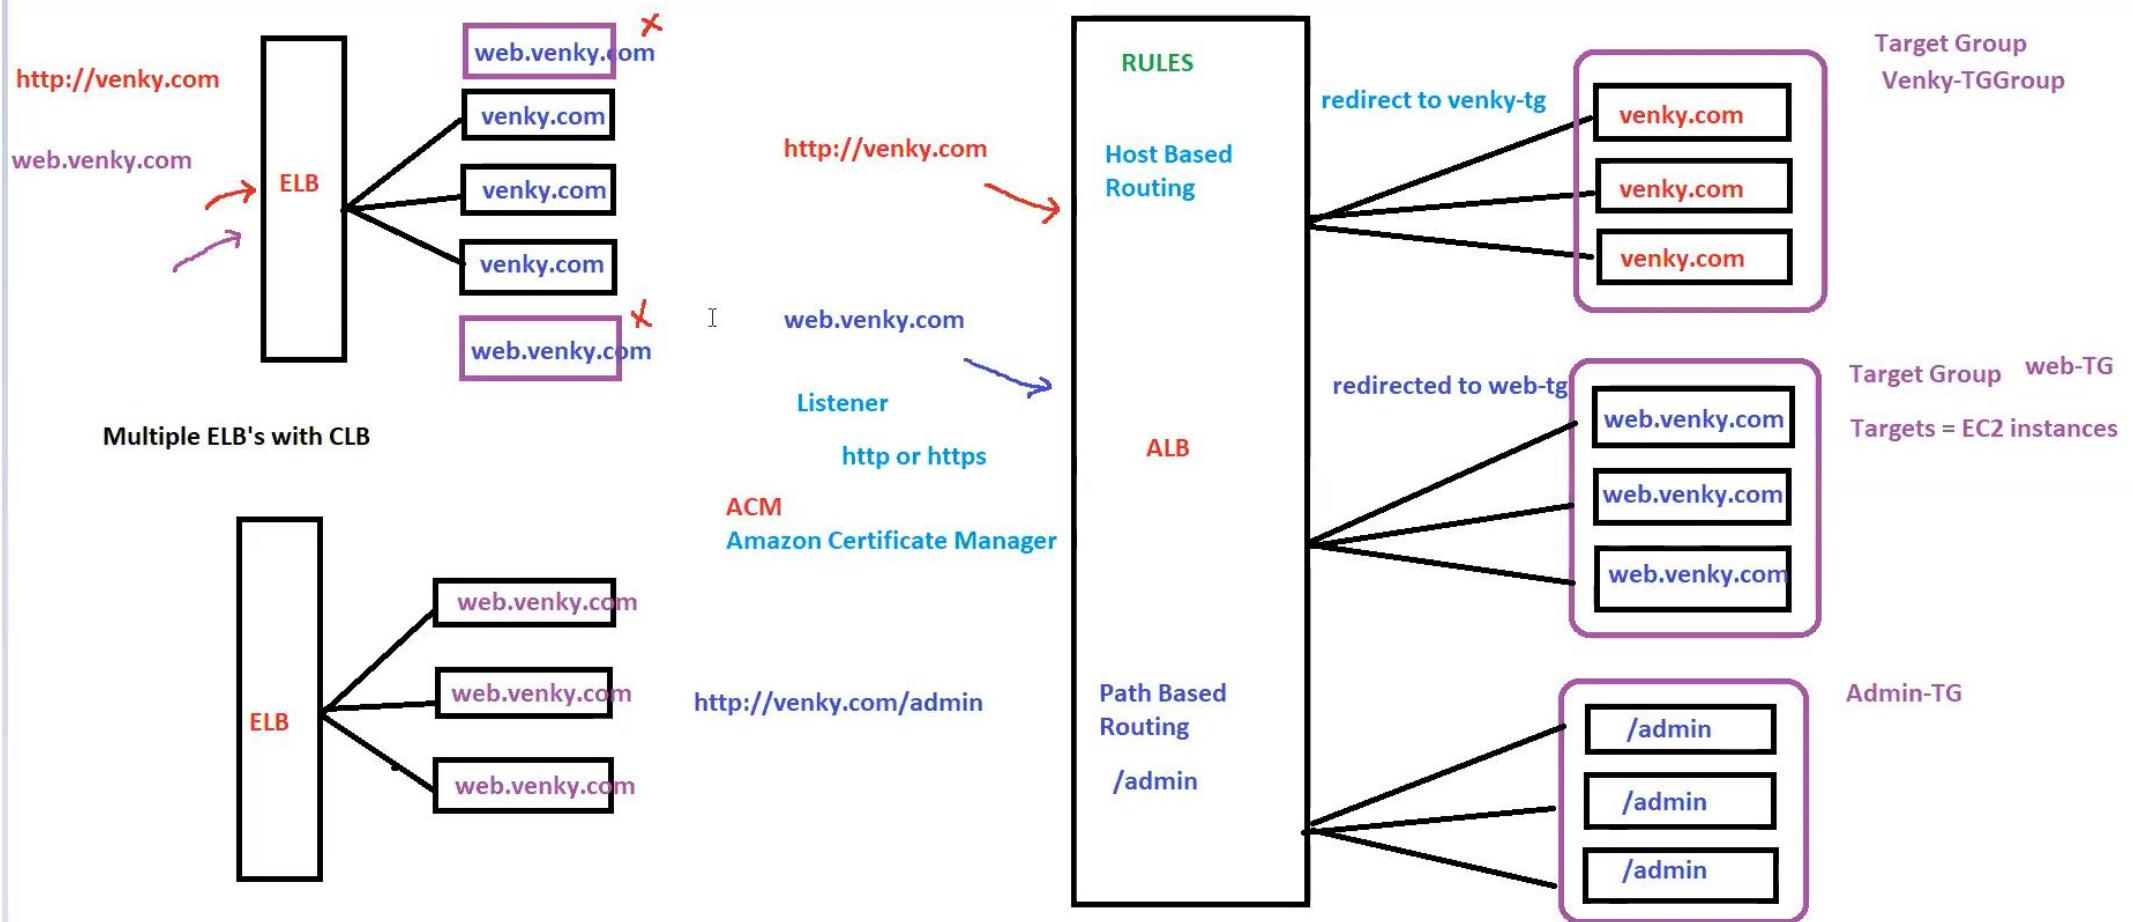
- GLB
- Latest Generation

Deploy, manage, and scale a fleet of 3rd party network virtual appliance in AWS
 . Firewall, Prevention Systems etc

GENEVE protocol on 6081







Public IP

Which can be accessed from internet

Public IP is optional

Public IP is Dynamic

If you STOP and START the EC2 instance,
Public IP will be Changed

AWS assign Public IP to the EC2 instance

Types of IP's

Private IP

Which cannot be accessed from the internet

Private IP is mandatory

Private IP are mostly Static

With in the VPC and through VPN

Elastic IP (EIP)

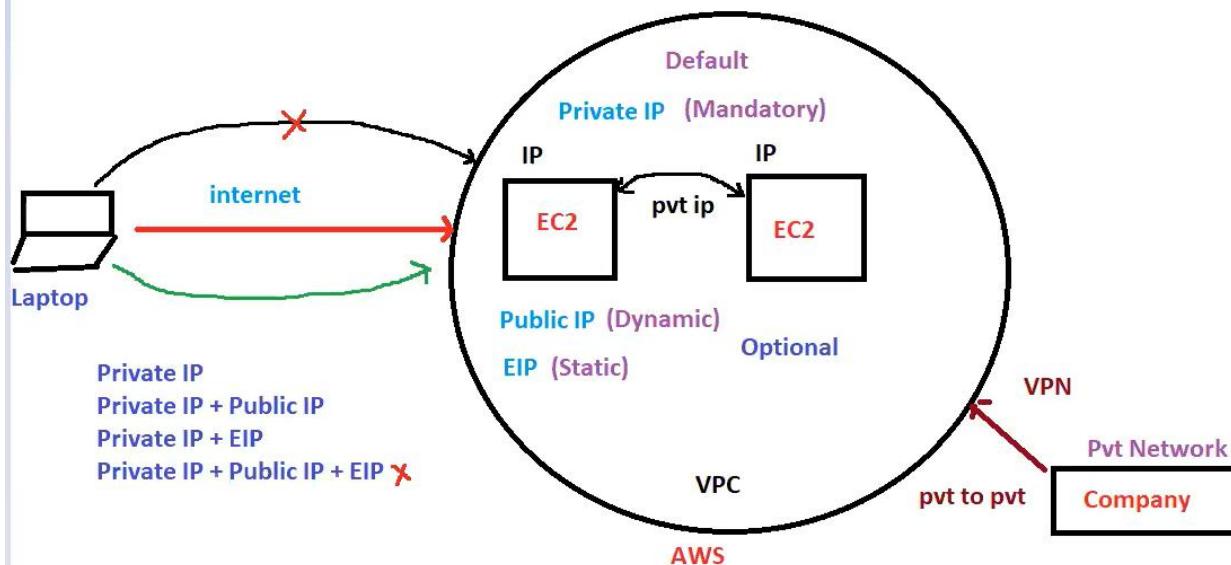
Same as Public IP, EIP is Static

If you STOP and START the EC2 instance EIP will not be changed

5 EIP's are FREE

EIP has to be associated to the EC2 instance

If you have not associated / Not used EIP to any instance, it will be charged if those EIP's are below 5 also, Dont keep it idle, Idle EIP's are Charged



Instance Meta-Data = Data about instance

From Console, you get the instance meta-data information

From CLI, follow this below URL

<http://169.254.169.254/latest/meta-data/> (use this inside EC2 instance only on CLI)

User-Data : BootStrap Script

The Script which you provide will run at the boot time of the EC2 instance

User-Data will run only for the very first time of launching EC2 instance

Linux = Shell Script

Windows = PowerShell

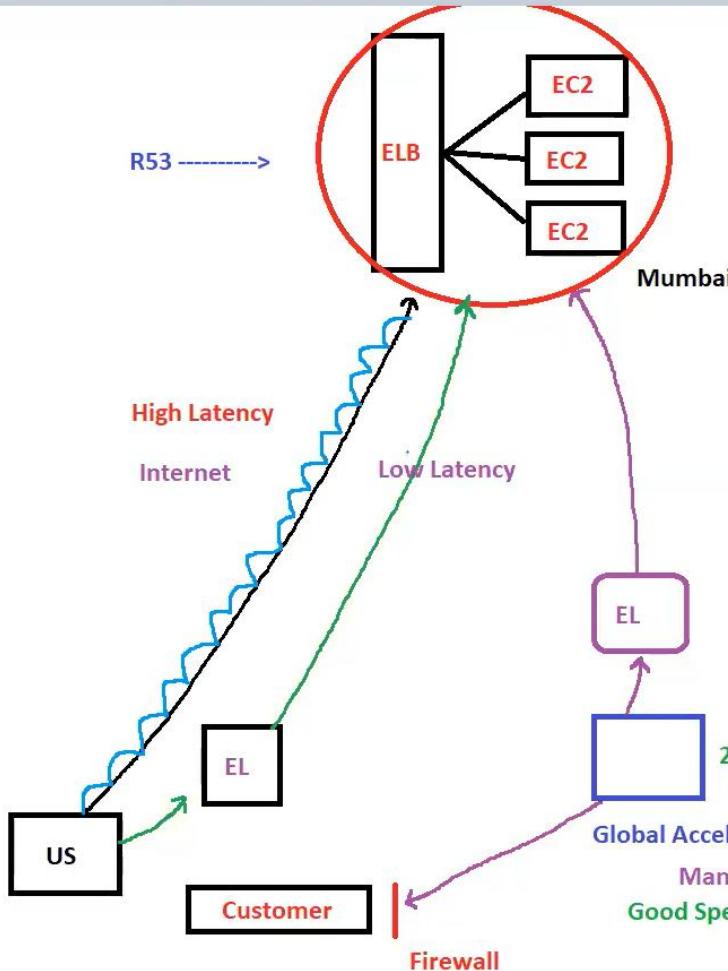
The 7 Steps

Select AMI	Select Instance Type	Instance Configuration	Select Storage	Select Security Group	Add Tags	Attach a PEM File	Launch
Linux Windows	t2.micro	How many instances ? Public IP user-data all configurations	EBS Volumes Root (30GB, 10GB) Additional Volumes gp2	Default SG	Name = LinuxInstance	Create it Or use it	



Unicast IP : One Server holds one IP address

AnyCast IP : All Servers hold same IP address and the client is routed to the nearest one



. Leverage the AWS internal network to route to your application

. 2 Anycast IP are created for your application

. The Anycast IP send traffic directly to the Edge Locations

. The Edge Locations send traffic to application

DDoS protection through WAF and AWS Shield

Billable

GA vs CloudFront

They both use the AWS global network and its Edge Locations around the world
Both Services are integrated with AWS Shield for DDoS protection

CloudFront

- . Improves performance for both Cache content (images and videos)
- . Dynamic Content and Static also
- . Content is Served from Edge Locations

Global Accelerator

- . Improve performance for a wide range of applications
- . It supports 2 Anycast Static IP
- . Whitelisting IP and Disaster Recovery

User ---> boom.com ---> R53 ---> ELB
User ---> boom.com ---> R53 ---> CF ---> ELB
user ---> boom.com ---> R53 ----> GA ----> EL ---> ELB

CloudWatch is all about Alarms, Events and Logs



Host Level Metrics
CPU
Network
Disk
Status Check

B: 6:00 6:05 6:10 6:15
D: 6:00 6:01 6:02 6:03



EC2
App Server
Instance id

Schedule / Cron Job
EventBridge

EVENTS / Event Bridge

Pending/Running, Stopping/Stopped, Shuttingdown/Terminated
Instance state change notification

Stopped
Pending/Running

Cloudwatch is used to monitor performance of AWS Resources

CloudWatch

CloudWatch can monitor all AWS Services

CloudWatch can monitor only Host level metrics (Default Metrics)

CPU, Network, Disk and Status Checks are AWS Provided Default Metrics

Custom Metrics (Memory)

Basic Monitoring = Every 5 mins interval data points, FREE

Detailed Monitoring = Every 1 min interval data points, Billable

NameSpace = Group of Metrics/ Collection of related metrics

CloudWatch can monitor Lambda Fns, Microservices, containers, K8s etc

Alarms

EC2 ---> CPU > 90% ---> Notification (SNS)

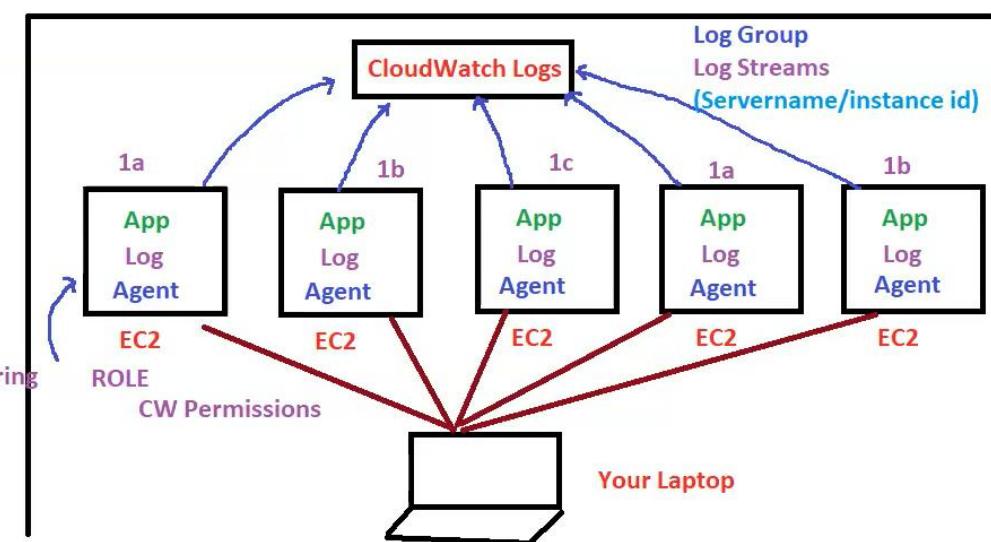
Alarms can also do some actions

(Terminate, STOP, Reboot, Recover)

Alarm has 3 States In Alarm, OK, Insufficient

Composite Alarms

- . CW alarms are on single metrics
- . Composite alarms are monitoring the states of multiple alarms
- . AND or OR Conditions can be used



CloudWatch Logs

Launch EC2 instance (Amazon Linux2)
Create a ROLE (permissions - CW)
and attach to the EC2 instance
Login to the EC2 instance
Install cloudwatch agent
configure the agent
start the agent



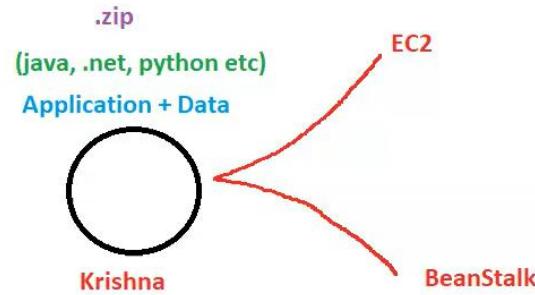
FYI: "File" parameter is related to the log path. Change the log path as per the requirement

```
sudo -s
yum install -y awslogs
cd /etc/awslogs
ls
cat awscli.conf
vi awscli.conf
--> Press i --> change the region
---> Press Esc , :wq! --> Enter
cat awscli.conf
cat awslogs.conf
systemctl start awslogsd
```

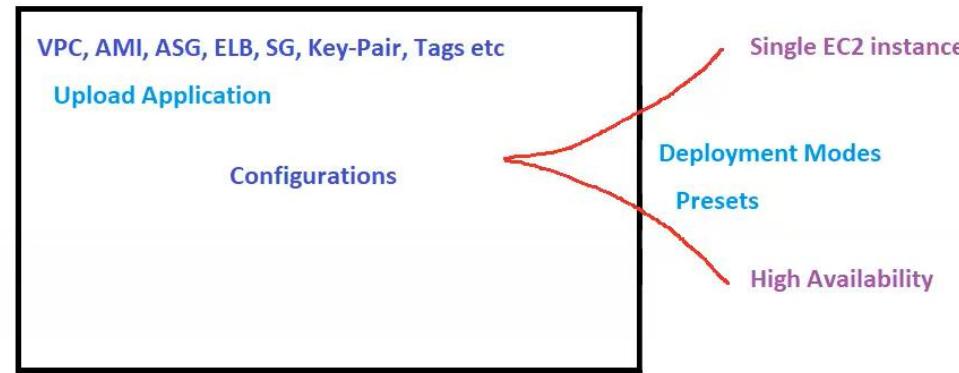
see the logs in cloudwatchlogs

Backbone of BeanStalk is EC2

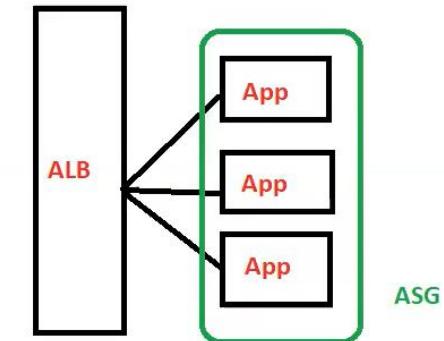
BeanStalk is Free, AWS Charges only for Resources launched by BeanStalk

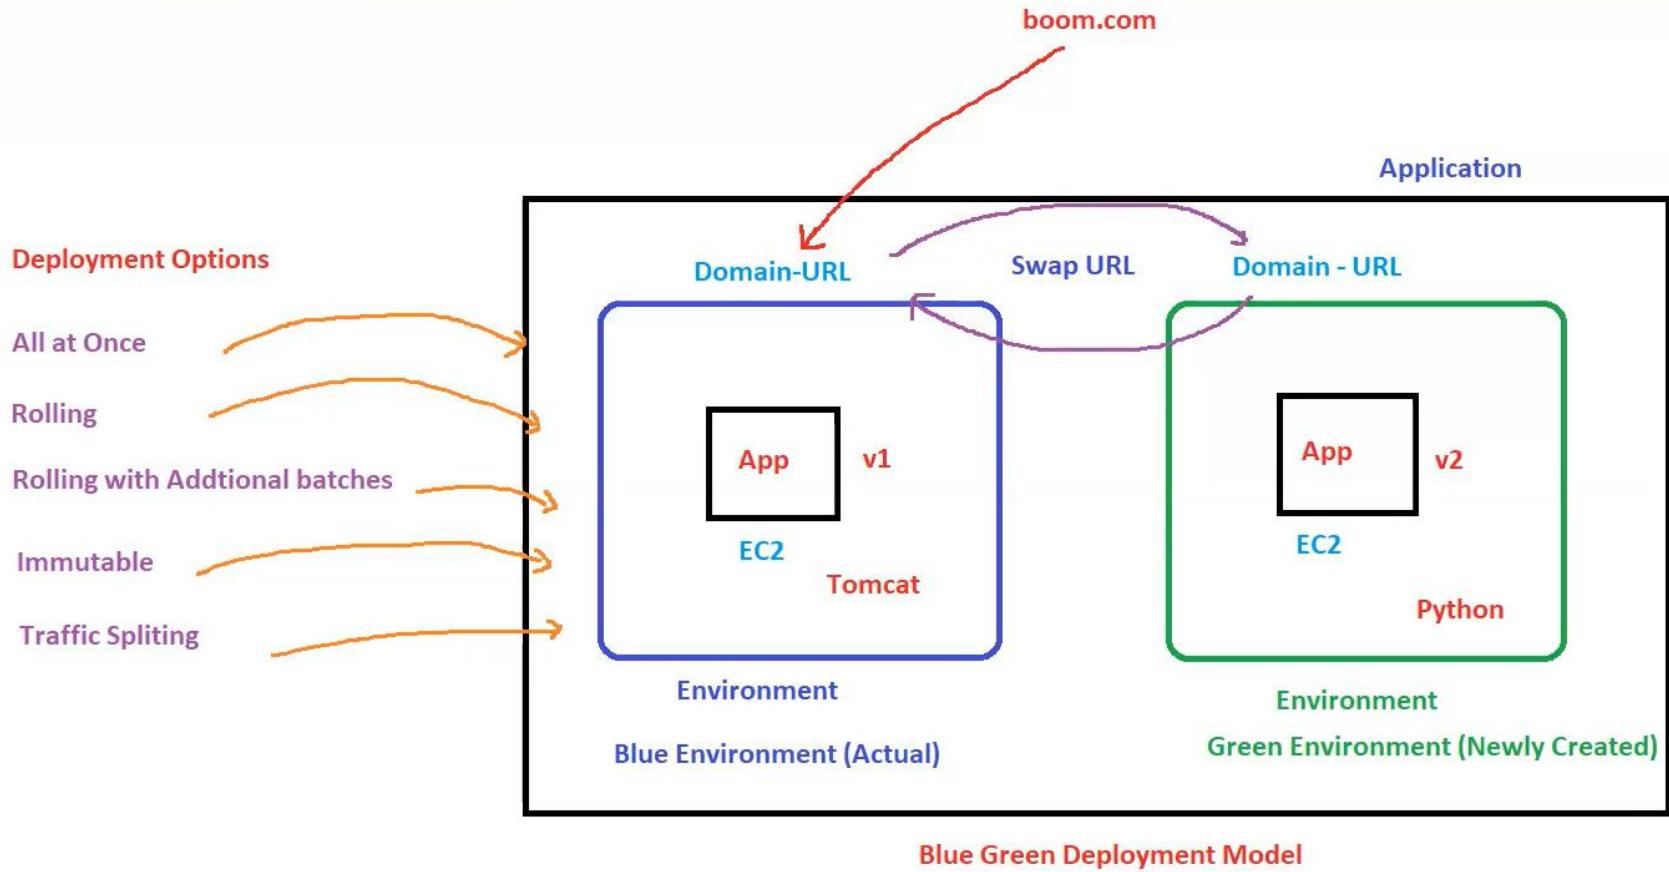


ElasticBeanStalk
ElasticBeanStalk is used for easy and quick deployment of applications in AWS
PAAS = Application + Data

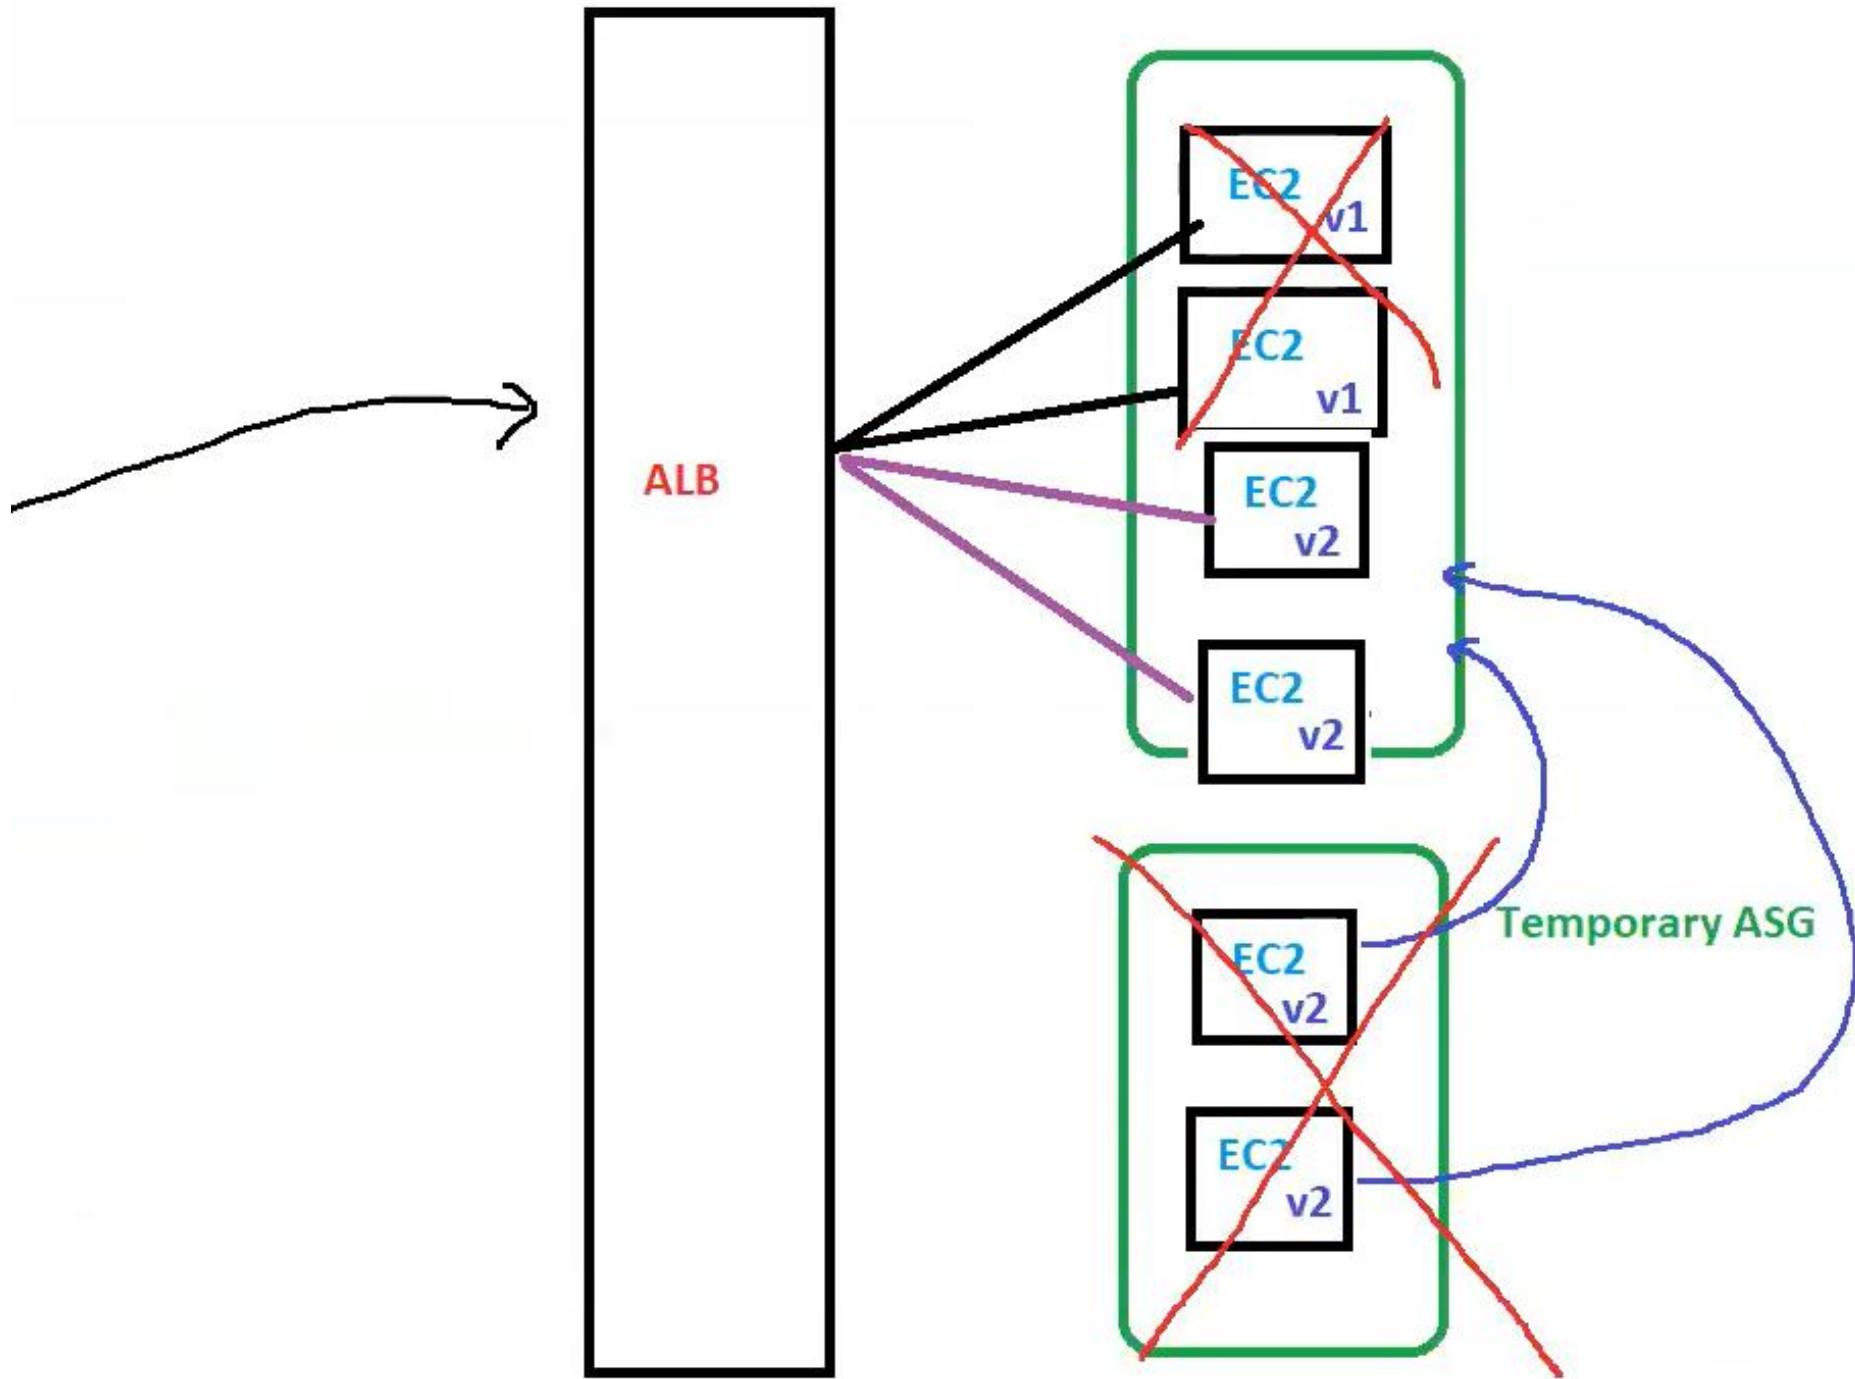


ElasticBeanStalk
Platform | Tomcat
 | Node Js
 | Python
 | GO
 | Docker
 | .NET
 | etc





I
Manual



S3 is Object Based Storage

In S3, we can store all FLAT Files

You can only upload, download and access files from S3

The Files in S3 cannot be executed

You cannot install OS, DB etc in S3

S3 is unlimited storage

S3 support Static Website Hosting

S3 is cheaper than EC2

S3 is Serverless

Bucket = Container of Objects

Object = File

KEY = Name of Object

S3 = Simple Storage Service

S3 is Global

Bkt1 ---> Mumbai
Bkt2 ---> Ireland



Bucket

Bucket is Regional

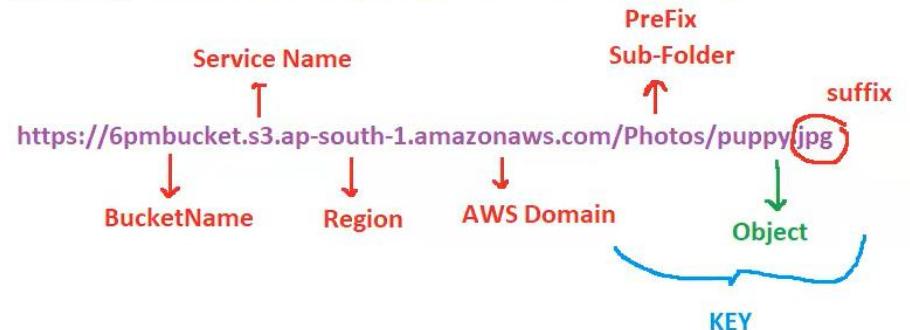
Bucket Names are universal / Unique

No nested buckets / Bucket under bucket cannot be created

You can create the buckets in different Regions

Max Number of Buckets you can created in your AWS account is 100 (Soft limit)

By default , buckets are Private, if required we can make it public

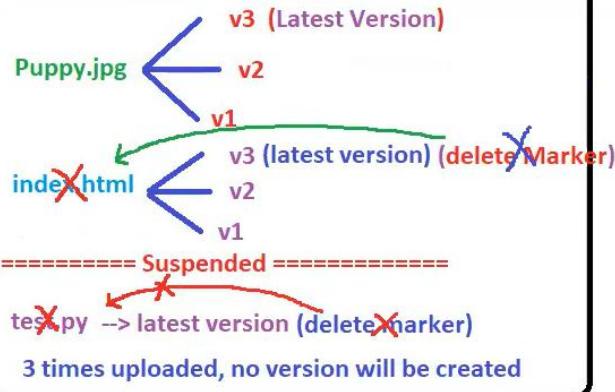


S3 is a WORM Model

Write Once and Read Many

S3 is mostly used for Read Purpose

Versioning is Enabled



6pmversionbucket

If you upload a object after versioning is suspended, the latest version will be created as usual
But, if you update the original object, versioning files are NOT CREATED

If you delete the original object, Delete marker is applied
If you delete the delete marker, Object will not be restored

In Suspended state, existing versioning files will not have any impact

S3 Versioning

Versioning is like a Backup tool

By Default, Versioning is not enabled on the bucket

Versioning is enabled on bucket level, but applied on Object Level

Version ID is always unique

Versioning files can be downloaded anytime

If you delete the original object, delete marker is applied on the latest version

If you want the object to be restored, delete the delete marker and your object is restored

If you want to restore the old version, download and upload it back to the bucket

Delete marker is applied only on the latest version, not for old/previous version

You cannot download the delete marker, you can only delete it

Based on the requirement we need to enable versioning, be careful on larger objects

Once you have enable versioning , you cannot disable it, you can SUSPEND it

S3 is unlimited storage

Min object size = 0 Bytes, Max Object Size = 5TB

You can have unlimited number of objects having 5TB each in single bucket

For Single PUT, you can upload max 5GB

Multi-Part Upload (MPU): Break the files into multiple chunks and upload it chunk by chunk : it can be done through CLI

AWS recommend, if you have object > 100MB use MPU

While uploading the object into S3, Selecting Storage class is mandatory

Standard Frequently Access (FA)

This is used for frequently access data

Default Storage Class

Regular Purpose

No Retrieval Charges

Availability = 99.99%

Durability = 11 9's

Min Object Size = 0 Bytes

Intelligent Tier

Unknown Access Patterns

Availability = 99.9%

Durability = 11 9's

Min Duration = 30 days

Standard Infrequently Access (IA)

This is used for infrequently access data

Cheaper than FA

Retival Charges apply

Once in a month

Availability = 99.9%

Durability = 11 9's

Min Object Size = 128KB

Min Duration = 30 days

Glacier

Infrequently access data

Archiving Purpose

Vault = container of Archives

Archive = zip/object

1 Archive can be upto 40TB

Unlimited number of archives

1000 Vaults

Retival Charges apply

Storage Classes

Availability = Anytime

Durability = Longtime

Reduced Redundancy Storage (RRS)

Frequently access but NOT CRITICAL

No Retirval charges

AWS doesnt recommend to use this Storage Class

Cheaper than others

Availability = 99.99%

Durability = 99.99%

One-Zone IA

Infrequently access but NOT CRITICAL

Retival charges apply

Availability = 99.5%

Durability = 11 9's

Min Object Size = 128KB

Min Duration = 30 days

Life Cycle Management

Life Cycle Rules

It is possible to move the objects from one storage class to another storage class automatically

Glacier Retival Options

Expedited = 1 to 5 mins

Standard = 3 to 5 hours

Bulk = 5 to 12 hours

Availability = 99.99%

Durability = 11 9's

Min Duration = 90 days

Deep Glacier

Min Duration = 180 days

LCM is created on bucket level,
but applied on Object Level

Transition and Expiration

You can setup S3 features for entire bucket
or for a prefix (sub-folder)

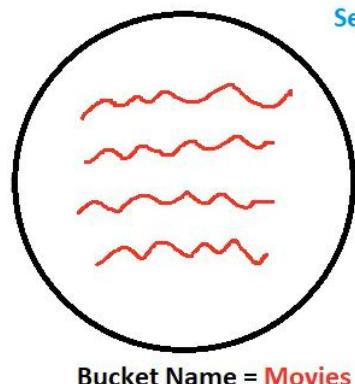
LCM Rule	Current Versions	Previous Versions
Transition		
FA --> IA (30 days) --> Glacier (60 days)		
--> Delete after 365 days	Expiration	

Life Cycle Management

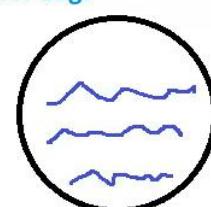
Object Lock
Permanently
Certain period of time

I

Object Level Logs = CloudTrail



Server Access Logs



Who is accessing your buckets
Server access logs are bucket level

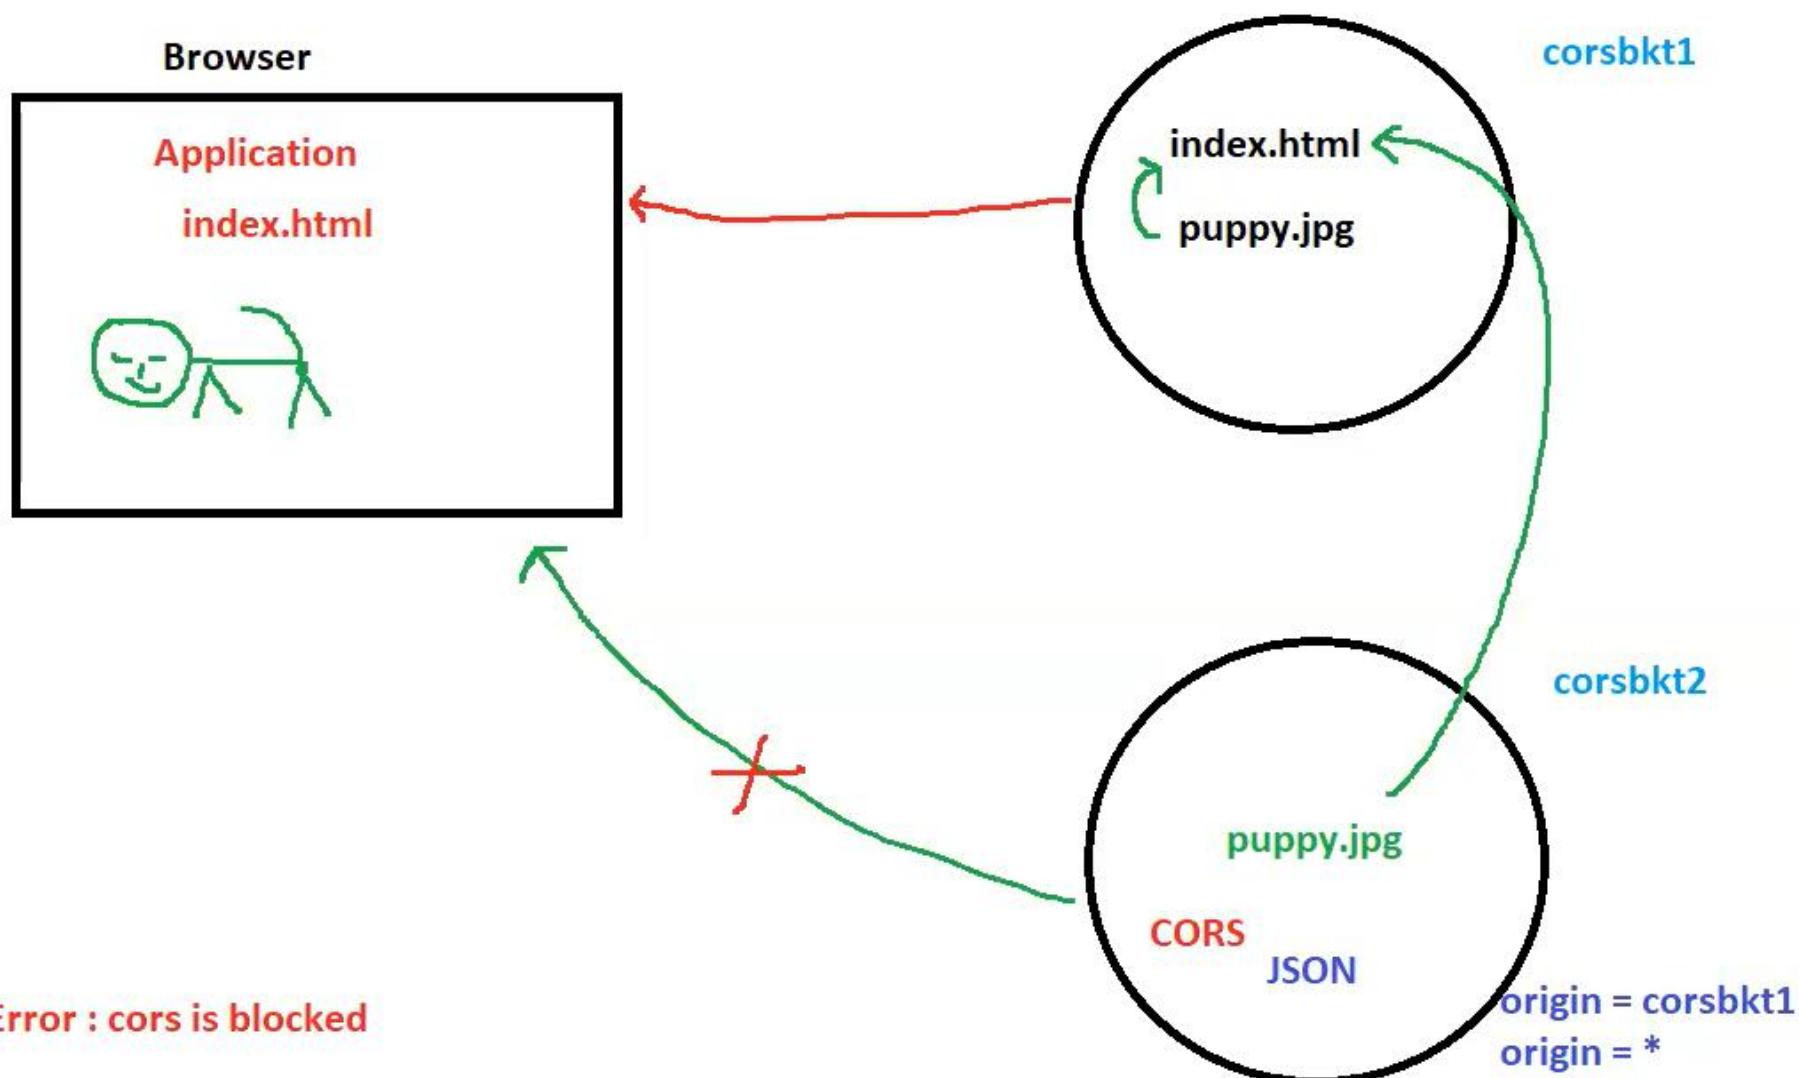
192.168.10.106ambucketpuppy.jpgGETSuccess200

Athena : Analyze the logs directly from S3

IP	Src	Dest	Object	method	Time

SQL Query

CORS = Cross Origin Resource Sharing



CRR = Cross Region Replication

SRR = Same Region Replication

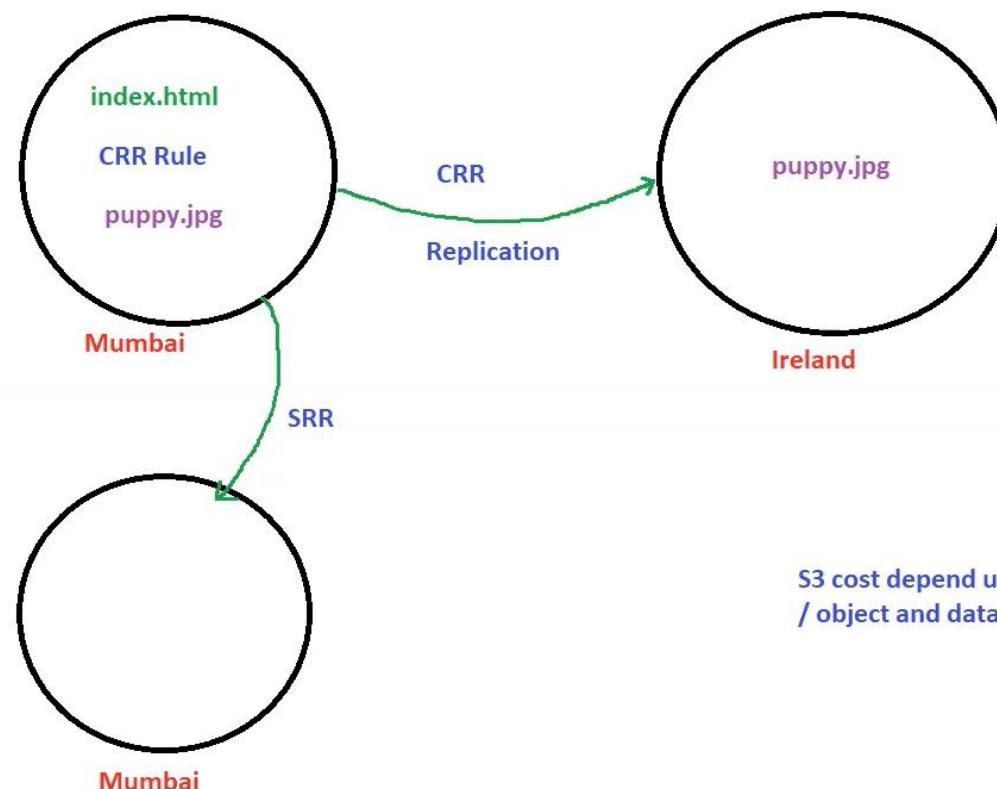
CRR/SRR is not enabled by default

CRR/SRR is enabled on bucket level

CRR/SRR can be for entire bucket or for a Prefix

Versioning is mandatory to have CRR/SRR

Existing objects can also be replicated at the time of creating CRR Rule with Batch Operation Process



CRR buckets can be in same AWS account or different AWS account as well

S3 cost depend upon the size of the bucket / object and data transfer

Encryption can be done in 2 ways

In-Transit Encryption : Encryption while data is moving / Transferring HTTPS

Data at Rest Encryption : Encryption while data is at Rest , KMS

Amazon Certificate Manager (ACM)

is where you can generate HTTPS certificates ([In-Transit](#))

Key Management Service (KMS)

is where you can create encryption keys for data at rest

Encryption

Amazon S3 has 3 types of Encryption

AES - 256

Advance Encryption Standard

Server Side Encryption

By default, Bucket Encryption is Enabled

SSE - S3 = (AWS Managed Key) ✓

SSE - KMS = (AWS KMS Key)

SSE- C = (Customer Provided Key)

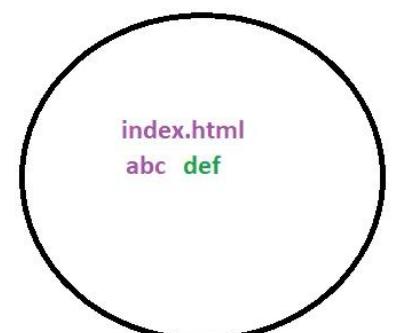
Client Side Encryption : Should be handled by Customer

In-Transit Encryption : HTTPS

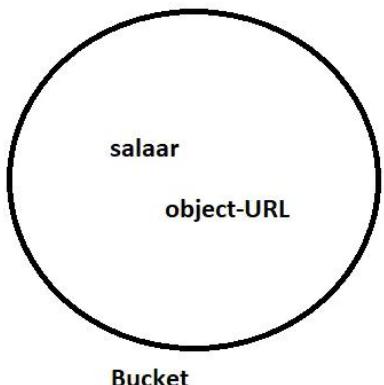
S3 Data Consistency Models - 2 Types

Read after write consistency for PUTS of New Objects

Eventually consistency for OVERWRITES of PUTS and DELETES



Pre-Signed URL



Endpoint

Endpoint is valid only for certain period of time

Ex : 5 mins

After 5 mins, URL will expire

Temporary Purpose

Transfer Acceleration

Ireland

Virginia

Tokyo

Sydney

internet 15 mins

CDN

Faster Upload

Hyderabad

I
ACL = Access Control List

S3 - Requester Pays

- . In General, bucket owners pay for all S3 storage and data transfer cost associated with their bucket
- . With Requester Pays Buckets, the requester instead of the bucket owner pays the costs of the requests and the data download from the bucket
- . Helpful if you want to share the large data sets
- . The requester must be authenticated in AWS so that AWS knows where to charge (in their AWS account). Cannot be anonymous

S3 Event Notifications

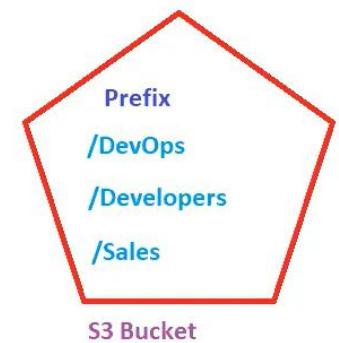
S3 Batch Operations

- . Perform bulk operations on existing S3 objects with a single request

S3 Access Points

Access Point simplify security management for S3 buckets

Users	Access Points
DevOps Users	DevOps AP
Dev Users	Dev AP
Sales Users	Sales AP



Instead of writing critical bucket policies, you can create access points to each prefix and give the DNS Name to the users to access their respective folders in Buckets

Access points can be Public (internet) or Private(VPC)

By default Object Nature is Private

If you create a Private Bucket, Objects are Private

If you create a Public Bucket, Objects are Still Private (If required make the object Public)

Encryption can be done in 2 ways

In-Transit Encryption : Encryption while data is moving / Transferring HTTPS

Data at Rest Encryption : Encryption while data is at Rest , KMS

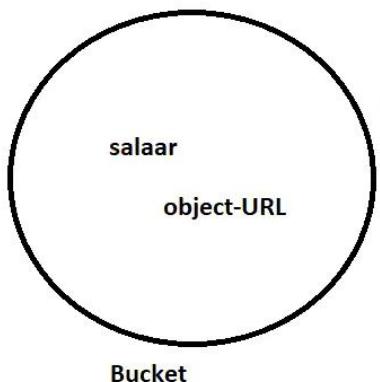
Amazon Certificate Manager (ACM)

is where you can generate HTTPS certificates (**In-Transit**)

Key Management Service (KMS)

is where you can create encryption keys for data at rest

Pre-Signed URL



Encryption

Amazon S3 has 3 types of Encryption

Server Side Encryption

SSE - S3 = (AWS Managed Key) ✓

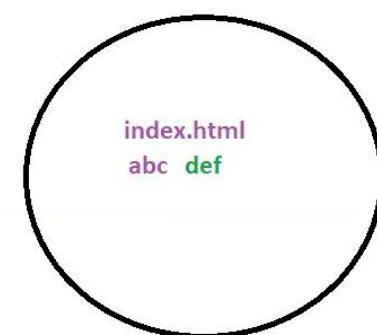
SSE - KMS = (AWS KMS Key)

SSE- C = (Customer Provided Key)

AES - 256

Advance Encryption Standard

By default, Bucket Encryption is Enabled



Client Side Encryption : Should be handled by Customer

In-Transit Encryption : HTTPS

S3 Data Consistency Models - 2 Types

Read after write consistency for PUTS of New Objects

Eventually consistency for OVERWRITES of PUTS and DELETES

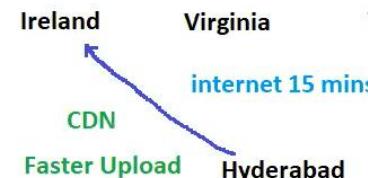
Endpoint

Endpoint is valid only for certain period of time
Ex : 5 mins

After 5 mins, URL will expire

Temporary Purpose

Transfer Acceleration



ACL = Access Control List

Provides governance, compliance and audit for your AWS account

- . CloudTrail is enabled by Default
- . Get an history of events / API calls made in your AWS account
- . Console
- . SDK's
- . CLI
- . AWS Services
- . Can put logs from Cloudtrail into Cloudwatch logs or S3
- . A Trail can be applied to all regions or a single region
- . If resource is deleted in AWS, investigate CloudTrail First

CloudTrail

CloudTrail Events

Management Events

Operations performed on resources in AWS accounts

Ex: attachpolicy, create vpc, etc setting up logs etc

Data Events

By default, data events are not logged (because of high volume operations)

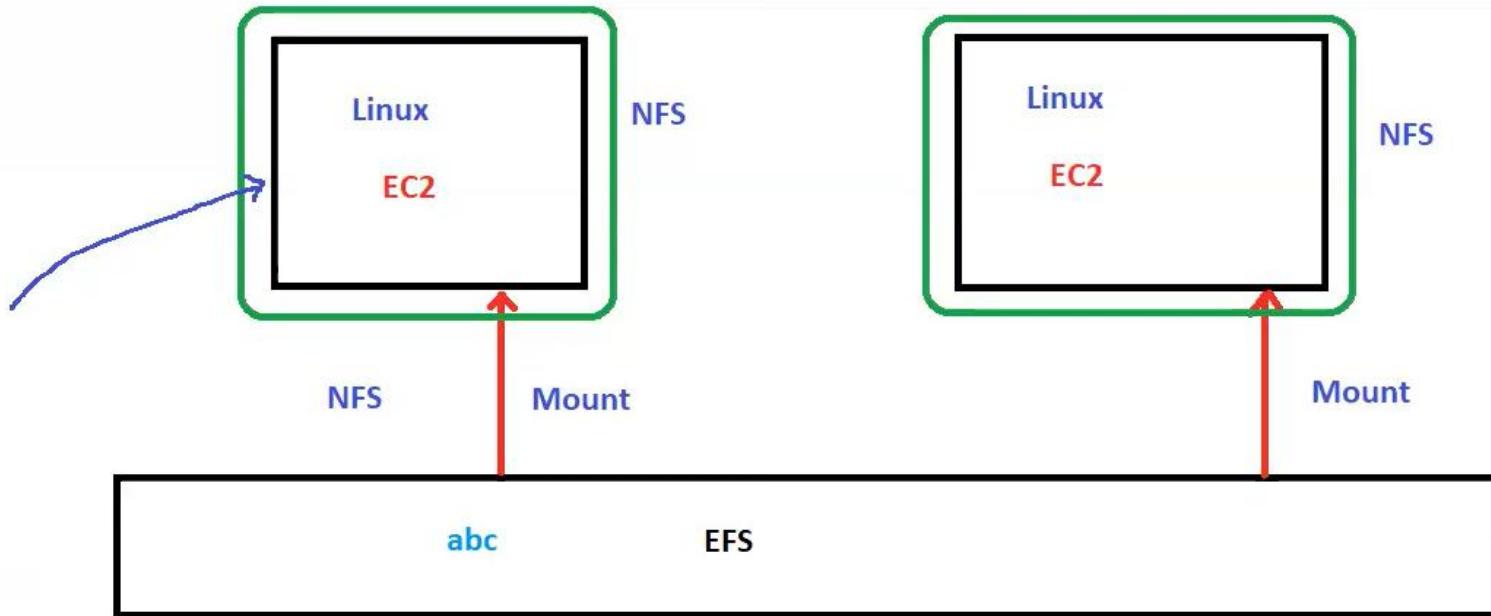
Ex: getobject, deleteobject, putobject etc

Insight Events

- . Enable CloudTrail insights to detect unusual activities
- . inaccurate provisioning
- . hitting service limits
- . bursts of IAM actions

CloudTrail insights analyze normal management events to create a baseline
. It detects write events to detect unusual patterns

Events are stored for 90 days in CloudTrail
To keep beyond this period, log them to S3 and use Athena



1. Create EFS
2. Launch EC2 instance and mount EFS and create some sample files
3. Launch another EC2 instance and mount the same EFS to the EC2 instance and check the files

```
yum install -y nfs-utils
mkdir efs
mount -t nfs4 filesystemdnsname:/ efs/
cd efs
create sample files
```

AWS CLI can be installed on Windows and Linux as well

For Windows ---> AWSCLI.msi

Linux ---> Use commands to install AWS CLI

Configuring the KEYS on the EC2 instance is BAD WAY

AWS Access

Console(GUI) username/pwd	Programmatical (CLI) Access Key and Secret Key
------------------------------	---

IAM User
EC2

Console
CLI

IAM ROLE
S3, EC2

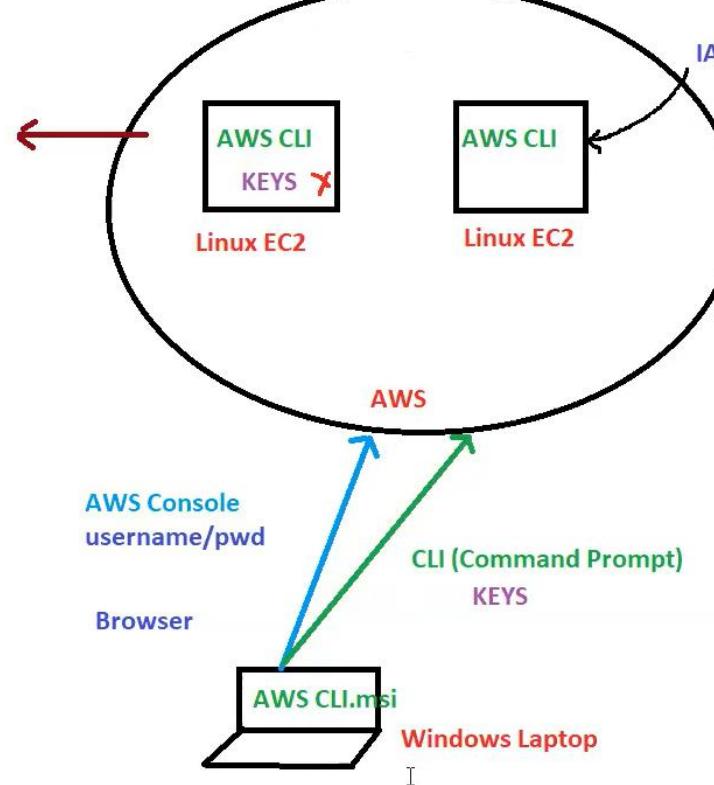
BAD WAY : Launch Linux EC2 instance and install AWS CLI and configure KEYS and run sample commands

GOOD WAY : Launch Linux EC2 instance and install AWS CLI and attach ROLE and run sample commands

Install AWS CLI

Open Command Prompt

```
aws configure
access key
secret key
region : ap-south-1
output = table/json/xml
```



AWS CLI can be installed on Windows and Linux as well

For Windows --> AWSCLI.msi

Linux --> Use commands to install AWS CLI

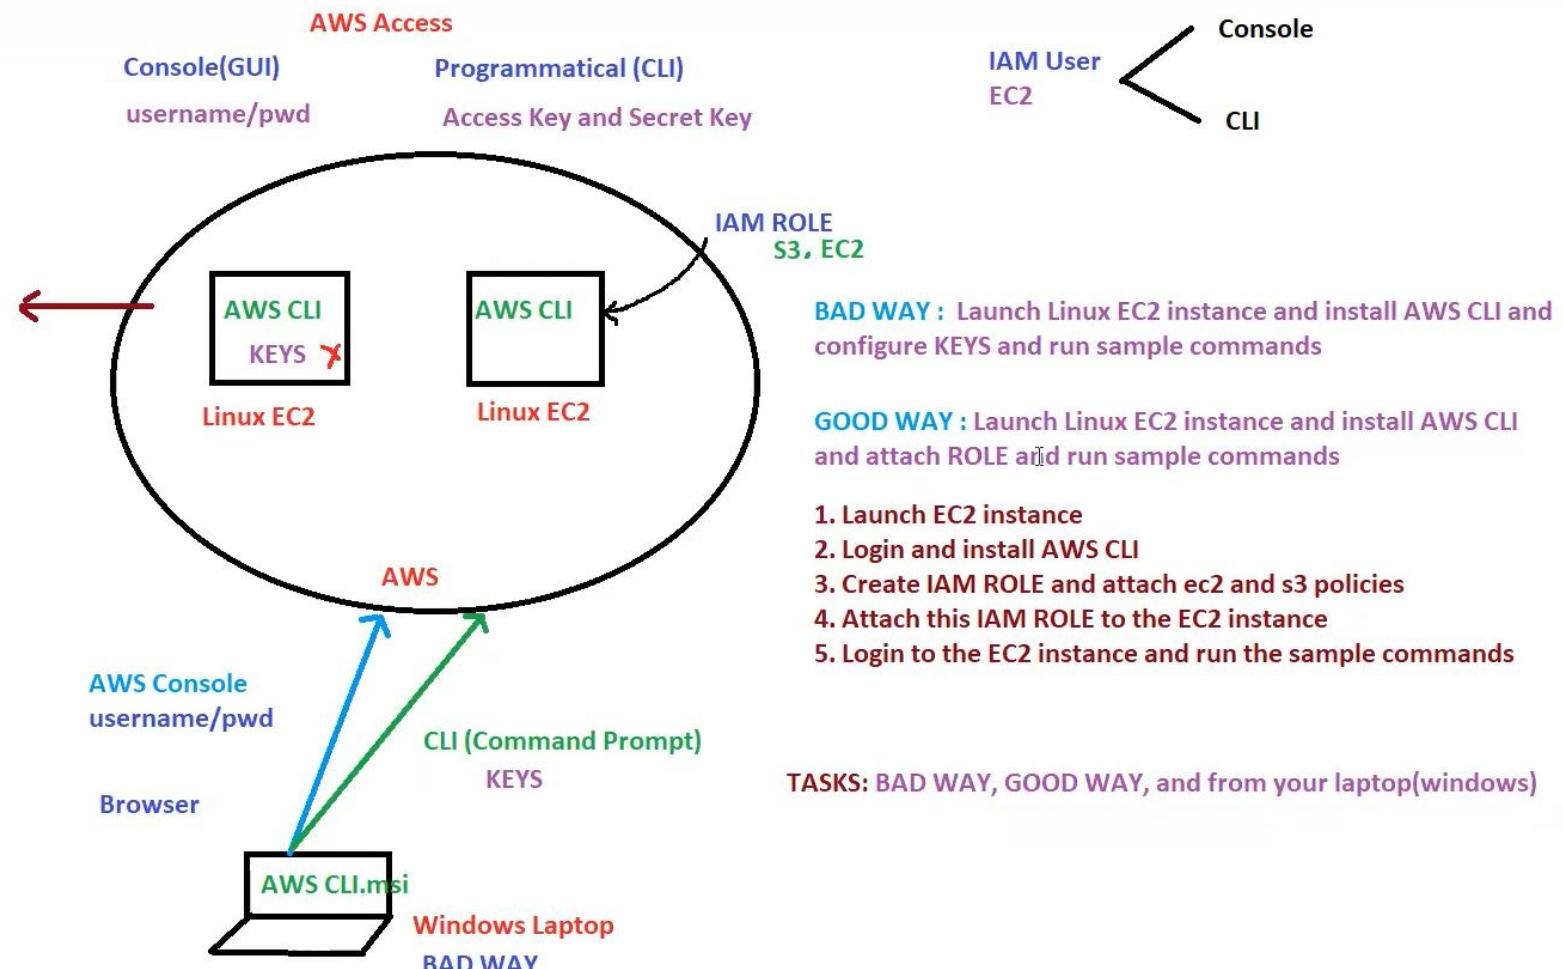
Configuring the KEYS on the EC2 instance is BAD WAY

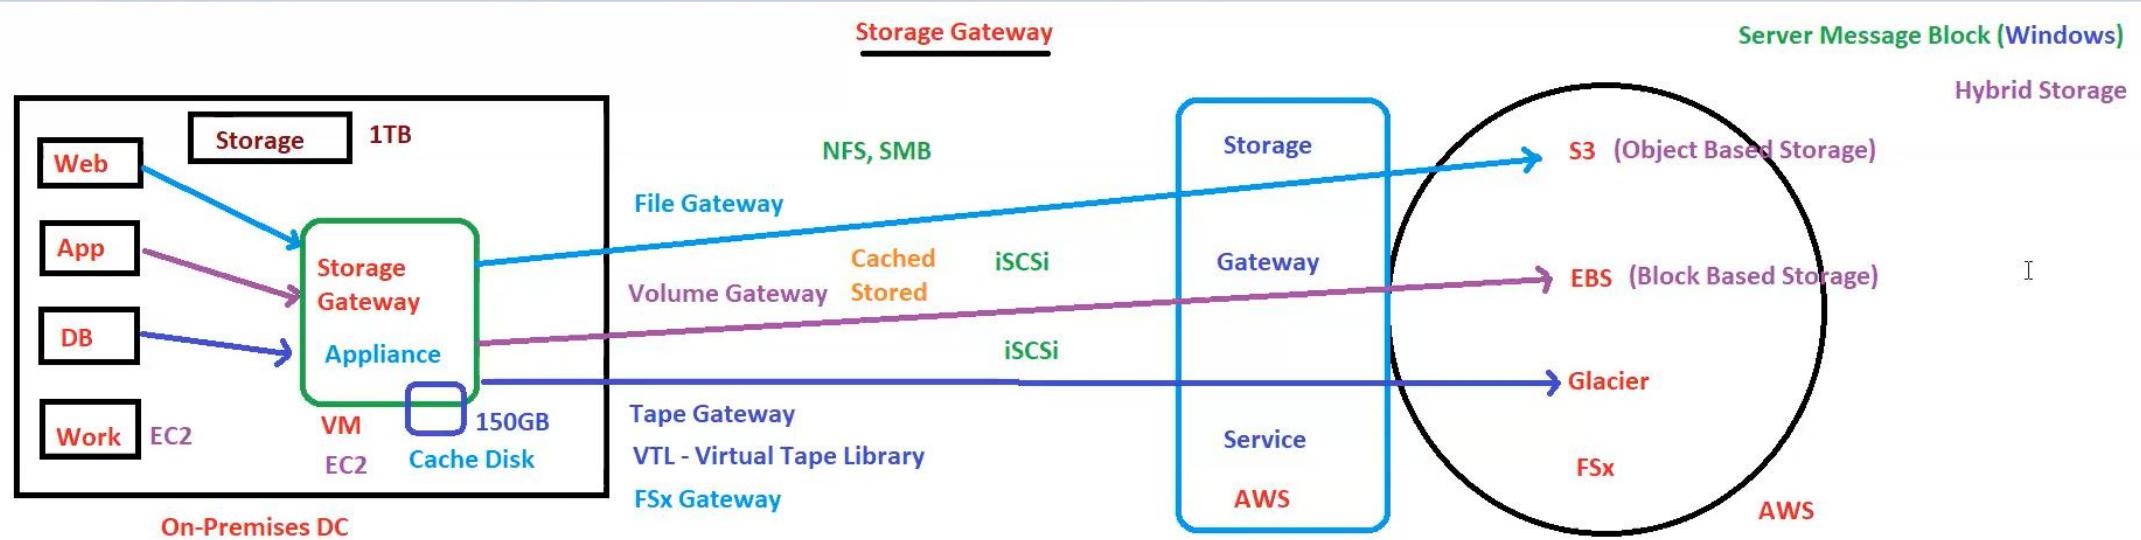
BAD WAY = Configure KEYS
GOOD WAY = USE ROLES

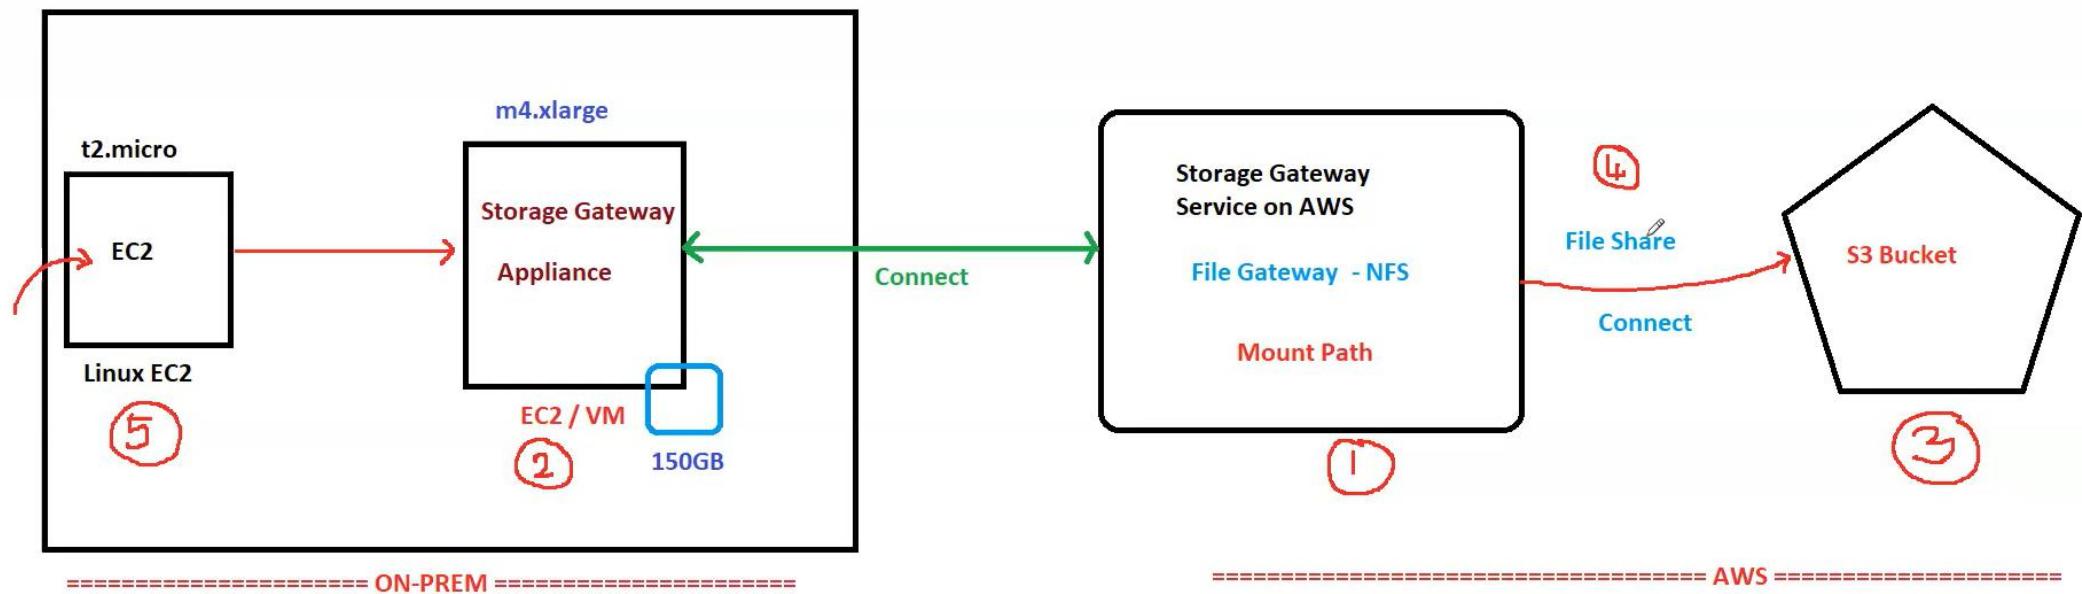
Install AWS CLI

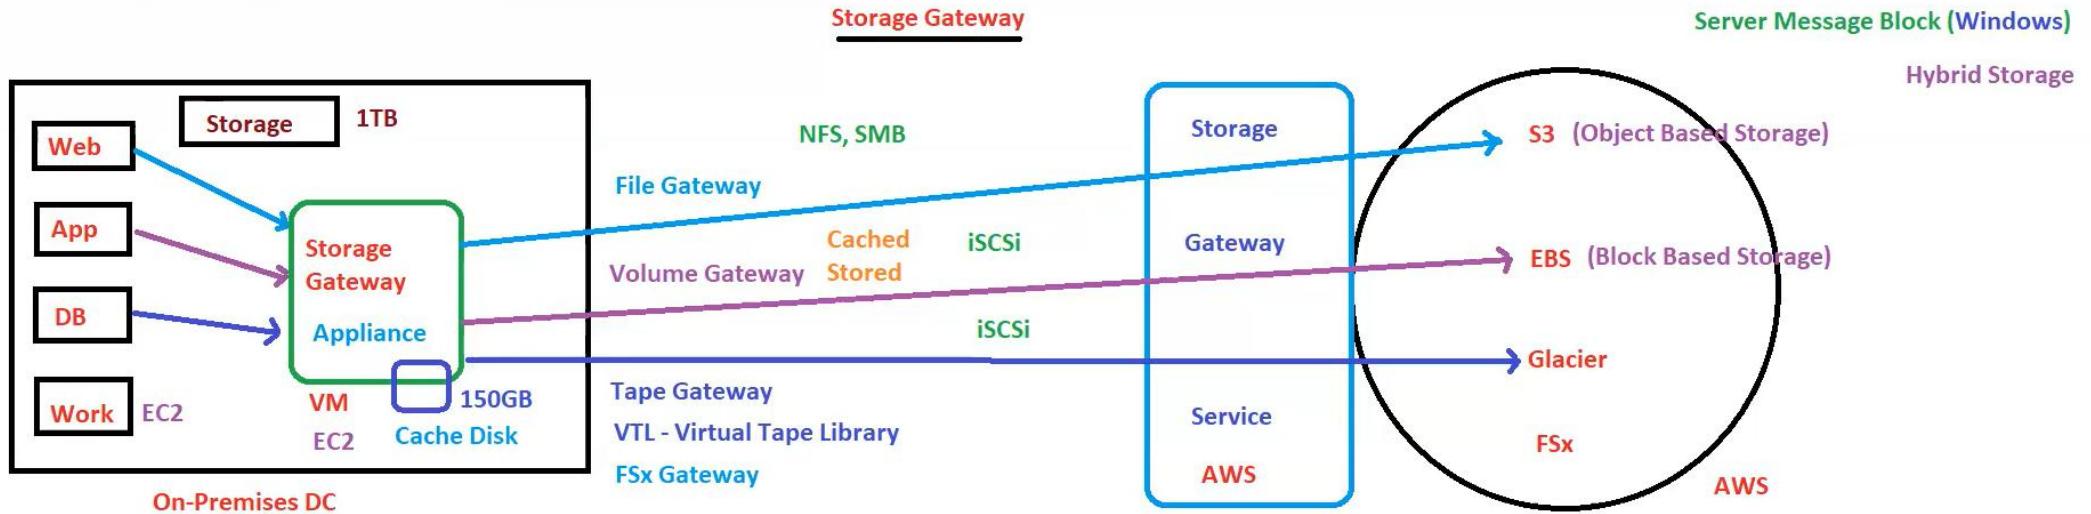
Open Command Prompt
aws configure
access key
secret key
region : ap-south-1
output = table/json/xml

cd ~/aws
cat credentials







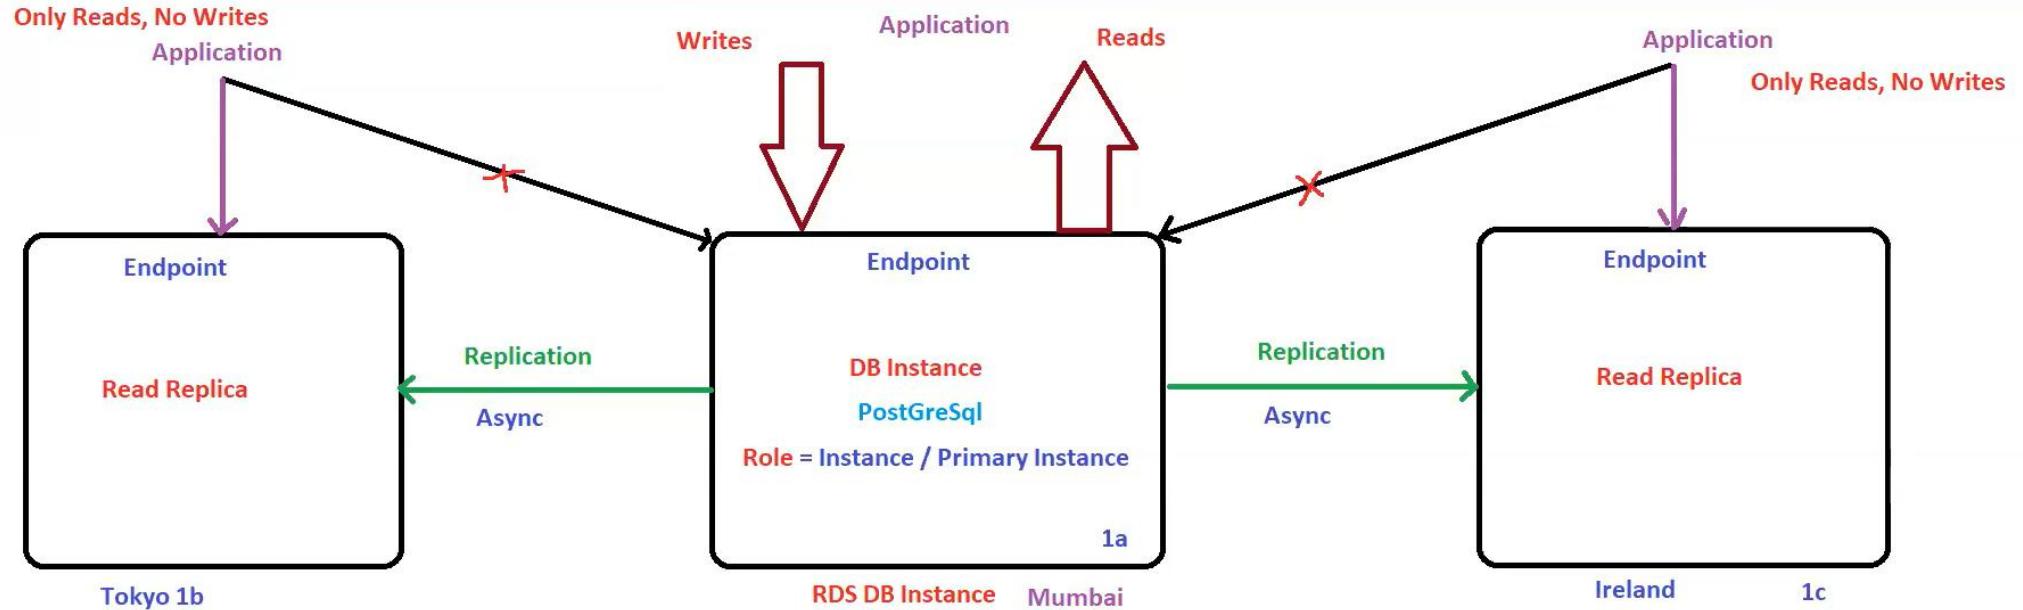


1. Go to AWS Storage Gateway Console and Create File Gateway
2. Launch EC2 instance with m5.xlarge having Root vol 80GB, add vol with 150GB
3. Enable All Traffic for default SG(Temporary, in real time use necessary protocols)
4. Continue on Storage Gateway Service Console, Connect to the Gateway: Provide EC2 instance Public IP (In real time, use EIP and VPC Hosted)
5. Wait until to get gateway active and allocate additional 150GB as Cache Storage (no log group, no alarms)
6. Create a S3 Private Bucket
7. In Storage Gateway Service Console, Create a File Share and while creating file share choose NFS, Allow Client 0.0.0.0/0, and rest keep it defaults

8. Launch RedHat EC2 instance t2.micro for testing

```
sudo -s
yum install -y nfs-utils
mkdir gateway
mount -t nfs -o nolock,hard 172.31.26.106:/6pm-filegateway-bkt gateway/
cd gateway
touch hello.txt
```

Check hello.txt in S3



Application

HostName / Endpoint
Username
Password
Port Number

Endpoints are provided by AWS

RDS is Regional

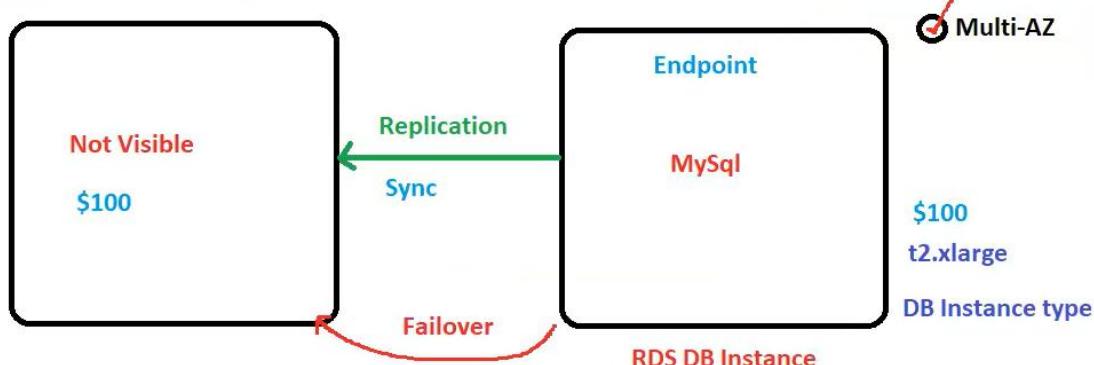
Read Replicas are used to increase the Performance
Max 15 Read Replicas
RR can be in diff AZ and cross region too
RR can be used for Read Purpose not Write Purpose
RR has its own Endpoints

It is possible to promote Read Replica to normal standalone instance

Multi-AZ
Multi-AZ is for high availability Purpose

Backups ---> Snapshots

Manual
Automatic (Schedule)



DB Operations will not have any failover , anything related to network, servers etc will have failover.

In Multi-AZ endpoints will not change after failover

You can enable Multi-AZ for RR also

For Multi-AZ charges are doubled

RDP, SSH are not possible to the RDS DB Instance

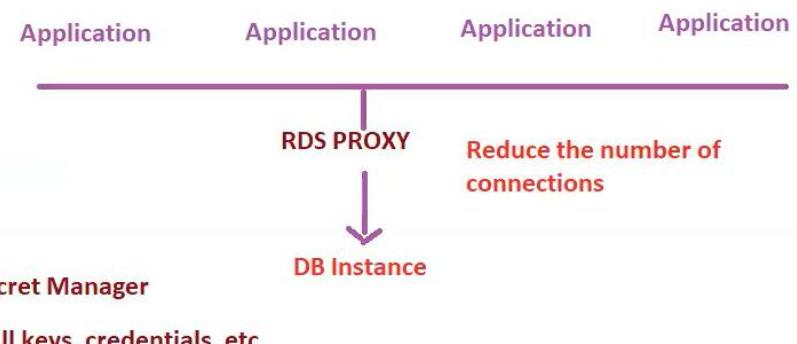
RDS DB Instances can be reserved

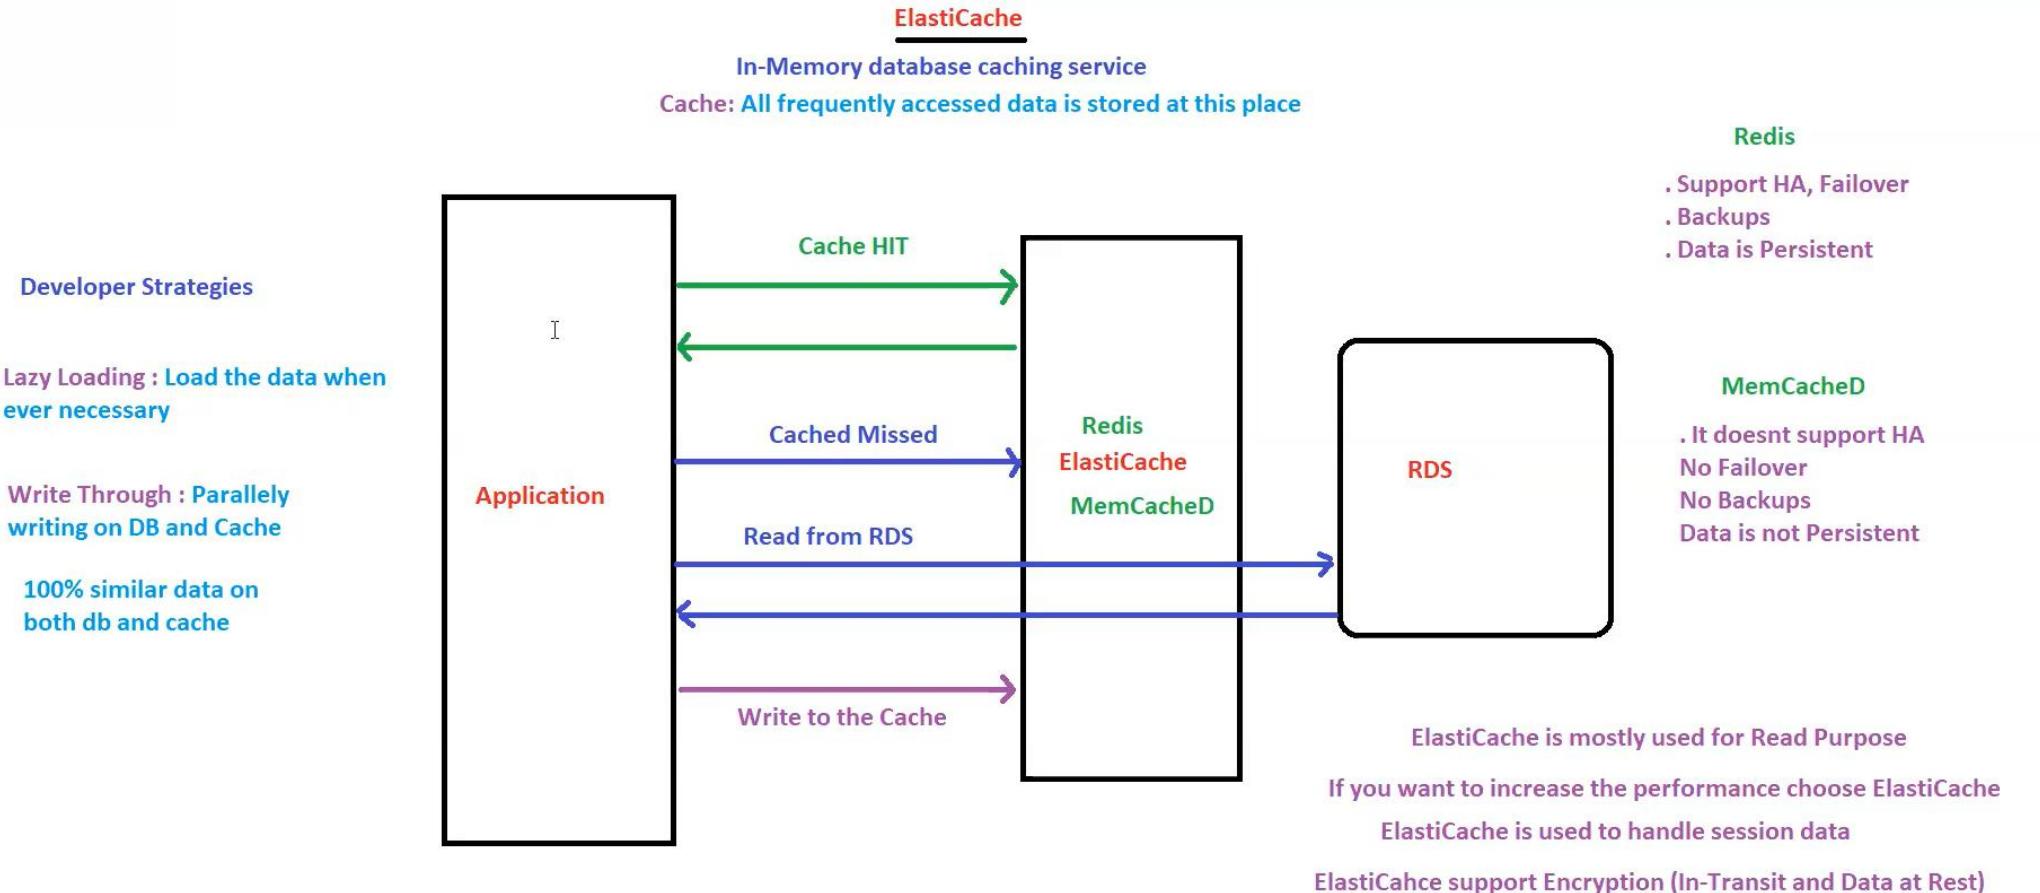
Parameter Group

Subnet Group

AWS Secret Manager

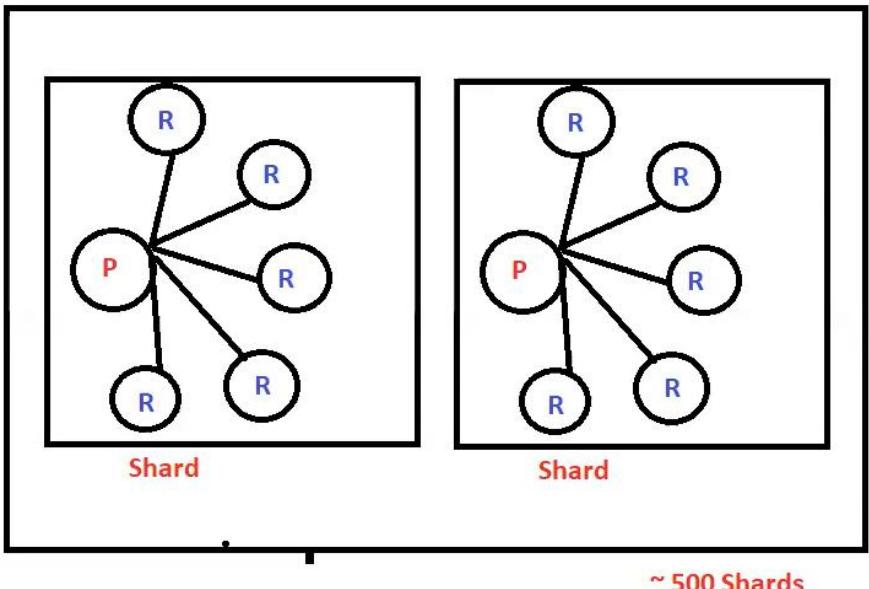
All keys, credentials, etc



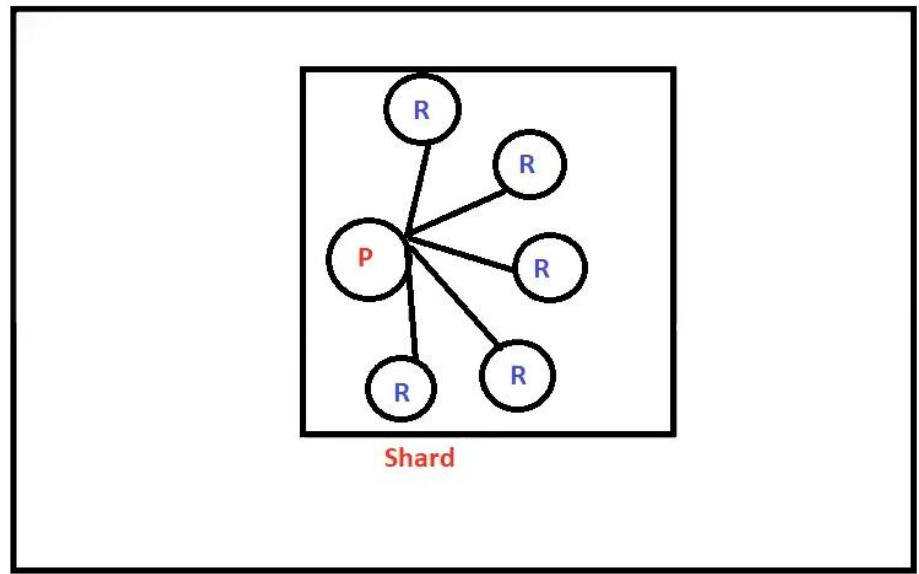


Redis

Cluster Mode Enabled



Cluster Mode Disabled



Shard = Collection of Nodes / Servers

Cluster = Collection of Shards

Each shard has 6 Nodes = 1 Primary , 5 Replica Nodes

App code

```
save_user(17, {"name": "Kevin McGehee"})  
user = get_user(17)
```

Write Through

```
def save_user(user_id, values):  
    record = db.query("update users ... where id = ?", user_id, values)  
    cache.set(user_id, record, 300) # TTL  
    return record
```

Lazy Load

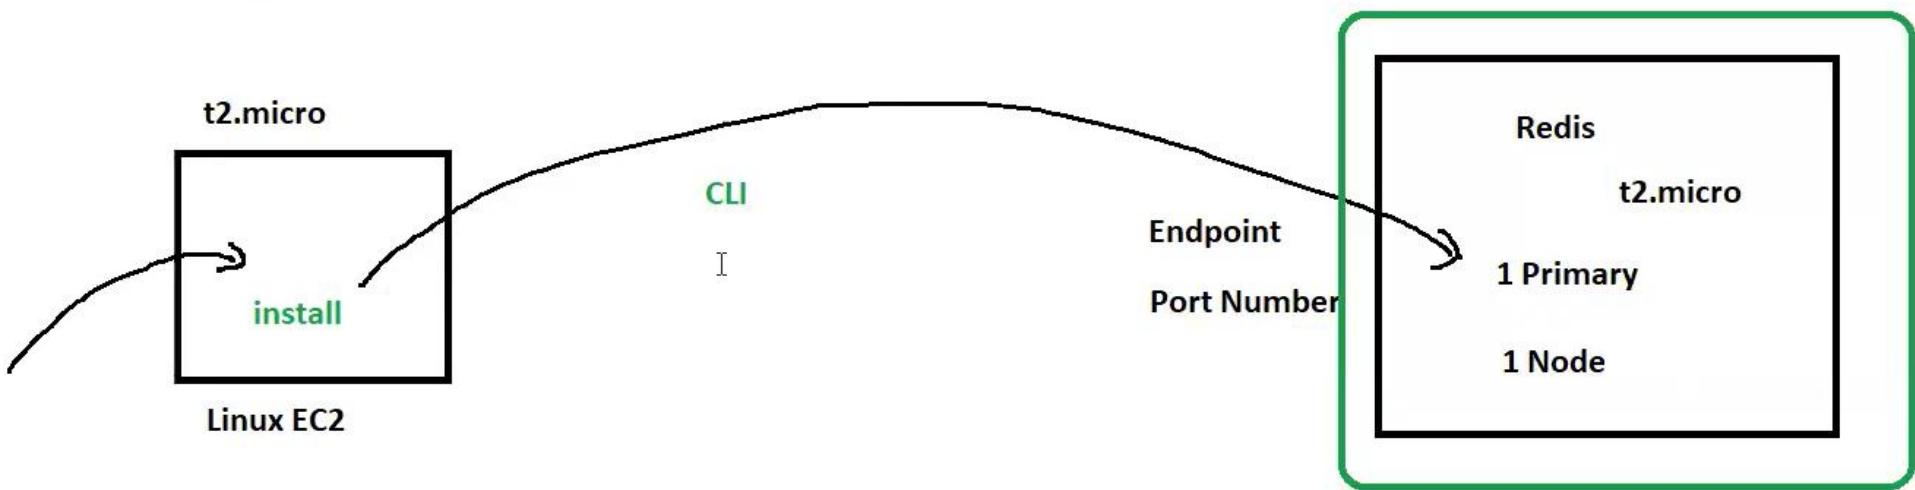
```
def get_user(user_id):  
    record = cache.get(user_id)  
    if record is None:  
        record = db.query('select * from users where id = ?', user_id)  
        cache.set(user_id, record, 300) # TTL  
    return record
```

Write Through

1. Updated DB
2. SET in Cache

Lazy Load

1. GET from cache.
2. If MISS get from DB
3. Then SET in Cache





AWS DynamoDB

This slide is a presentation slide for AWS DynamoDB, featuring a large central image of a blue cylinder representing the database, followed by six numbered sections on the left side.

- 1 **INTRODUCTION**
AWS DynamoDB
- 2 **Traditional Architecture**
 - Traditional applications leverage RDBMS databases
 - These databases have the disadvantage of being monolithic, inflexible, who decide who has the right to access the data
 - Many requirements about how the data should be modeled/transformed
 - Ability to do joins, aggregations, computation
 - Vertical scaling (means usually getting a more powerful CPU/RAM)
- 3 **NoSQL databases**
 - The majority of NoSQL databases are built with a schema-free approach
 - NoSQL databases are designed to handle large amounts of data
 - To solve the problems, we could "scale up" our systems by expanding our existing hardware. This presents a significant challenge because it requires us to constantly add more hardware to our system to handle the increasing load of data
 - Another approach is to use a distributed database system, which allows us to distribute the data across multiple machines
 - Examples of distributed databases include Google BigTable, Amazon DynamoDB, and Apache HBase
 - Benefits of distributed databases:
 - Scalability
 - High availability
 - Data consistency
 - Data partitioning
- 4 **NoSQL databases**
 - NoSQL means, Not Only SQL
 - NoSQL databases are non-relational databases (not fixed column sets)
 - NoSQL databases include MongoDB, DynamoDB etc
 - NoSQL databases do not support joins
 - NoSQL databases are good for handling unstructured data, such as JSON or XML
 - NoSQL databases don't perform aggregations such as "SUM"
 - NoSQL databases scale horizontally
- 5 **SQL vs NOSQL databases**
 - SQL databases are designed for structured data, while NoSQL databases are designed for unstructured data
 - SQL databases are relational, meaning they store data in tables, while NoSQL databases are non-relational, meaning they store data in documents, key-value pairs, or graphs
 - SQL databases use ACID transactions, while NoSQL databases use eventual consistency
 - SQL databases are good for transactional processing, while NoSQL databases are good for analytical processing
 - SQL databases are good for structured data, while NoSQL databases are good for unstructured data
- 6 **NoSQL databases**
 - There are four main types of NoSQL databases:
 - Document (MongoDB)
 - Key-Value (Redis, Memcached, Amazon DynamoDB)
 - Column Family (Apache HBase, Amazon SimpleDB)
 - Graph (Neo4j, Amazon Neptune)

Traditional Architecture



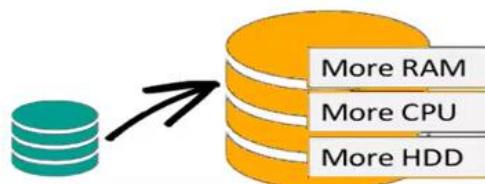
- Traditional application leverage RDBMS databases
- These databases have the SQL query language
- Strong requirements about how the data should be modeled(tables, rows, columns, joins, indexes, etc)
- Ability to do join, aggregations, computations
- Vertical scaling (means usually getting a more powerful CPU/RAM/IO)

NoSQL databases

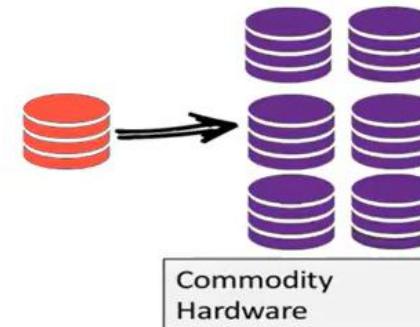


- The concept of NoSQL databases became popular with Internet giants like Google, Facebook, Amazon, etc. who deal with huge volumes of data.
- The system response time becomes slow when you use RDBMS for massive volumes of data.
- To resolve this problem, we could "scale up" our systems by upgrading our existing hardware. This process is expensive.
- The alternative for this issue is to distribute database load on multiple hosts whenever the load increases. This method is known as "scaling out."

Scale-Up (vertical scaling):



Scale-Out (horizontal scaling):





NoSQL databases

- NoSQL means, Not Only SQL.
- NoSQL databases are non-relational databases(no fixed columns etc) and are distributed(horizontal scaling).
- NoSQL databases include MongoDB, DynamoDB etc
- NoSQL databases do not support joins
- All the data that is needed for a query is present in one row
- NoSQL databases don't perform aggregations such as “SUM”.
- NoSQL databases scale horizontally

SQL vs NOSQL databases



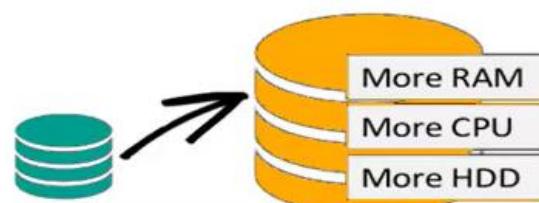
SQL	NoSQL
SQL is generally used in RDBMS	NoSQL is used for Non-Relational database system
Structured data can be stored in tables	Using JSON data, unstructured data can be stored
The Schemas are static	The Schemas are dynamic
Schemas are rigid and bound to relationships	Schemas are non-rigid, they are flexible
Helpful to design complex queries	No interface to prepare complex queries
Here we call tables, rows and columns	Here we call collections, collections has documents
MySQL, Oracle, Sqlite, Postgres and MS-SQL	MongoDB, BigTable, Redis, RavenDb, Cassandra, Hbase, Neo4j and CouchDb

NoSQL databases

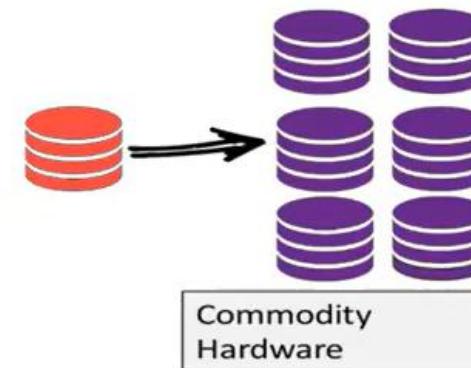


- The concept of NoSQL databases became popular with Internet giants like Google, Facebook, Amazon, etc. who deal with huge volumes of data.
- The system response time becomes slow when you use RDBMS for massive volumes of data.
- To resolve this problem, we could "scale up" our systems by upgrading our existing hardware. This process is expensive.
- The alternative for this issue is to distribute database load on multiple hosts whenever the load increases. This method is known as "scaling out."

Scale-Up (vertical scaling):



Scale-Out (horizontal scaling):



SQL vs NOSQL databases

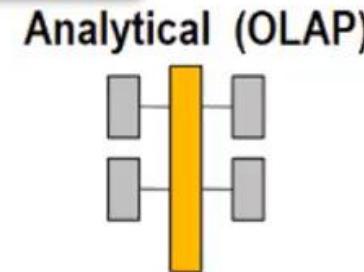
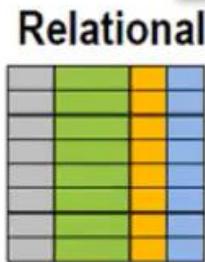


SQL	NoSQL
SQL is generally used in RDBMS	NoSQL is used for Non-Relational database system
Structured data can be stored in tables	Using JSON data, unstructured data can be stored
The Schemas are static	The Schemas are dynamic
Schemas are rigid and bound to relationships	Schemas are non-rigid, they are flexible
Helpful to design complex queries	No interface to prepare complex queries
Here we call tables, rows and columns	Here we call collections, collections has documents
MySQL, Oracle, Sqlite, Postgres and MS-SQL	MongoDB, BigTable, Redis, RavenDb, Cassandra, Hbase, Neo4j and CouchDb

NoSQL databases

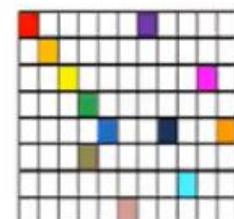


SQL Database

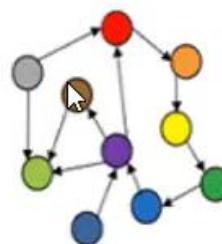


NoSQL Database

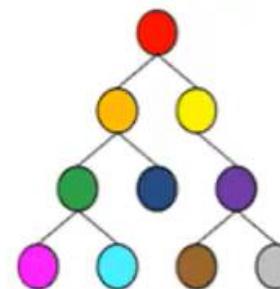
Column-Family



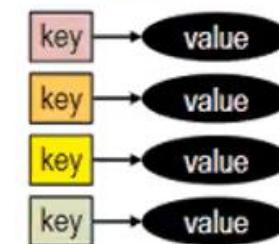
Graph



Document



Key-Value



DynamoDB vs RDS



- DynamoDB offers “push button” scaling , meaning that you can scale your database on the fly, without any down time.
- DynamoDB is Serverless
- RDS is not so easy and you usually have to use a bigger instance size or to add a read replica.

Basics of DynamoDB

- Stored on SSD storage
- Spread across 3 geographically distinct data centers
- Eventual Consistent **Reads**(Default)
- Strongly Consistent **Reads**

Attributes

PK	Columns	SK	Employee Table	
Items →	EmpNo	EmpName	Phone	Salary
Rows	1000	Hari	984651461	20L
TTL = 30 days	1001	Prasad	84588945	30L
	1002	Madhavi	986556213	40L
	1003	Nishan	8787784584	50L
	1000	Manoj	9845732345	

Table = Attributes + Items

Consistency Models



Primary Key = The value should be always unique and not null

Primary Key = Partition Key

In DynamoDB, initially you create table with only one single column, and that column should be Partition Key

Primary Key is mandatory in DynamoDB while creating a table

Composite Key = Partition Key + Sort Key

Sort Key is Optional, PK is Mandatory

If you have a primary key for a single column = Duplicate values are NOT allowed

If you have a PK + SK = Primary Key duplicate values are Allowed

EmployeeNumber = Primary Key / Hash Attribute

Phone Number = Sort Key / Range Attribute

Document
Key-Value

DynamoDB data is stored in JSON format

DynamoDB Streams = Change Logs

Kinesis Data Streams

Captures item level changes

Indexes

It helps to increase the performance on the table / Retrieving the data

DynamoDB performance depend upon RCU and WCU

LSI (Local Secondary Index) = Partition key + Any Column as SK

5 RCU, 5 WCU --> default

GSI (Global Secondary Index) = Any Column as Partition key + Any Column as SK

**1 RCU = 5.2 million reads
1 WCU = 2.5 million writes**

LSI can be created only at the time of creating dynamodb table, later you cannot modify or delete it

GSI can be created, modified and delete anytime

Provisioned Capacity Units

Read Capacity Units

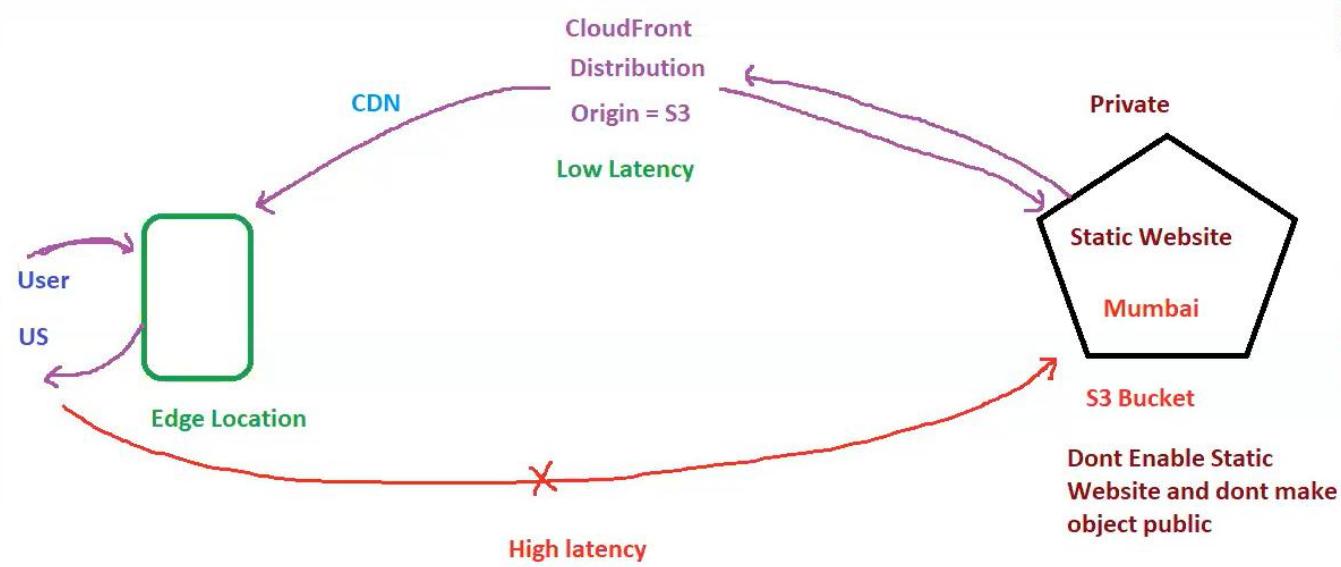
RCU

Write Capacity Units

WCU

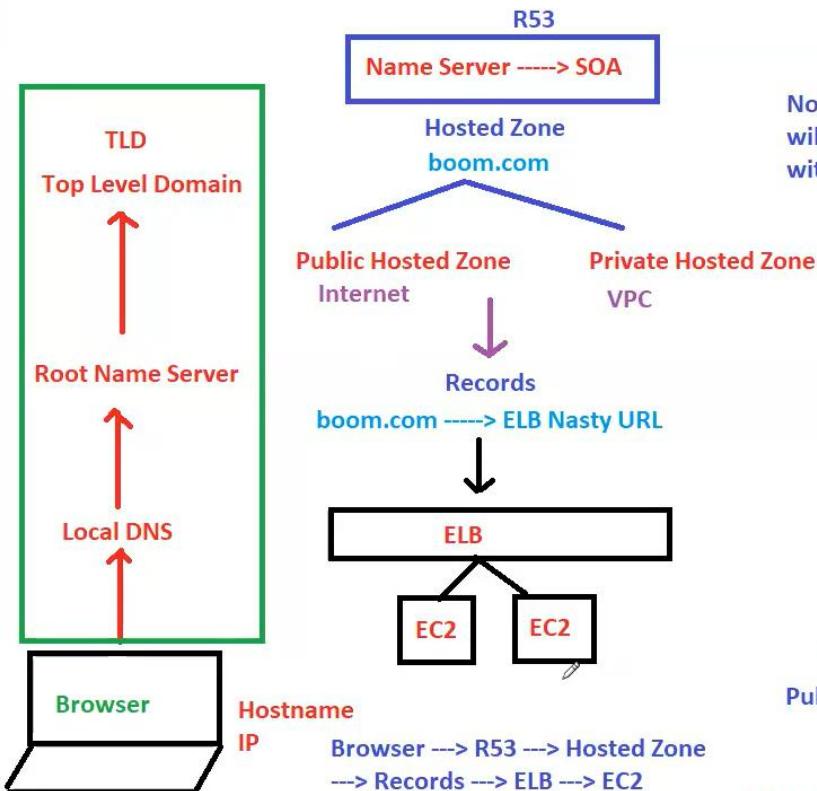
**For ECR, 1 RCU = 2 Reads Per Second for 4KB Size
For SCR, 1 RCU = 1 Reads Per Second for 4KB Size**

1 WCU = 1 Write Per Second for 1KB Size



1. Create a Private S3 Bucket
 2. Upload a Static Website into S3
 3. Dont enable static website option and dont make the objects public
 4. Create CloudFront Distribution
 - Origin ---> S3 Bucket
 - OAC --> Create origin access setting
root object = index.html
- rest everything leave it defaults
- Copy the policy and update in S3 bucket policy
Access the website with CloudFront URL

DNS = Keep tracks of all hostnames and IP address
= It converts hostname to ip and ip to hostname



Route53
DNS Service from AWS, DNS Port Number is 53

Purchase Domain

Route53
No Need to create HZ, it will create automatically with NS and SOA Records

Go-Daddy (Any)
Manually create HAZ in R53, it will automatically create NS and SOA Records

Route53 Features

1. Domain Registration
2. DNS Routing
3. Health Checks
4. Routing Policies



Route53 is Global

DNS is all about Records

Hosted Zone = Container of Records

Hosted Zone (Domain Name = boom.com)

Public Hosted Zone

NS Record = Pool of Servers (4 Servers)

SOA Record = Admin of Hosted Zone

You cannot delete NS and SOA Records

NS and SOA are default Records automatically created when hosted zone created

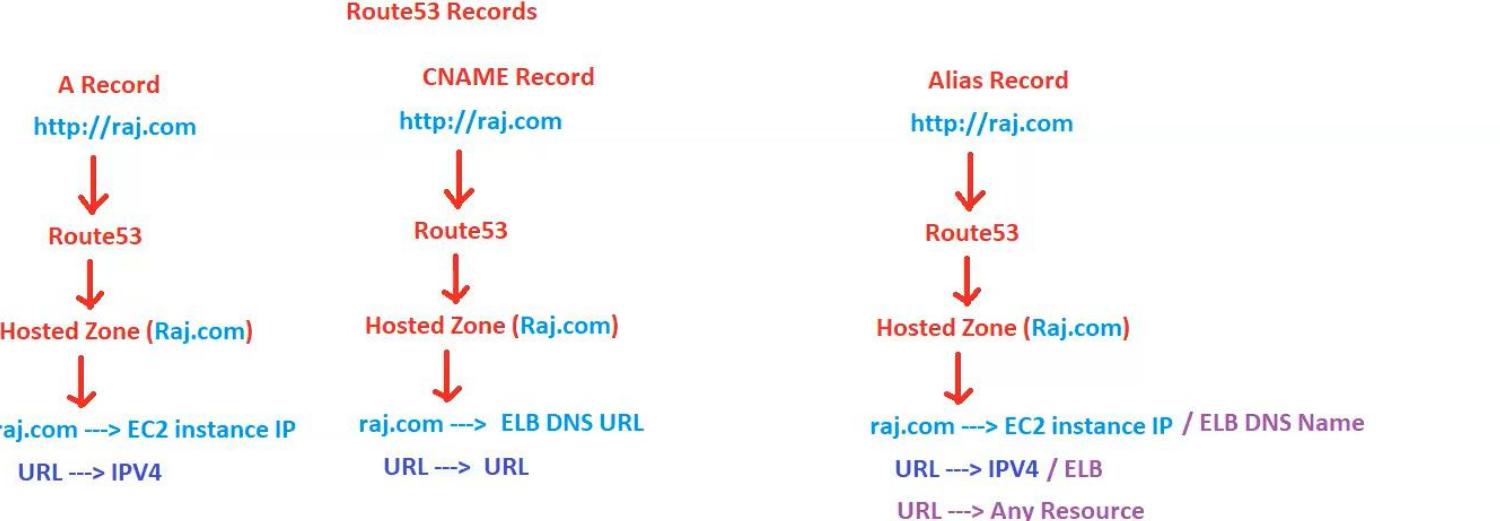
A Record = URL to IPV4

AAAA Record = URL to IPV6 ✗

CName Record = URL to URL

Alias Record = URL to Any Resource

MX Record = Emails



`http://raj.com`

Main Domain
Naked Domain
Zone Apex Record

`admin.raj.com`

`hello.raj.com` ---> Sub-Domains
`photos.raj.com`

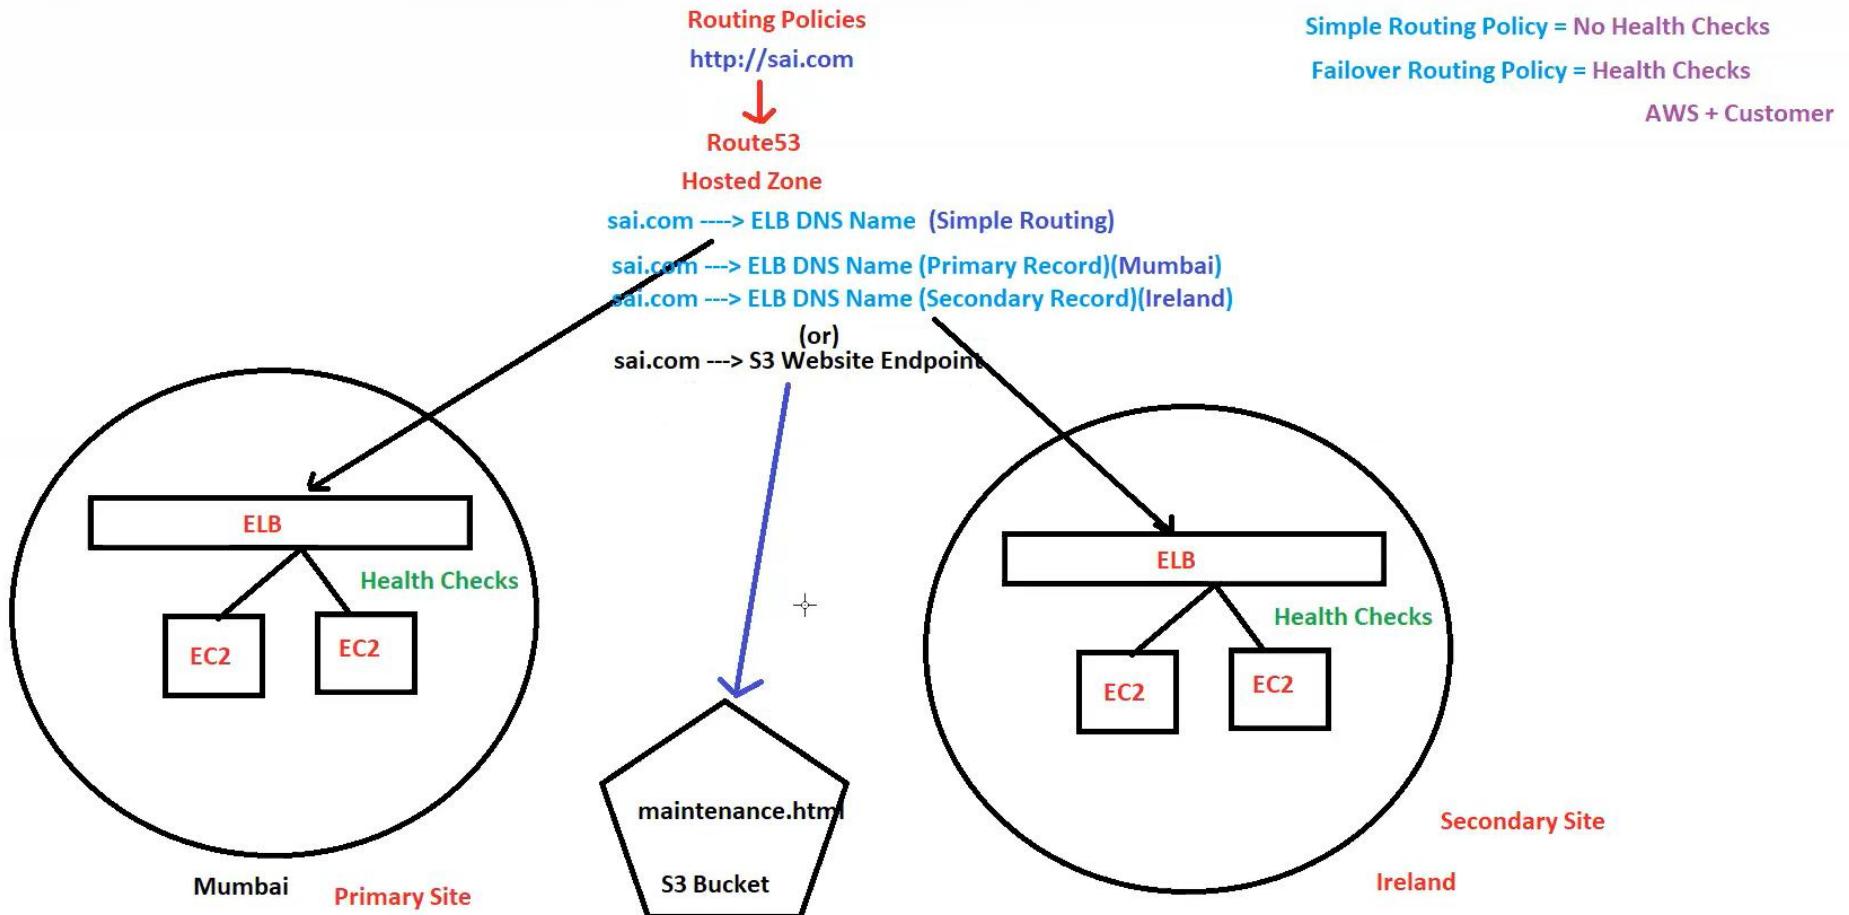
CNAME Records are Billable, Alias Records are FREE

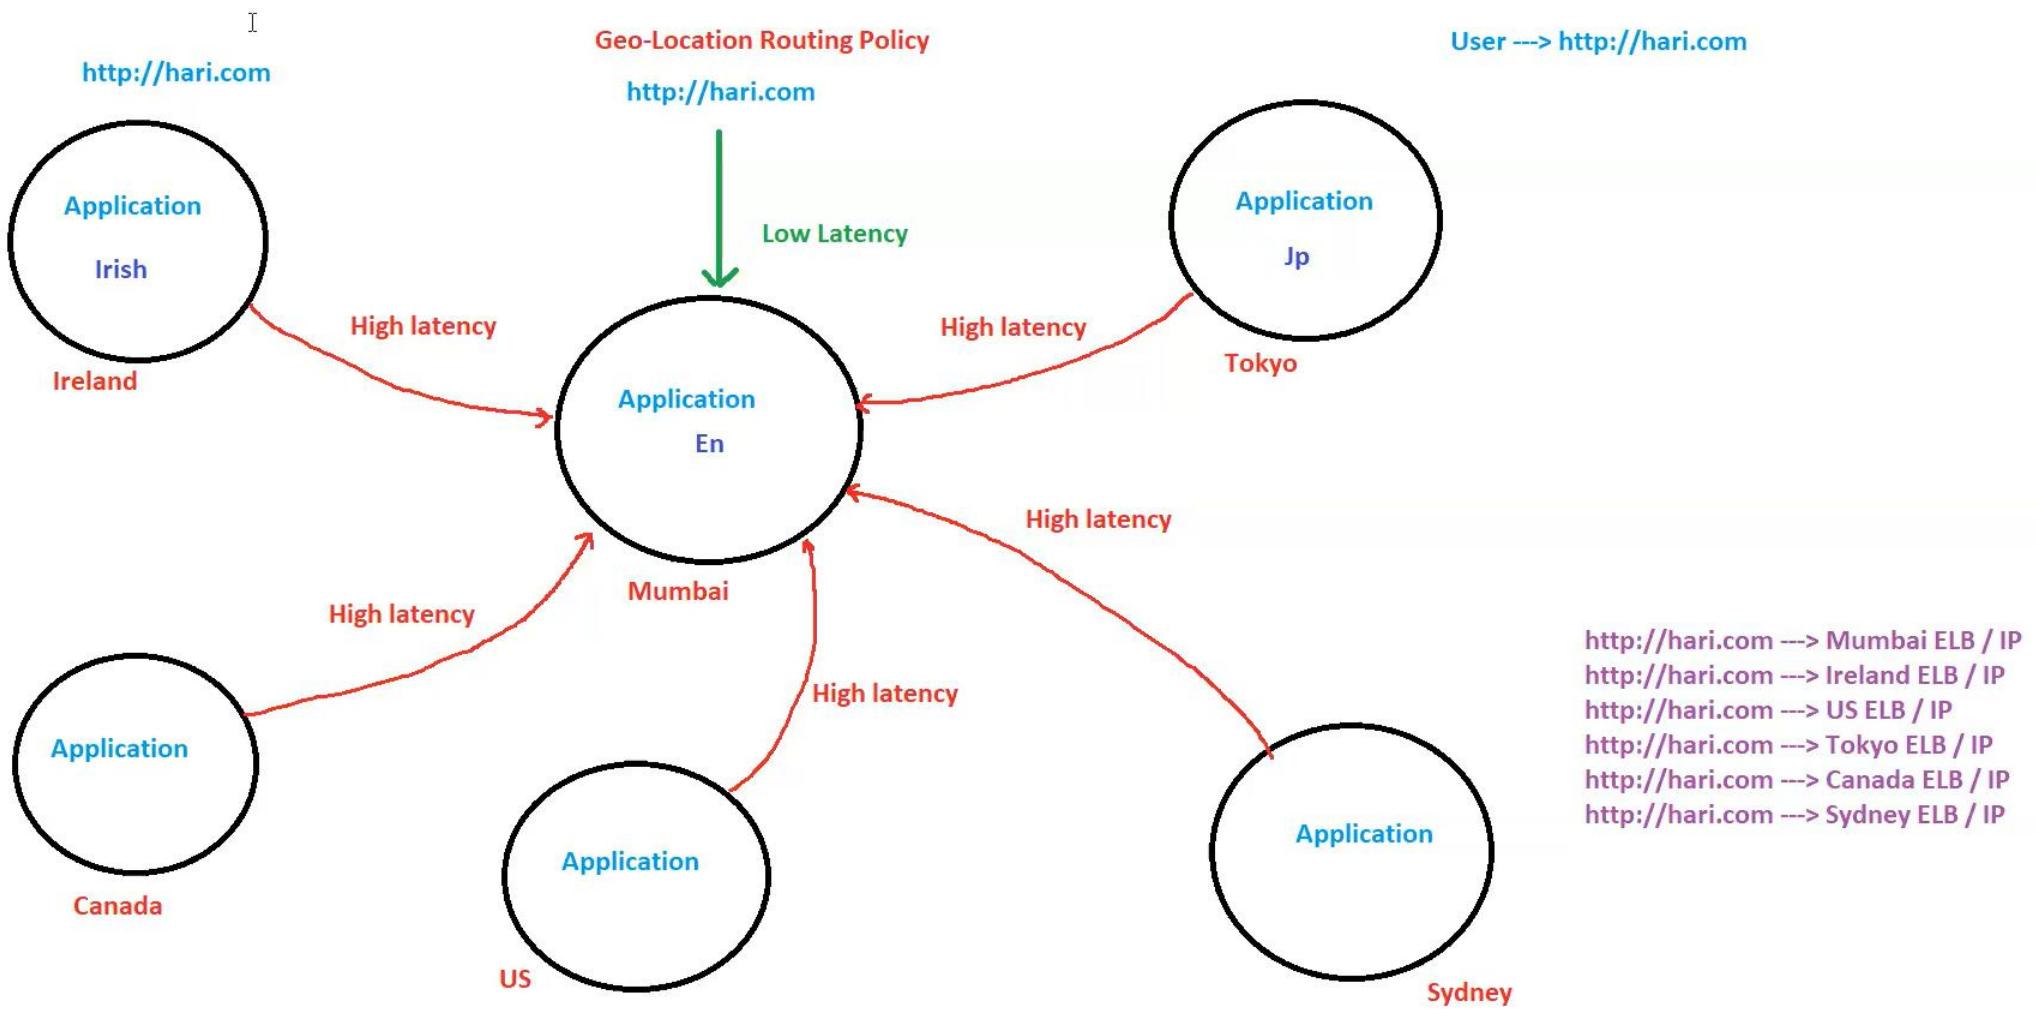
For Naked Domain, `raj.com` cannot use CNAME records, instead use Alias

For all Sub-Domains you can use CNAME

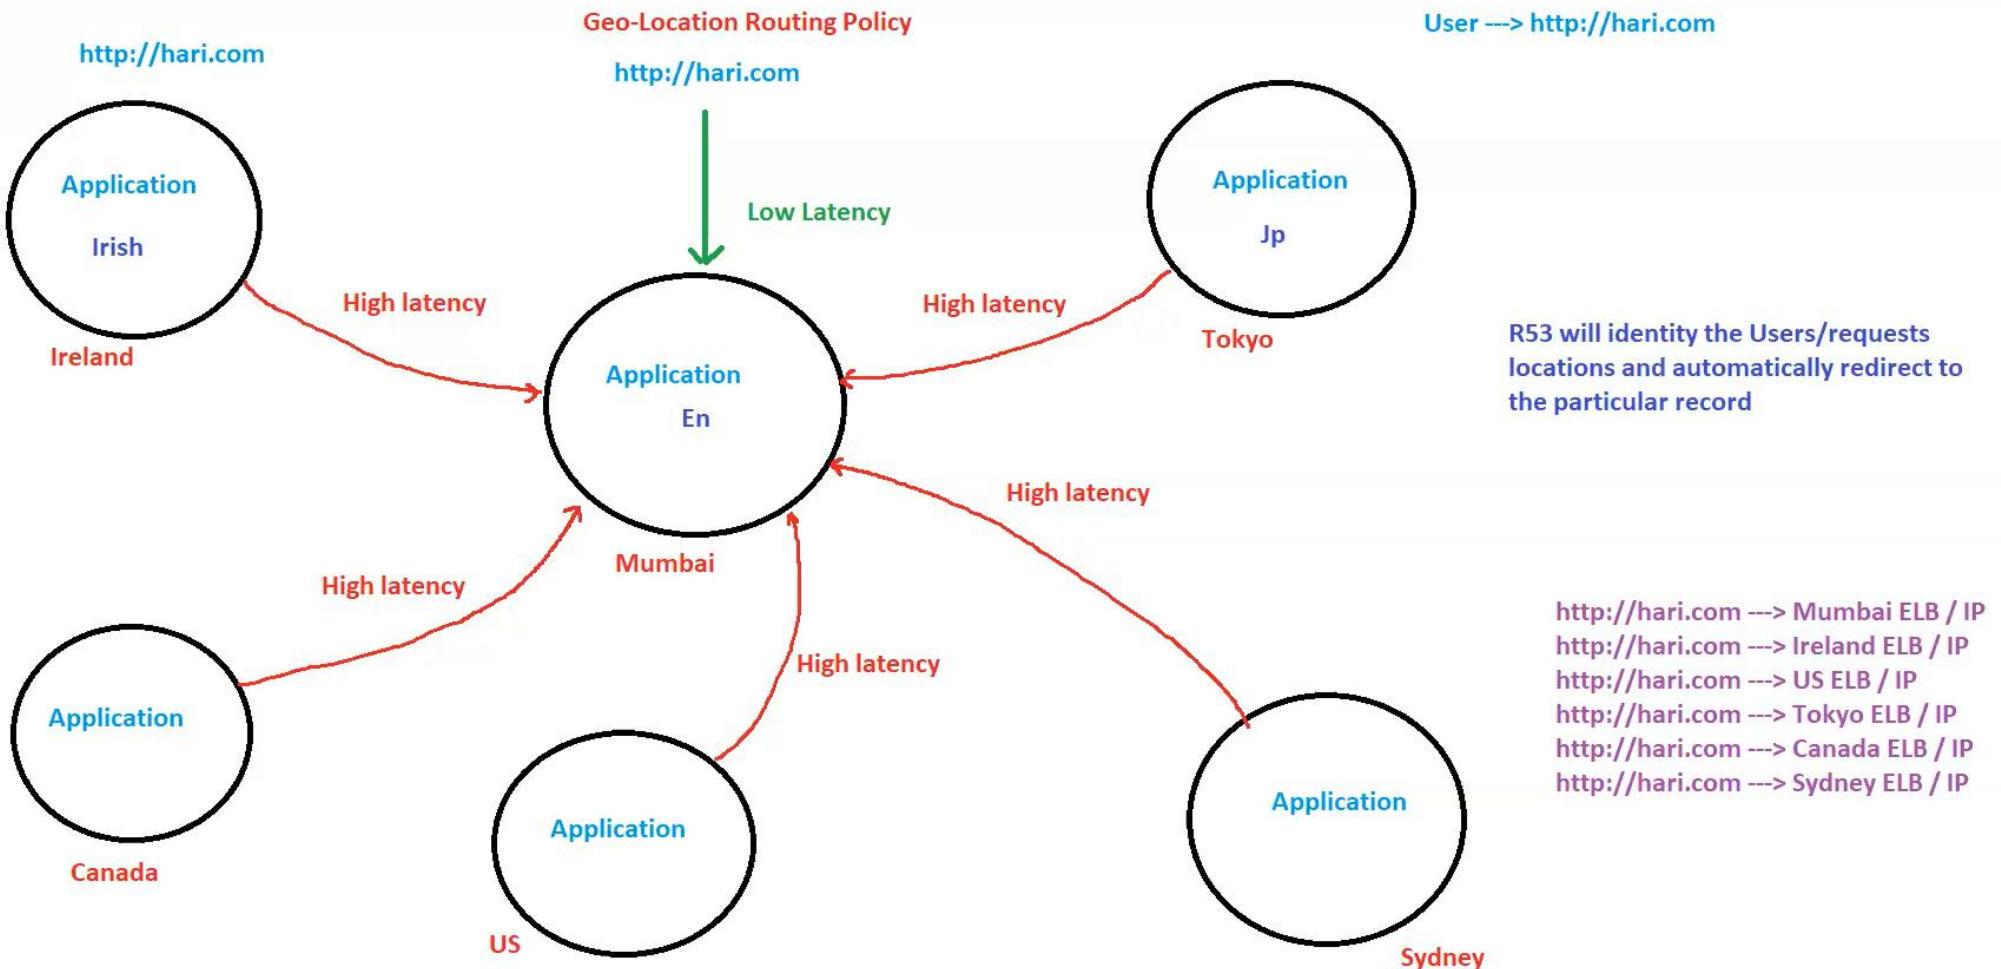
Always choose Alias over CNAME

A Record + Alias Record





<http://hari.com> ---> Mumbai ELB / IP
<http://hari.com> ---> Ireland ELB / IP
<http://hari.com> ---> US ELB / IP
<http://hari.com> ---> Tokyo ELB / IP
<http://hari.com> ---> Canada ELB / IP
<http://hari.com> ---> Sydney ELB / IP



VPC
Internet Gateway
Public and Private Subnet
NAT Gateway
Router with Routing tables

IGW ---> Provides internet to VPC

Public Subnet --> which is exposed to internet

3

Private Subnet --> which is NOT exposed to internet

Public Subnet traffic is routed to IGW

Private Subnet traffic is routed to NAT Gateway

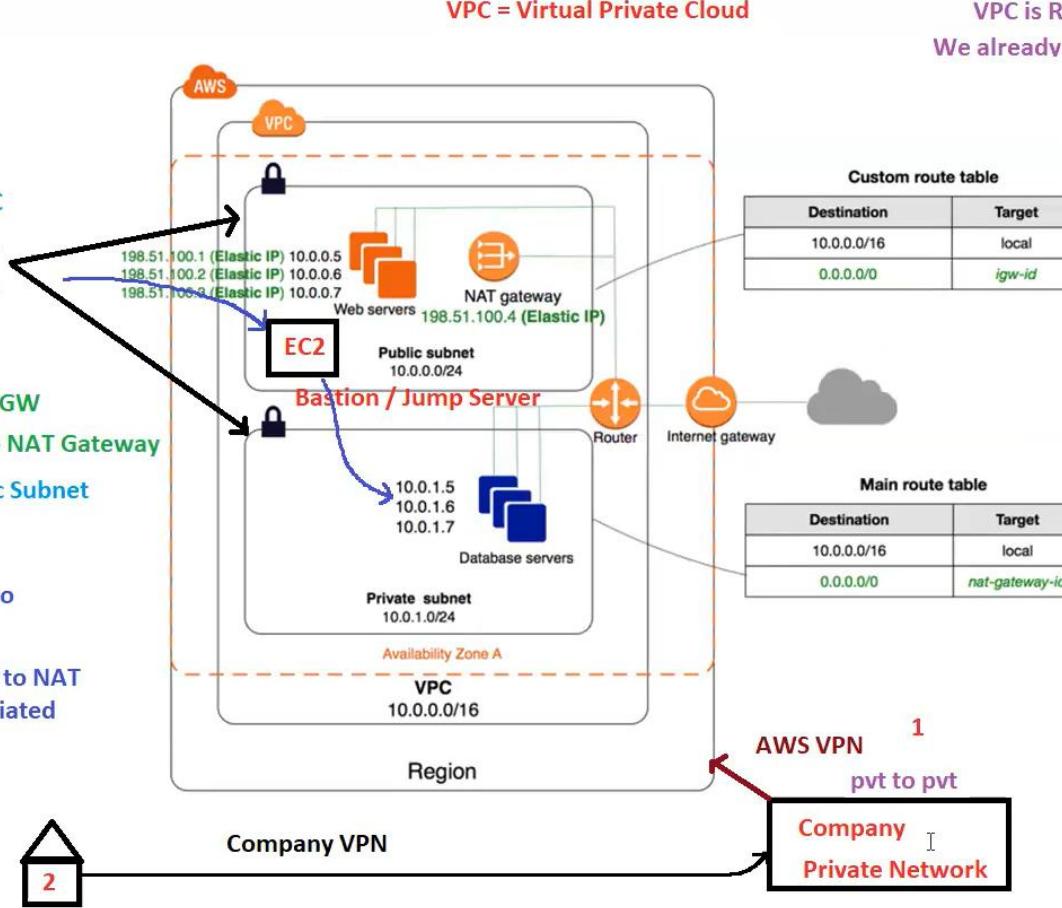
NAT Gateway should be in Public Subnet

Router

Public RT --> All traffic is routed to IGW, Public Subnet is associated

Private RT --> All traffic is routed to NAT Gateway, Private Subnet is associated

1. Actual Real time
2. Work from home
3. Learning



VPC is Regional, Max 5 VPC's per Region
We already have a default VPC provided by AWS

NAT = Network Address Translator

NAT is used to provide internet access to Private Subnet

NAT will convert Private IP to Public IP

1. Create VPC
2. Create IGW and attach it to the VPC
3. Create Public Subnet
4. Create Private Subnet
5. Create NAT Gateway (Public Subnet, EIP)
6. Create Public RT --> All Traffic is routed to IGW, Public Subnet is associated
7. Create Private RT --> All Traffic is routed to NAT Gateway, Private Subnet is associated
8. Create a Security Group
9. Launch EC2 instance in Public Subnet (Bastion Server) and another in Private subnet (PrivateServer)
10. Connect to Bastion Server and then to Private Server

NAT and IGW are Services not Servers

NAT Instance ---> Outdated

NAT Gateway --> Managed by AWS

VPC Endpoints: It is used to access only AWS Services without NAT

S3 VPCE
DynamoDB VPCE
etc

NAT = Full internet access

Private Purpose

192
172
10

Rest are all used for Public IP's

1 Subnet should be in 1AZ at the same time

1 AZ can have multiple Subnets

2^{32-n}

$2^{32-24} = 2^8 = 256 - 5 = 251$

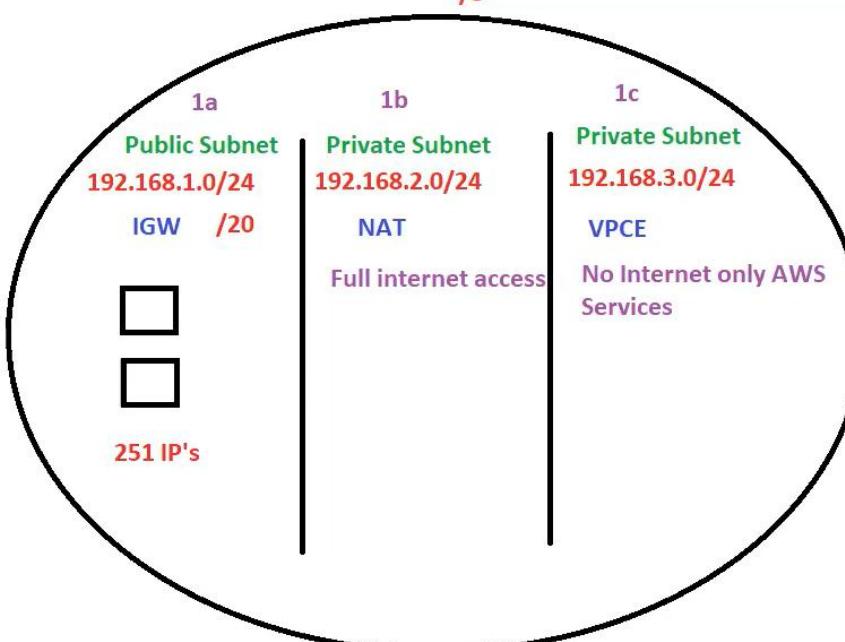
$2^{32-20} = 2^{12} = 4096 - 5 = 4091$

$2^{32-16} = 2^{16} = 65536 - 5 = 65531$

$2^{32-32} = 2^0 = 1$

CIDR = Classless InterDomain Routing

Base IP Subnet Mask
VPC (192.168.0.0/16) /8



5 IP's are Reserved for each Subnet

.0 = Network
.1 = Router
.2 = DNS
.3 = Future
.255 = BroadCasting

1. Create VPC (192.168.0.0/16)
2. Create IGW and attach it to the VPC
3. Create Public Subnet (192.168.1.0/24)
4. Create Private Subnet (192.168.2.0/24)
5. Create NAT Gateway (Public Subnet, EIP)
6. Create Public RT ---> 0.0.0.0/0 routed to IGW, Public Subnet is associated
7. Create Private RT ---> 0.0.0.0/0 is routed to NAT Gateway, Private Subnet is associated
8. Create a Security Group
9. Launch EC2 instance in Public Subnet (Bastion) and another in Private Subnet (Private Server)
10. Connect to bastion and then to Private Server

In Bastion Server

```
sudo -s  
vi 6PMBATCHNEW.pem  
---> press i  
---> copy paste the pem file data  
---> press esc, :wq!
```

```
chmod 400 6PMBATCHNEW.pem  
ssh -i "6PMBATCHNEW.pem" ec2-user@192.168.2.11
```

from here you are in mumbaiprivate server

```
sudo -s  
Try internet is working or not by installing a sample command yum install -y wget  
Switch OFF internet by removing the NAT entry in Private RT  
Switch ON internet by adding the NAT entry in Private RT  
Install AWS CLI  
Create a IAM ROLE with S3 Permissions and attach ROLE to the EC2 instance  
Create a sample S3 bucket from AWS CLI (MumbaiPrivate Server) [this got created because we have NAT entry]  
Remove NAT entry in Private RT  
List S3 bucket from AWS CLI --> cannot access because no internet as you removed NAT entry  
Create a VPCE  
And List S3 bucket [this will work because we have now VPCE]
```

There are three types of VPC endpoints – Interface endpoints, Gateway Load Balancer endpoints, and Gateway endpoints.

